

200+ Backend Interview Problems

Curated by

[Mission_Compile \(Bhavesh Vaswani\)](#)

Insta ID: mission_compile

Copyright © 2025 Bhavesh Vaswani. All rights reserved.

No part of this document may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without prior written permission of the author.

A brief about the document:

This is not just a list of questions—it is a practical guide to help you learn, prepare, and master backend engineering interview concepts with real-world scenarios.

- Contains **200+ backend interview problems** covering system design, databases, caching, message queues, APIs, microservices, security, networking, data processing, Kubernetes, cloud architecture, and more.
- Covers **50+ essential topics** including horizontal/vertical scaling, load balancing, CAP theorem, distributed systems, authentication, GraphQL, Redis persistence, SQL injection prevention, HTTP/2 vs HTTP/3, distributed locks, ETL pipelines, load testing, blue-green deployments, chaos engineering, and deployment strategies.
- Detailed **problem-solution format** with step-by-step explanations, real-world examples, and color-coded keywords for quick reference.
- Organized by topics with **clear visual separation**, making it easy to navigate and revise specific concepts.
- Includes practical scenarios like **when to use Redis vs Kafka**, how to implement circuit breakers, database sharding strategies, and API security best practices.

- Contains **trade-off discussions** to help you make informed architectural decisions during interviews.
- Real-world examples showing **how companies like Netflix, Amazon, and Uber solve similar problems at scale**.

Perfect for preparing backend engineering interviews, revising system design concepts, or learning distributed systems from scratch! 🚀



Table of Contents

- [Horizontal vs Vertical Scaling](#) - 3 problems
- [Load Balancing Strategies](#) - 4 problems
- [High Availability & Fault Tolerance](#) - 3 problems
- [CAP Theorem](#) - 2 problems
- [Idempotency](#) - 3 problems
- [Circuit Breaker Pattern](#) - 2 problems
- [Rate Limiting & Throttling](#) - 3 problems
- [Graceful Degradation](#) - 2 problems
- [Distributed Tracing](#) - 2 problems
- [Retry Logic & Exponential Backoff](#) - 2 problems
- [Database Indexing & Query Optimization](#) - 4 problems
- [Database Sharding](#) - 3 problems
- [Read Replicas & Replication](#) - 3 problems
- [Connection Pooling](#) - 2 problems
- [ACID Transactions](#) - 3 problems
- [NoSQL Data Modeling](#) - 3 problems
- [Eventually Consistent Systems](#) - 2 problems
- [N+1 Query Problem](#) - 2 problems
- [SQL vs NoSQL Trade-offs](#) - 3 problems
- [Consistent Hashing](#) - 2 problems
- [Caching Strategies](#) - 4 problems
- [Cache-Aside Pattern](#) - 2 problems
- [Distributed Caching \(Redis\)](#) - 3 problems

[Redis Data Structures](#) - 3 problems
[Cache Stampede Problem](#) - 2 problems
[Kafka Architecture](#) - 3 problems
[RabbitMQ Patterns](#) - 3 problems
[Event-Driven Architecture](#) - 3 problems
[Dead Letter Queues](#) - 2 problems
[Kafka vs RabbitMQ](#) - 2 problems
[Message Deduplication](#) - 2 problems
[REST API Best Practices](#) - 4 problems
[API Versioning](#) - 2 problems
[Authentication \(OAuth 2.0, JWT\)](#) - 4 problems
[Authorization \(RBAC\)](#) - 2 problems
[API Security](#) - 3 problems
[Microservices Communication](#) - 3 problems
[Saga Pattern](#) - 2 problems
[API Gateway Pattern](#) - 3 problems
[Serverless Architecture](#) - 3 problems
[Auto Scaling Strategies](#) - 2 problems
[CDN & Edge Caching](#) - 3 problems
[Asynchronous Processing](#) - 2 problems
[Bloom Filters](#) - 2 problems

****Total: 50+ Topics | 200+ Problems****

Problems & Solutions



Horizontal vs Vertical Scaling



Problem 1: When to Scale Horizontally vs Vertically

Question: INTERVIEWER: Your application is running on a single server that's hitting CPU limits. Should you upgrade the server or add more servers?

Solution:

Scaling isn't just about adding resources — it's about choosing the right growth strategy.
Here's when to scale horizontally vs vertically 

1 Vertical Scaling (Scale Up)

 Add more CPU, RAM, or storage to your existing server.

Example: Upgrade from 4-core to 16-core server when database queries need more processing power.

AWS instance resize adding RAM upgrading CPU

2 Horizontal Scaling (Scale Out)

- 👉 Add more servers to distribute load across multiple machines.

Example: 10 servers handling 10K req/s each instead of 1 server struggling with 100K req/s.

Auto-scaling groups Kubernetes pods load balancers

3 When to Choose Vertical Scaling

- 👉 Use when your application can't be distributed (legacy monoliths, single-threaded apps).

Example: Old ERP system that requires all data in one place with high RAM.

Monolithic databases single-instance applications

4 When to Choose Horizontal Scaling

- 👉 Use for stateless applications that can run independently on multiple machines.

Example: REST API servers that don't share state — each handles requests independently.

Microservices stateless APIs containerized apps

5 Vertical Scaling Limits

- 👉 Physical hardware limits — can't infinitely upgrade a single machine.

Example: You hit cloud provider's largest instance type (96 cores, 384GB RAM) — now what?

Hardware constraints cost inefficiency at scale

6 Horizontal Scaling Benefits

- 👉 No theoretical limit — keep adding servers as traffic grows.

Example: Black Friday? Spin up 100 more servers, then scale down after.

Cloud elasticity high availability fault tolerance

7 Database Considerations

- 👉 Databases often need vertical scaling first (faster disk I/O), then horizontal (sharding).

Example: Upgrade to SSD and more RAM before sharding your database.

Database-specific scaling patterns

8 Cost Trade-offs

- 👉 Vertical scaling is expensive at high end; horizontal scaling is more cost-effective long-term.

Example: 1 massive server costs more than 10 medium servers with same total capacity.

TCO analysis cloud pricing models

💡 Problem 2: Scaling a Monolith to Handle 10x Traffic

Question: INTERVIEWER: Your monolithic app handles 10K users smoothly but crashes with 100K users. How do you scale it?

Solution:

Scaling monoliths requires both vertical and horizontal strategies working together.

1 Add Load Balancer

- 👉 Distribute traffic across multiple instances of your monolith.

Example: Like having multiple checkout counters instead of one — customers get served faster.

Nginx HAProxy AWS ALB

2 Stateless Application Design

- 👉 Remove server-side session state so any instance can handle any request.

Example: Like ATMs that work anywhere — your account isn't tied to one specific machine.

JWT tokens Redis sessions database-backed sessions

3 Vertical Scaling as Quick Fix

- 👉 Temporarily upgrade server resources while you work on horizontal scaling.

Example: Like upgrading from a small restaurant to a bigger one before opening franchises.

Instance type upgrade more RAM/CPU

4 Database Connection Pooling

- 👉 Reuse database connections efficiently as you add more app servers.

Example: Like sharing company cars instead of everyone driving separately — fewer resources wasted.

PgBouncer HikariCP connection pool configuration

5 Cache Aggressively

- 👉 Reduce database load by caching frequent queries.

Example: Like keeping frequently asked phone numbers on speed dial.

Redis Memcached CDN for static assets

6 Offload Heavy Tasks

- 👉 Move background jobs to separate workers using message queues.

Example: Like having a kitchen prepare food while waiters serve — specialization improves speed.

Kafka RabbitMQ AWS SQS with worker processes

7 Auto-Scaling Groups

- 👉 Automatically add/remove servers based on traffic patterns.

Example: Like hiring temporary staff during holiday rush and letting them go after.

AWS Auto Scaling Kubernetes HPA

Problem 3: Database That Can't Scale Horizontally

Question: INTERVIEWER: Your legacy database doesn't support sharding. How do you scale it to handle 100x more traffic?

Solution:

When horizontal scaling isn't possible, you optimize what you have and scale reads separately.

1 Vertical Scaling

- 👉 Upgrade to the largest instance your database supports.

Example: Like moving from a sedan to a bus when you can't buy multiple cars.

2 Read Replicas

- 👉 Create read-only copies to offload SELECT queries from the primary.

Example: Like having multiple library copies so more people can read simultaneously.

PostgreSQL streaming replication

MySQL replication

3 Aggressive Caching

- 👉 Cache query results to avoid hitting the database entirely.

Example: Like keeping frequently used files on your desktop instead of searching every time.

Redis

application-level cache

4 Query Optimization

- 👉 Add indexes, optimize slow queries, remove N+1 problems.

Example: Like organizing your closet so you find clothes faster instead of buying more closets.

5 Archive Old Data

- 👉 Move historical data to separate storage to keep active database small.

Example: Like storing old emails in archive folders — keeps inbox fast.

Data archival cold storage

6 Connection Pooling

- 👉 Maximize concurrent connections with efficient pooling.

Example: Like a taxi service reusing cars efficiently instead of buying one per customer.

PgBouncer RDS Proxy

7 Partition Tables

- 👉 Split large tables into smaller partitions within same database.

Example: Like organizing files by year so you search smaller folders.

PostgreSQL partitioning range partitioning



Load Balancing Strategies

💡 Problem 1: API Crashes Under 1M Requests Per Second

Question: INTERVIEWER: How do you design an API that survives 1 million requests per second?

Solution:

1 Million RPS isn't about code, it's about architecture.

Here are the 8 key layers to scale your API architecture.

1 Load Balancer

- 👉 Distributes traffic across multiple servers so no single machine melts.

Example: 1M req/s split across 200 servers = only 5k req/s each.

Nginx HAProxy AWS ALB/NLB

2 Horizontal Scaling

- 👉 Add more servers when traffic spikes instead of upgrading one big server.

Example: Black Friday? Spin up 50 more API nodes in seconds.

Auto-scaling groups Kubernetes ECS

3 Caching Layer

- 👉 Serve frequent reads from Redis/Memcached to avoid DB overload.

Example: User profile cached → avoids 10M database hits/day.

Redis Memcached CDN

4 CDN for Static Content

- 👉 Images and static assets load from edge servers near the user.

Example: A user in Delhi gets images from a Delhi CDN node.

CloudFront Cloudflare Akamai

5 Async Processing (Queues)

- 👉 Push heavy tasks to Kafka/SQS so API responds instantly.

Example: Payment API returns fast → receipt email sent in background.

Kafka RabbitMQ AWS SQS

6 Database Sharding

- 👉 Split huge datasets across multiple DB shards to scale reads/writes.

Example: Users A–M on shard 1, N–Z on shard 2.

Vitess Citus application-level sharding

7 Rate Limiting

- 👉 Block or throttle abusive clients to protect server capacity.

Example: "100 requests/sec limit" prevents bots from killing the API.

Redis rate limiting

API Gateway throttling

8 Lightweight Payloads

- 👉 Reduce JSON response size to cut latency and bandwidth.

Example: Return only "id, name, price" instead of 20 unnecessary fields.

Payload optimization

GraphQL for selective queries



Problem 2: Choosing the Right Load Balancing Algorithm

Question: INTERVIEWER: You have 10 backend servers. How does the load balancer decide which server gets the next request?

Solution:

Load balancing algorithms determine how traffic is distributed across servers.

1 Round Robin

- 👉 Distributes requests sequentially to each server in rotation.

Example: Request 1 → Server A, Request 2 → Server B, Request 3 → Server C, repeat.

Simple

works well when all servers have equal capacity

2 Weighted Round Robin

- 👉 Servers with higher capacity get more requests based on assigned weights.

Example: Server A (weight 3) gets 3 requests for every 1 request Server B (weight 1) gets.

Use when servers have different CPU/RAM specs

3 Least Connections

- 👉 Routes requests to the server handling the fewest active connections.

Example: Server A has 50 connections, Server B has 30 → new request goes to Server B.

Best for long-lived connections like WebSockets

4 IP Hash

- 👉 Uses client's IP address to consistently route to same server.

Example: User from IP 192.168.1.1 always goes to Server C (maintains session affinity).

Sticky sessions stateful applications

5 Least Response Time

- 👉 Routes to server with fastest response time and fewest active connections.

Example: Server A responds in 50ms, Server B in 100ms → prefer Server A.

Performance-optimized routing

6 Random

- 👉 Randomly selects a server for each request.

Example: Each request has equal probability of hitting any server.

Simple good for stateless microservices

Problem 3: Layer 4 vs Layer 7 Load Balancing

Question: INTERVIEWER: Should you use a Layer 4 or Layer 7 load balancer for your application?

Solution:

Layer 4 and Layer 7 load balancers operate at different network layers with different capabilities.

1 Layer 4 (Transport Layer)

- 👉 Routes traffic based on IP address and TCP/UDP port.

Example: Like a mail sorter who only looks at zip codes — fast but doesn't read letter contents.

AWS NLB HAProxy in TCP mode

2 Layer 7 (Application Layer)

- 👉 Routes based on HTTP headers, URL paths, cookies, or content.

Example: Like a smart receptionist who reads your request and directs you to the right department.

AWS ALB Nginx HAProxy in HTTP mode

3 When to Use Layer 4

- 👉 Need maximum speed and lowest latency (no content inspection overhead).

Example: High-throughput TCP services, gaming servers, IoT data streams.

Lower latency higher throughput simpler

4 When to Use Layer 7

- 👉 Need content-based routing, SSL termination, or path-based routing.

Example: Route /api/ to API servers, /static/* to CDN, /admin/* to admin servers.*

More features flexible routing SSL offloading

5 SSL/TLS Handling

- 👉 Layer 4 passes encrypted traffic through (end-to-end encryption); Layer 7 terminates SSL at load balancer, inspects content, re-encrypts to backend.

Example: Layer 4 preserves end-to-end encryption; Layer 7 trades security for content inspection.

Security vs performance trade-off

6 Health Checks

- 👉 Layer 4 uses simple TCP connection test (is port open?); Layer 7 uses HTTP health check (does /health return 200 OK?).

Example: Layer 7 provides smarter health monitoring than Layer 4.

Layer 7 detects application-level failures



Load Balancing Strategies

💡 Problem 4: Health Checks and Failover Strategy

Question: INTERVIEWER: One of your backend servers crashes. How does the load balancer detect this and stop sending traffic to it?

Solution:

Health checks ensure load balancers only route traffic to healthy servers.

1 Active Health Checks

- 👉 Load balancer periodically pings each server to verify it's healthy.

Example: Every 10 seconds, send GET /health → expect 200 OK response.

Proactive monitoring

2 Passive Health Checks

- 👉 Monitor real traffic and mark server unhealthy after consecutive failures.

Example: If 3 consecutive requests to Server B fail → remove it from pool.

React to actual traffic patterns

3 Health Check Intervals

- 👉 Configure how often to check server health (balance between quick detection and server load).

Example: Check every 5 seconds with 2-second timeout, mark unhealthy after 3 failures.

Tuning for responsiveness vs overhead

4 Graceful Degradation

- 👉 Continue serving traffic with remaining healthy servers when some fail.

Example: 10 servers → 2 crash → remaining 8 handle traffic until auto-scaling adds replacements.

5 Circuit Breaker Integration

- 👉 Temporarily stop routing to consistently failing server to let it recover.

Example: Like giving an injured player a break instead of forcing them to play and worsen injury.

Prevents cascading failures

6 Connection Draining

- 👉 Allow in-flight requests to complete before removing server from pool.

Example: Server marked unhealthy → wait 30 seconds for existing requests to finish before fully removing.

Prevents request interruption during deployments



High Availability & Fault Tolerance



Problem 1: Designing a System with 99.99% Uptime

Question: INTERVIEWER: Your e-commerce platform must have 99.99% uptime (only 52 minutes downtime per year). How do you architect for this?

Solution:

High availability requires eliminating single points of failure across all layers.

1 Multi-Region Deployment

👉 Deploy application across multiple AWS/Azure regions for geographic redundancy.

Example: Primary in US-East, failover in EU-West — if entire US region fails, EU takes over.

Active-active or active-passive setup

2 Load Balancer Redundancy

👉 Use multiple load balancers with health checks and DNS failover.

Example: Two load balancers behind Route 53 — if one fails, DNS routes to backup.

Eliminates load balancer as single point of failure

3 Database Replication

👉 Maintain synchronized database replicas with automatic failover.

Example: Primary DB in us-east-1, read replicas in us-west-2 — if primary fails, promote replica.

RDS Multi-AZ

PostgreSQL streaming replication

4 Stateless Application Servers

👉 Design servers with no local state so any instance failure doesn't lose data.

Example: Session data in Redis, not server memory — users stay logged in even if server crashes.

Horizontal scaling easy recovery

5 Health Monitoring & Auto-Recovery

- 👉 Continuously monitor services and automatically restart or replace failed instances.

Example: Kubernetes detects crashed pod → launches replacement in 10 seconds.

Self-healing infrastructure

6 Circuit Breakers

- 👉 Prevent cascading failures by stopping requests to failing dependencies.

Example: Payment service down → circuit opens, show "payment temporarily unavailable" instead of crashing entire checkout.

Resilience patterns

7 Backup & Disaster Recovery

- 👉 Regular automated backups with tested restore procedures.

Example: Daily DB snapshots to S3, tested monthly disaster recovery drills.

RTO and RPO planning

8 Zero-Downtime Deployments

- 👉 Use blue-green or rolling deployments to update without downtime.

Example: Deploy new version to "green" environment, switch traffic when ready, keep "blue" as instant rollback.

Continuous deployment without service interruption

💡 Problem 2: Handling Partial System Failures

Question: INTERVIEWER: Your recommendation service is down, but users should still be able to browse and buy products. How do you handle this?

Solution:

Graceful degradation ensures core functionality works even when some services fail.

1 Service Isolation

- 👉 Each microservice fails independently without bringing down others.

Example: Like a restaurant — if dessert kitchen fails, you can still order main courses.

Microservices architecture bulkheads

2 Fallback Responses

- 👉 Return cached or default data when service is unavailable.

Example: Recommendation service down → show "Popular items" instead of personalized recommendations.

Graceful degradation

3 Timeouts & Retries

- 👉 Set aggressive timeouts to fail fast and retry with exponential backoff.

Example: Wait max 500ms for recommendation service → if timeout, show default recommendations.

Prevent waiting indefinitely for failed services

4 Circuit Breaker Pattern

- 👉 Stop calling failing service temporarily to let it recover.

Example: After 5 consecutive failures, stop calling recommendation service for 60 seconds — prevents overload.

Hystrix Resilience4j

5 Feature Flags

- 👉 Dynamically disable non-critical features during incidents.

Example: Turn off "related products" section if service is struggling — keeps checkout working.

LaunchDarkly feature toggles

6 Async Background Jobs

- 👉 Move non-critical operations to background queues that can be retried later.

Example: Order confirmation email fails → queue it for retry, but complete the purchase.

Kafka SQS delayed job processing

💡 Problem 3: Implementing Bulkhead Pattern

Question: INTERVIEWER: How do you prevent one slow service from exhausting all resources?

Solution:

Bulkhead pattern isolates resources to contain failures.

1 Resource Isolation

- 👉 Separate thread pools/connection pools for different services.

Example: Payment service gets 10 threads, Search gets 20 threads → isolated.

Thread pools resource isolation

2 Preventing Cascading Failures

- 👉 Slow payment service doesn't block search functionality.

Example: Payment threads exhausted → search still works with its own threads.

Failure containment independent pools

3 Connection Pool Bulkheads

- 👉 Separate database connection pools per service/feature.

Example: Analytics queries have separate pool → don't block transactional queries.

Connection pooling query isolation

4 Semaphore Pattern

- 👉 Limit concurrent requests to external service.

Example: Max 5 concurrent calls to third-party API → prevents overwhelming it.

Semaphores concurrency limits

5 Queue-Based Bulkheads

- 👉 Use separate queues for different workloads.

Example: High-priority orders queue separate from bulk reports queue.

Queue separation priority handling

6 Kubernetes Resource Limits

- 👉 Set CPU/memory limits per pod to prevent resource hogging.

Example: Each microservice pod limited to 2 CPU cores, 4GB RAM.

Resource quotas Kubernetes limits

Problem 3: Single Point of Failure Elimination

Question: INTERVIEWER: Your system has a single Redis instance for caching. What happens if it crashes?

Solution:

Identify and eliminate every single point of failure in your architecture.

1 Redis Clustering

- 👉 Distribute cache across multiple Redis nodes with automatic failover.

Example: 3-node Redis cluster — if one node fails, others continue serving requests.

Redis Sentinel Redis Cluster

2 Cache-Aside Pattern with Fallback

- 👉 If cache is down, application falls back to database directly.

Example: Try Redis → if timeout, query database directly → slower but system stays up.

Degraded performance but no complete failure

3 Message Queue Redundancy

- 👉 Use clustered message brokers with replication.

Example: 3 Kafka brokers with replication factor 3 — losing one broker doesn't lose messages.

Kafka clusters RabbitMQ mirrored queues

4 Database High Availability

- 👉 Primary database with automatic failover to replica.

Example: PostgreSQL primary with 2 streaming replicas — if primary fails, promote replica in seconds.

AWS RDS Multi-AZ Patroni for PostgreSQL

5 Network Redundancy

- 👉 Multiple network paths and availability zones.

Example: Deploy across 3 availability zones — if one AZ loses connectivity, traffic routes to others.

Multi-AZ deployment cross-region replication

6 Monitoring & Alerting

- 👉 Detect failures instantly and alert team for manual intervention if needed.

Example: Redis down → PagerDuty alert fires within 30 seconds → team investigates.

Prometheus Datadog CloudWatch



CAP Theorem



Problem 1: Choosing Between Consistency and Availability

Question: INTERVIEWER: You're building a social media 'likes' counter. Should you prioritize consistency or availability?

Solution:

CAP theorem states you can only have 2 of 3: Consistency, Availability, Partition Tolerance.

1 CP Systems (Consistency + Partition Tolerance)

- 👉 Always return correct data, even if it means refusing requests during network partitions.

Example: Banking system — better to show "temporarily unavailable" than show wrong account balance.

MongoDB with majority writes HBase ZooKeeper

2 AP Systems (Availability + Partition Tolerance)

- 👉 Always respond to requests, even if data might be slightly stale or inconsistent.

Example: Social media likes — showing 99 likes instead of 100 for a few seconds is acceptable.

Cassandra DynamoDB with eventual consistency Riak

3 CA Systems (Consistency + Availability)

- 👉 Traditional RDBMS that assume no network partitions (single datacenter).

Example: PostgreSQL in single AZ — consistent and available, but entire DB goes down if network partitions.

PostgreSQL MySQL without replication

4 Eventual Consistency Trade-off

- 👉 Accept temporary inconsistency in exchange for high availability.

Example: Like news spreading through social media — everyone eventually hears it, but not instantly.

DNS CDN caching distributed counters

5 Strong Consistency Cost

- 👉 Requires coordination between nodes, increasing latency.

Example: Like getting everyone in a meeting to agree before announcing decision — slower but accurate.

Distributed transactions Paxos/Raft consensus

6 Real-World Example: Amazon Shopping Cart

- 👉 Availability prioritized — even during network issues, you can add items to cart.

Example: Occasionally you might add same item twice, but that's better than checkout being down.

AP system design

💡 Problem 2: Network Partition Handling

Question: INTERVIEWER: Your distributed database splits into two partitions due to network failure. How do you handle reads and writes?

Solution:

Network partitions are inevitable in distributed systems — must choose how to handle them.

1 Quorum-Based Writes

- 👉 Require majority of nodes to acknowledge write before confirming success.

Example: 5-node cluster → write succeeds only if 3+ nodes confirm (majority quorum).

Cassandra

DynamoDB with quorum consistency

2 Split-Brain Prevention

- 👉 Ensure only one partition can accept writes during network split.

Example: Partition with majority of nodes stays active, minority partition refuses writes.

Avoids conflicting updates

3 Read Repair

- 👉 Detect and fix inconsistencies when reading data from multiple replicas.

Example: Query 3 nodes → 2 return version A, 1 returns old version B → update B to A.

Cassandra read repair

anti-entropy

4 Version Vectors / Vector Clocks

- 👉 Track causality of updates to resolve conflicts during partition recovery.

Example: Like Git merge — system knows which updates happened before/after others.

DynamoDB

Riak

5 Last-Write-Wins (LWW)

- 👉 Simplest conflict resolution — newest timestamp wins.

Example: Two users edit same document during partition → most recent edit overwrites older one.

Risk: may lose legitimate updates

6 Application-Level Conflict Resolution

- 👉 Let application decide how to merge conflicting updates.

Example: Shopping cart — if partition causes duplicate items, merge by adding quantities.

Custom business logic for conflicts

💡 Problem 3: HTTP/2 vs HTTP/3 Protocol Differences

Question: INTERVIEWER: Your app uses HTTP/1.1 — what do you gain by upgrading to HTTP/2 or HTTP/3?

Solution:

HTTP/2 and HTTP/3 solve head-of-line blocking and improve performance over slow networks.

1 HTTP/1.1 Limitations

- 👉 One request per TCP connection, or head-of-line blocking in pipelining.

Example: 6 parallel connections limit → page with 100 assets loads slowly.

HTTP/1.1 connection limits

2 HTTP/2 Multiplexing

- 👉 Multiple requests/responses over single TCP connection.

Example: 100 assets → all loaded over 1 connection → no connection limit bottleneck.

Multiplexing single connection

3 HTTP/2 Server Push

- 👉 Server proactively sends resources before client requests.

Example: Server pushes CSS, JS immediately after HTML → faster page load.

Server push preload

4 HTTP/2 Header Compression (HPACK)

- 👉 Compress repeated headers to reduce overhead.

Example: Cookie header repeated in 100 requests → HPACK compresses → saves bandwidth.

HPACK compression

5 HTTP/3 QUIC Protocol

- 👉 Built on UDP instead of TCP → faster connection establishment.

Example: TCP requires 3-way handshake (1-2 RTTs); QUIC: 0-1 RTT → faster.

QUIC UDP 0-RTT

6 HTTP/3 No Head-of-Line Blocking

- 👉 TCP packet loss blocks all streams; QUIC isolates streams.

Example: 1 lost packet in HTTP/2 blocks all requests; HTTP/3 only blocks affected stream.

No HOL blocking independent streams

7 When to Upgrade

- 👉 HTTP/2 for multiplexing; HTTP/3 for mobile/lossy networks.

Example: E-commerce site with many assets → HTTP/2; Global app → HTTP/3 for better mobile performance.

Migration use cases

Problem 3: HTTP/2 vs HTTP/3 Protocol Differences

Question: INTERVIEWER: Your app uses HTTP/1.1 — what do you gain by upgrading to HTTP/2 or HTTP/3?

Solution:

HTTP/2 and HTTP/3 solve head-of-line blocking and improve performance over slow networks.

1 HTTP/1.1 Limitations

- 👉 One request per TCP connection, or head-of-line blocking in pipelining.

Example: 6 parallel connections limit → page with 100 assets loads slowly.

HTTP/1.1 connection limits

2 HTTP/2 Multiplexing

- 👉 Multiple requests/responses over single TCP connection.

Example: 100 assets → all loaded over 1 connection → no connection limit bottleneck.

Multiplexing single connection

3 HTTP/2 Server Push

- 👉 Server proactively sends resources before client requests.

Example: Server pushes CSS, JS immediately after HTML → faster page load.

Server push preload

4 HTTP/2 Header Compression (HPACK)

- 👉 Compress repeated headers to reduce overhead.

Example: Cookie header repeated in 100 requests → HPACK compresses → saves bandwidth.

HPACK compression

5 HTTP/3 QUIC Protocol

- 👉 Built on UDP instead of TCP → faster connection establishment.

Example: TCP requires 3-way handshake (1-2 RTTs); QUIC: 0-1 RTT → faster.

QUIC UDP 0-RTT

6 HTTP/3 No Head-of-Line Blocking

- 👉 TCP packet loss blocks all streams; QUIC isolates streams.

Example: 1 lost packet in HTTP/2 blocks all requests; HTTP/3 only blocks affected stream.

No HOL blocking independent streams

7 When to Upgrade

- 👉 HTTP/2 for multiplexing; HTTP/3 for mobile/lossy networks.

Example: E-commerce site with many assets → HTTP/2; Global app → HTTP/3 for better mobile performance.

Migration use cases



Idempotency



Problem 1: Preventing Duplicate Payment Processing

Question: INTERVIEWER: User double-clicks the 'Pay Now' button. How do you prevent charging them twice?

Solution:

Idempotency ensures repeating the same operation produces the same result without side effects.

1 Idempotency Key

👉 Assign a unique ID to each request so backend recognizes duplicates.

Example: User clicks "Pay Now" twice → both requests use same payment_id → processed only once.

Stripe idempotency keys

UUID per transaction

2 Database Unique Constraint

👉 Database rejects duplicate entries based on unique key.

Example: orders table has UNIQUE constraint on order_id → second insert with same order_id fails gracefully.

PostgreSQL UNIQUE

INSERT ... ON CONFLICT

3 Distributed Lock

👉 Acquire lock before processing to prevent concurrent duplicates.

Example: Redis SET NX with order_id as key → only first request gets lock and processes payment.

Redis SETNX

DynamoDB conditional writes

4 Check-Before-Process

- 👉 Query if request was already processed before executing.

Example: Before processing payment, check if transaction_id already exists in database.

Simple but race conditions possible without locks

5 Frontend Debouncing

- 👉 Disable button immediately after first click to prevent duplicate submissions.

Example: Like elevator button — pressing multiple times doesn't call multiple elevators.

Client-side prevention but not sufficient alone

6 API Response Caching

- 👉 Cache API responses by idempotency key for safe retries.

Example: First request processed → cache response for 24 hours → retries return cached response.

Stripe caches idempotent responses for 24 hours

💡 Problem 2: Idempotent REST API Design

Question: INTERVIEWER: Which HTTP methods should be idempotent, and how do you implement idempotency in POST requests?

Solution:

Idempotency is critical for reliable API design, especially for retries and network failures.

1 GET Requests (Idempotent by Design)

- 👉 Multiple GET requests return same data without modifying state.

Example: Calling GET /users/123 ten times returns same user data every time.

Read operations are naturally idempotent

2 PUT Requests (Idempotent by Design)

- 👉 Updating a resource with same data multiple times produces same result.

Example: `PUT /users/123 {"name": "John"}` — calling twice still results in name = "John".

Full resource replacement

3 DELETE Requests (Idempotent)

- 👉 Deleting same resource multiple times has same effect (resource is gone).

Example: `DELETE /users/123` — first call deletes user, subsequent calls return 404 but system state unchanged.

Deleting already-deleted resource is safe

4 POST Requests (NOT Idempotent by Default)

- 👉 POST creates new resources, so repeating creates duplicates.

Example: `POST /orders` creates new order each time — duplicates are problem.

Need explicit idempotency handling

5 Idempotency-Key Header

- 👉 Client sends unique idempotency key with POST requests.

Example: `POST /orders` with `Idempotency-Key: abc123` → server stores key and result → retries return cached result.

Stripe-style idempotency

6 UUID-Based Resource Creation

- 👉 Client generates UUID and uses PUT instead of POST.

Example: `PUT /orders/550e8400-e29b-41d4-a716-446655440000` → creates or updates order with that ID.

Client-generated IDs

7 Response Status Codes

- 👉 Return appropriate status codes for idempotent operations.

Example: First POST returns 201 Created, duplicate returns 200 OK with original resource.

Clear communication of idempotency

💡 Problem 3: Message Queue Idempotency

Question: INTERVIEWER: A message in your Kafka queue gets processed twice due to retry. How do you handle this?

Solution:

Message queue idempotency prevents duplicate processing of messages.

1 Message Deduplication

- 👉 Store message IDs in database or cache to detect duplicates.

Example: Like a bouncer with a guest list — checks if person already entered before letting them in.

Redis set with message IDs database unique constraint

2 At-Least-Once with Idempotent Processing

- 👉 Accept that messages may be delivered multiple times, design processing to be idempotent.

Example: Email notification job → check if email already sent before sending again.

Kafka at-least-once delivery with idempotent consumers

3 Exactly-Once Semantics

- 👉 Use messaging system's built-in exactly-once delivery guarantees.

Example: Kafka transactions ensure message is processed exactly once across consumer group.

Kafka exactly-once semantics SQS FIFO with deduplication

4 Database Transaction with Message Ack

- 👉 Process message and acknowledge it within same database transaction.

Example: Debit account AND mark message as processed in single transaction → atomic operation.

Transactional outbox pattern

5 Versioning / Timestamps

- 👉 Include version or timestamp in messages to detect stale duplicates.

Example: Two messages to update user email → process only the one with newer timestamp.

Conflict resolution

6 Bloom Filters for Fast Duplicate Detection

- 👉 Use probabilistic data structure to quickly check if message was seen before.

Example: Like a quick memory check — "have I seen this message ID before?"

Memory-efficient deduplication for high-throughput systems



Circuit Breaker Pattern

💡 Problem 1: Preventing Cascading Failures

Question: INTERVIEWER: Your payment service is down and causing timeouts across your entire application. How do you prevent this from bringing down everything?

Solution:

Circuit breakers protect your system from cascading failures when dependencies fail.

1 Circuit Breaker States

- 👉 Circuit breakers have three states: Closed (normal), Open (failing), and Half-Open (testing recovery).

Example: Like an electrical circuit breaker that trips when overloaded, then tests before fully reconnecting.

Hystrix Resilience4j custom implementation

2 Closed State (Normal Operation)

- 👉 All requests pass through normally while monitoring for failures.

Example: Payment service working fine → all requests go through → error rate monitored.

Success threshold: < 50% error rate

3 Open State (Fail Fast)

- 👉 Immediately reject requests without calling the failing service.

Example: Payment service down → circuit opens → return cached response or error immediately.

Prevents wasting time on doomed requests

4 Half-Open State (Testing Recovery)

- 👉 Allow limited requests through to test if service has recovered.

Example: After 30 seconds, try 5 requests → if successful, close circuit; if failed, reopen.

Gradual recovery testing

5 Fallback Mechanisms

- 👉 Provide alternative responses when circuit is open.

Example: Payment service down → show "Pay with alternative method" or use cached data.

Graceful degradation

6 Failure Thresholds

- 👉 Define error rate percentage or count that triggers circuit to open.

Example: Open circuit after 50% error rate over 10 requests, or 5 consecutive failures.

Prevents false positives from single failures

7 Timeout Configuration

- 👉 Set aggressive timeouts to detect failures quickly.

Example: 3-second timeout → if service doesn't respond, count as failure.

Fast failure detection

8 Monitoring & Alerts

- 👉 Track circuit state changes and alert when circuits open.

Example: Circuit opened → alert team → investigate root cause before customers complain.

Prometheus metrics Datadog monitoring

💡 Problem 2: Circuit Breaker for External API Calls

Question: INTERVIEWER: Your app calls a third-party weather API that's flaky. How do you ensure it doesn't slow down your entire application?

Solution:

Circuit breakers isolate external dependencies to prevent them from impacting your system.

1 Per-Dependency Circuit Breakers

- 👉 Use separate circuit breakers for each external dependency.

Example: Weather API circuit independent of Maps API circuit → one fails, others keep working.

Isolated failure domains

2 Request Timeout

- 👉 Set short timeouts for external API calls to fail fast.

Example: Weather API timeout = 2 seconds → don't wait 30 seconds for slow response.

Prevent indefinite blocking

3 Retry with Exponential Backoff

- 👉 Retry failed requests with increasing delays between attempts.

Example: Retry after 1s, then 2s, then 4s → gives failing service time to recover.

Avoid overwhelming failing service

4 Cached Fallback Data

- 👉 Return stale but cached data when external API is down.

Example: Weather API down → show yesterday's weather data instead of error.

Better user experience than complete failure

5 Bulkhead Pattern

- 👉 Isolate thread pools for different external dependencies.

Example: 10 threads for weather API, 10 for maps API → one exhausted pool doesn't affect other.

Resource isolation

6 Health Check Endpoint

- 👉 Periodically check external API health to preemptively open circuit.

Example: Poll /health every 30s → if unhealthy, open circuit before user requests fail.

Proactive monitoring



Rate Limiting & Throttling

💡 Problem 1: Protecting API from Abuse

Question: INTERVIEWER: Your public API is getting hammered by bots making 10,000 requests per second. How do you protect it?

Solution:

Rate limiting prevents abuse and ensures fair usage of your API resources.

1 Fixed Window Rate Limiting

- 👉 Count requests in fixed time windows (e.g., per minute).

Example: Allow 100 requests per minute → reset counter at start of each minute.

Simple but burst traffic issues at window boundaries

2 Sliding Window Rate Limiting

- 👉 Use rolling time window for smoother rate limiting.

Example: Count requests in last 60 seconds (sliding window) → prevents burst at window boundaries.

More accurate prevents window boundary exploitation

3 Token Bucket Algorithm

- 👉 Tokens refill at steady rate, requests consume tokens.

Example: Bucket holds 100 tokens, refills 10/second → allows bursts up to 100, maintains average rate.

Allows bursts while maintaining average rate

4 Leaky Bucket Algorithm

- 👉 Process requests at fixed rate, queue excess requests.

Example: Like a funnel → water (requests) pours in fast, but drains out at steady rate.

Smooths traffic spikes strict rate enforcement

5 Rate Limit Tiers

- 👉 Different rate limits for different user tiers.

Example: Free tier: 100 req/hour, Pro: 1000 req/hour, Enterprise: unlimited.

Monetization fair resource allocation

6 Distributed Rate Limiting with Redis

- 👉 Use centralized Redis to track rate limits across multiple servers.

Example: All API servers check same Redis counter → consistent rate limiting.

INCR command with EXPIRE synchronized across cluster

7 Rate Limit Headers

- 👉 Return rate limit info in response headers.

Example: X-RateLimit-Remaining: 45, X-RateLimit-Reset: 1609459200.

Client knows when to retry

8 IP-Based vs User-Based Limiting

- 👉 Rate limit by IP address for anonymous users, by user ID for authenticated.

Example: Anonymous users: 10 req/min per IP, Logged in: 100 req/min per user.

Flexible limiting strategies

💡 Problem 2: Handling Rate Limit Violations

Question: INTERVIEWER: A user exceeds their rate limit. What should your API return and how should it behave?

Solution:

Proper rate limit handling improves user experience and protects your API.

1 HTTP 429 Status Code

👉 Return 429 Too Many Requests when rate limit exceeded.

Example: User exceeds limit → return 429 with clear error message.

Standard HTTP status for rate limiting

2 Retry-After Header

👉 Tell client when they can retry.

Example: Retry-After: 60 → client should wait 60 seconds before retrying.

Helps clients implement proper backoff

3 Error Response Body

👉 Include helpful error message explaining rate limit.

Example: {"error": "Rate limit exceeded. Try again in 30 seconds."}

Actionable error information

4 Soft Throttling vs Hard Blocking

👉 Soft throttle by slowing down, hard block by rejecting.

Example: Soft: add 100ms delay per request over limit; Hard: reject immediately.

Tiered service quality

5 Burst Allowance

- 👉 Allow short bursts above sustained rate.

Example: Average 10 req/sec allowed, but burst of 50 requests tolerated briefly.

Token bucket allows this flexibility

6 Warning Thresholds

- 👉 Warn users before they hit rate limit.

Example: At 80% of limit, return X-RateLimit-Warning header.

Proactive communication

💡 Problem 3: Advanced Rate Limiting Algorithms

Question: INTERVIEWER: Your API has different rate limits for free vs paid users — how do you implement that?

Solution:

Advanced rate limiting balances fairness, resource protection, and business requirements.

1 Token Bucket Algorithm

- 👉 Bucket fills with tokens at fixed rate, each request consumes token.

Example: 100 tokens/minute bucket → burst of 100 requests allowed, then 100/min sustained.

Token bucket burst handling

2 Leaky Bucket Algorithm

- 👉 Requests queued, processed at fixed rate.

Example: Queue holds 50 requests → process 10/sec → smooths traffic spikes.

Leaky bucket traffic smoothing

3 Fixed Window Counter

- 👉 Count requests per fixed time window.

Example: 1000 requests per hour → resets at top of hour → potential burst at boundary.

Fixed window simple counter

4 Sliding Window Log

- 👉 Track timestamp of each request, count in sliding window.

Example: Track last 100 request times → allow request if count in last minute < 100.

Sliding window accurate limiting

5 User-Tier Based Limits

- 👉 Different limits per user tier.

Example: Free: 100/hr, Pro: 1K/hr, Enterprise: 10K/hr → stored in Redis.

Tiered limits business logic

6 Distributed Rate Limiting

- 👉 Coordinate limits across multiple API servers.

Example: Redis tracks global count → all API servers check/increment → consistent limits.

Distributed Redis coordination

7 Rate Limit Headers

- 👉 Inform clients of their limit status.

Example: X-RateLimit-Limit: 1000, X-RateLimit-Remaining: 847 → client can self-throttle.

HTTP headers client awareness

Problem 3: Advanced Rate Limiting Algorithms

Question: INTERVIEWER: Your API has different rate limits for free vs paid users — how do you implement that?

Solution:

Advanced rate limiting balances fairness, resource protection, and business requirements.

1 Token Bucket Algorithm

- 👉 Bucket fills with tokens at fixed rate, each request consumes token.

Example: 100 tokens/minute bucket → burst of 100 requests allowed, then 100/min sustained.

Token bucket burst handling

2 Leaky Bucket Algorithm

- 👉 Requests queued, processed at fixed rate.

Example: Queue holds 50 requests → process 10/sec → smooths traffic spikes.

Leaky bucket traffic smoothing

3 Fixed Window Counter

- 👉 Count requests per fixed time window.

Example: 1000 requests per hour → resets at top of hour → potential burst at boundary.

Fixed window simple counter

4 Sliding Window Log

- 👉 Track timestamp of each request, count in sliding window.

Example: Track last 100 request times → allow request if count in last minute < 100.

Sliding window accurate limiting

5 User-Tier Based Limits

- 👉 Different limits per user tier.

Example: Free: 100/hr, Pro: 1K/hr, Enterprise: 10K/hr → stored in Redis.

Tiered limits business logic

6 Distributed Rate Limiting

- 👉 Coordinate limits across multiple API servers.

Example: Redis tracks global count → all API servers check/increment → consistent limits.

Distributed Redis coordination

7 Rate Limit Headers

- 👉 Inform clients of their limit status.

Example: X-RateLimit-Limit: 1000, X-RateLimit-Remaining: 847 → client can self-throttle.

HTTP headers client awareness

💡 Problem 3: DDoS Protection Strategy

Question: INTERVIEWER: Your API is under DDoS attack with millions of requests from different IPs. How do you defend it?

Solution:

DDoS protection requires multiple layers of defense beyond simple rate limiting.

1 CDN & Edge Protection

- 👉 Use CDN to absorb and filter DDoS traffic at edge.

Example: Cloudflare blocks malicious traffic before it reaches your servers.

Cloudflare AWS Shield Akamai

2 IP Reputation & Blocking

- 👉 Block known malicious IPs automatically.

Example: IP from known botnet → blocked at firewall level.

[WAF rules](#) [IP blocklists](#)

3 CAPTCHA for Suspicious Traffic

- 👉 Challenge suspicious requests with CAPTCHA.

Example: 100 requests in 10 seconds from one IP → require CAPTCHA verification.

[reCAPTCHA](#) [hCaptcha](#)

4 Connection Limiting

- 👉 Limit concurrent connections per IP.

Example: Max 10 concurrent connections per IP → prevents connection exhaustion.

[nginx limit_conn](#) [AWS WAF](#)

5 Request Size Limits

- 👉 Reject oversized requests to prevent payload attacks.

Example: Reject requests larger than 10MB.

[Prevent payload-based attacks](#)

6 Geo-Blocking

- 👉 Block traffic from regions you don't serve.

Example: US-only service → block traffic from other countries during attack.

[CloudFront geo-restrictions](#)

7 Auto-Scaling with Limits

- 👉 Scale automatically but set max limits to control costs.

Example: Auto-scale up to 100 servers max → prevents unlimited cloud bills during DDoS.

Budget protection during attacks

8 API Key Requirements

👉 Require API keys for all requests to enable accountability.

Example: Anonymous requests → reject; API key required → can track and block abusers.

Accountability and blocking granularity

💡 Problem 4: Designing Secure VPC Architecture

Question: INTERVIEWER: How do you design a VPC that's secure, scalable, and allows for disaster recovery?

Solution:

VPC design establishes network boundaries, security zones, and connectivity patterns.

1 Public vs Private Subnets

👉 Public subnets for internet-facing resources, private for internal.

Example: ALB in public subnet → app servers in private subnet → no direct internet access.

Subnet separation security zones

2 Multi-AZ Deployment

👉 Spread resources across availability zones.

Example: 3 AZs → app servers in each → survive AZ failure.

Multi-AZ high availability

3 NAT Gateway for Outbound Traffic

- 👉 Private subnets access internet via NAT gateway.

Example: App server downloads packages → routed through NAT gateway in public subnet.

NAT Gateway outbound access

4 Security Groups (Stateful Firewall)

- 👉 Control inbound/outbound traffic at instance level.

Example: App server SG allows port 80/443 from ALB SG only → no direct internet access.

Security Groups firewall rules

5 Network ACLs (Stateless Firewall)

- 👉 Additional subnet-level firewall rules.

Example: Block known malicious IP ranges at NACL level → never reach instances.

NACLs subnet firewall

6 VPC Peering for Multi-Account

- 👉 Connect VPCs across accounts or regions.

Example: Production VPC peers with monitoring VPC → centralized logging.

VPC Peering multi-account

7 VPN/Direct Connect for Hybrid

- 👉 Secure connection between on-premise and cloud.

Example: Corporate network → VPN to AWS VPC → access internal resources securely.

VPN Direct Connect hybrid cloud

💡 Problem 4: Designing Secure VPC Architecture

Question: INTERVIEWER: How do you design a VPC that's secure, scalable, and allows for disaster recovery?

Solution:

VPC design establishes network boundaries, security zones, and connectivity patterns.

1 Public vs Private Subnets

- 👉 Public subnets for internet-facing resources, private for internal.

Example: ALB in public subnet → app servers in private subnet → no direct internet access.

Subnet separation security zones

2 Multi-AZ Deployment

- 👉 Spread resources across availability zones.

Example: 3 AZs → app servers in each → survive AZ failure.

Multi-AZ high availability

3 NAT Gateway for Outbound Traffic

- 👉 Private subnets access internet via NAT gateway.

Example: App server downloads packages → routed through NAT gateway in public subnet.

NAT Gateway outbound access

4 Security Groups (Stateful Firewall)

- 👉 Control inbound/outbound traffic at instance level.

Example: App server SG allows port 80/443 from ALB SG only → no direct internet access.

Security Groups firewall rules

5 Network ACLs (Stateless Firewall)

- 👉 Additional subnet-level firewall rules.

Example: Block known malicious IP ranges at NACL level → never reach instances.

NACLs subnet firewall

6 VPC Peering for Multi-Account

- 👉 Connect VPCs across accounts or regions.

Example: Production VPC peers with monitoring VPC → centralized logging.

VPC Peering multi-account

7 VPN/Direct Connect for Hybrid

- 👉 Secure connection between on-premise and cloud.

Example: Corporate network → VPN to AWS VPC → access internal resources securely.

VPN Direct Connect hybrid cloud

💡 Problem 4: Gateway vs Application-Level Rate Limiting

Question: INTERVIEWER: Your API is being hammered by bots — rate-limit at gateway or app level?

Solution:

Multi-layer rate limiting provides defense in depth against abuse.

1 Gateway-Level Rate Limiting (First Line of Defense)

- 👉 Implement rate limiting at API Gateway before requests hit application.

Example: AWS API Gateway blocks 10K req/sec from single IP → app servers never see the traffic.

API Gateway throttling WAF rate limits edge protection

2 Application-Level Rate Limiting (Business Logic)

- 👉 Enforce per-user or per-feature rate limits based on business rules.

Example: Free tier: 100 API calls/day, Premium: 10K calls/day → enforced at app layer.

Custom rate limiting tiered limits business rules

3 Distributed Rate Limiting

- 👉 Share rate limit state across application instances via Redis.

Example: 3 app servers share Redis counter → consistent rate limit across cluster.

Redis rate limiting centralized counters

4 IP-Based vs Token-Based Limiting

- 👉 Use IP limiting at gateway, token/user limiting at application.

Example: Gateway: max 1K req/sec per IP; App: max 100 req/min per API key.

Layered defense different granularity

5 Adaptive Rate Limiting

- 👉 Dynamically adjust limits based on system load.

Example: CPU >80% → reduce rate limits by 50% → protect backend from overload.

Dynamic limits load-based throttling

6 CAPTCHA Integration

- 👉 Challenge suspicious traffic patterns with CAPTCHA.

Example: Same IP makes 100 requests in 10 seconds → require CAPTCHA verification.

reCAPTCHA bot detection

💡 Problem 4: Gateway vs Application-Level Rate Limiting

Question: INTERVIEWER: Your API is being hammered by bots — rate-limit at gateway or app level?

Solution:

Multi-layer rate limiting provides defense in depth against abuse.

1 Gateway-Level Rate Limiting (First Line of Defense)

👉 Implement rate limiting at API Gateway before requests hit application.

Example: AWS API Gateway blocks 10K req/sec from single IP → app servers never see the traffic.

API Gateway throttling WAF rate limits edge protection

2 Application-Level Rate Limiting (Business Logic)

👉 Enforce per-user or per-feature rate limits based on business rules.

Example: Free tier: 100 API calls/day, Premium: 10K calls/day → enforced at app layer.

Custom rate limiting tiered limits business rules

3 Distributed Rate Limiting

👉 Share rate limit state across application instances via Redis.

Example: 3 app servers share Redis counter → consistent rate limit across cluster.

Redis rate limiting centralized counters

4 IP-Based vs Token-Based Limiting

👉 Use IP limiting at gateway, token/user limiting at application.

Example: Gateway: max 1K req/sec per IP; App: max 100 req/min per API key.

Layered defense different granularity

5 Adaptive Rate Limiting

- 👉 Dynamically adjust limits based on system load.

Example: CPU >80% → reduce rate limits by 50% → protect backend from overload.

Dynamic limits

load-based throttling

6 CAPTCHA Integration

- 👉 Challenge suspicious traffic patterns with CAPTCHA.

Example: Same IP makes 100 requests in 10 seconds → require CAPTCHA verification.

reCAPTCHA

bot detection



Graceful Degradation

💡 Problem 1: Serving Core Features During Outages

Question: INTERVIEWER: Your recommendation engine is down. How do you ensure users can still shop and checkout?

Solution:

Graceful degradation prioritizes core features when non-critical services fail.

1 Feature Prioritization

- 👉 Identify core vs nice-to-have features.

Example: Checkout is core, recommendations are nice-to-have → disable recommendations first.

Business impact analysis

2 Static Fallback Content

- 👉 Show static content when dynamic service fails.

Example: Recommendation engine down → show static "Trending items" list.

Simple reliable fallback

3 Cached Fallback

- 👉 Serve stale cached data when live service unavailable.

Example: Product reviews service down → show cached reviews from 1 hour ago.

Better user experience than errors

4 Feature Toggles

- 👉 Remotely disable features without deploying code.

Example: Service struggling → flip feature flag to disable recommendations → instant relief.

LaunchDarkly Unleash custom flags

5 Reduced Functionality Mode

- 👉 Offer limited functionality instead of complete failure.

Example: Search down → enable browse by category only.

Partial service better than none

6 User Communication

- 👉 Display banners explaining reduced functionality.

Example: "We're experiencing high traffic. Some features may be slower than usual."

Transparency builds trust

💡 Problem 2: Database Read Replica Failure Handling

Question: INTERVIEWER: Your read replicas are lagging 5 minutes behind primary. How do you handle user queries?

Solution:

Handle replication lag gracefully to balance consistency and availability.

1 Read-After-Write Consistency

- 👉 Route read requests to primary after write to ensure consistency.

Example: User updates profile → next read goes to primary, not lagging replica.

Session-based routing to primary

2 Staleness Tolerance Detection

- 👉 Determine which queries can tolerate stale data.

Example: Product catalog can be 5 minutes stale, but order status needs real-time data.

Business logic determines acceptable staleness

3 Fallback to Primary

- 👉 Query primary database when replicas are too stale.

Example: Replica lag > 10 seconds → route critical reads to primary.

Consistency over performance

4 Display Staleness Warnings

- 👉 Show users when data might be outdated.

Example: "This information is up to 2 minutes old" banner.

Transparency about eventual consistency

5 Session Affinity

- 👉 Route user's reads to same replica for consistent view.

Example: User always reads from replica-2 during their session → sees their own writes.

Consistent read view within session

6 Monitoring & Alerting

- 👉 Detect failures instantly and alert team for manual intervention if needed.

Example: Redis down → PagerDuty alert fires within 30 seconds → team investigates.

Proactive issue detection



Distributed Tracing

💡 Problem 1: Debugging Slow Requests Across Microservices

Question: INTERVIEWER: A user reports checkout takes 8 seconds, but you have 12 microservices involved. How do you find the bottleneck?

Solution:

Distributed tracing tracks requests across multiple services to identify performance issues.

1 Trace ID Propagation

- 👉 Generate unique trace ID for each request and propagate it across all services.

Example: Request enters API gateway → gets trace_id=abc123 → all 12 microservices log with same trace_id.

OpenTelemetry Jaeger Zipkin

2 Span Creation

- 👉 Each service creates a span representing its portion of work.

Example: Checkout service span: 2.5s, Payment service span: 3.2s, Inventory span: 0.8s.

[Detailed timing breakdown](#)

3 Parent-Child Span Relationships

- 👉 Link spans to show service call hierarchy.

Example: API span is parent → Checkout span is child → Payment span is grandchild.

[Call graph visualization](#)

4 Span Attributes & Tags

- 👉 Add metadata to spans for debugging context.

Example: user_id, product_id, http_status_code, error_message attached to spans.

[Rich context for debugging](#)

5 Distributed Context Propagation

- 👉 Pass trace context through HTTP headers, message queues, and RPC calls.

Example: traceparent header: 00-abc123-def456-01 propagated across services.

[W3C Trace Context standard](#)

6 Sampling Strategy

- 👉 Trace subset of requests to balance performance and visibility.

Example: Trace 10% of normal requests, 100% of errors → reduces storage cost.

[Cost vs visibility trade-off](#)

7 Waterfall Visualization

- 👉 Display timeline showing which services took how long.

Example: Jaeger UI shows 8-second checkout: 2s gateway, 3s payment, 2s inventory, 1s notification.

[Jaeger UI](#) [Zipkin UI](#)

8 Error Tracking

- 👉 Automatically capture and link errors to traces.

Example: Payment failure → trace shows exact line of code and all related spans.

[Full context for error investigation](#)

💡 Problem 2: Tracing Asynchronous Operations

Question: INTERVIEWER: An order confirmation email is delayed by 2 hours. How do you trace async operations through your message queue?

Solution:

Tracing async operations requires propagating context through message queues.

1 Trace Context in Message Headers

- 👉 Include trace_id and span_id in message queue headers.

Example: Kafka message includes {trace_id: abc123, parent_span_id: def456} in headers.

[Maintains trace continuity](#)

2 Producer Span

- 👉 Create span when publishing message to queue.

Example: Order service publishes to Kafka → creates producer span with timestamp.

[Track message production time](#)

3 Consumer Span

- 👉 Create child span when consuming message.

Example: Email service consumes message 2 hours later → span shows 2-hour queue delay.

End-to-end visibility

4 Queue Latency Tracking

- 👉 Measure time message spent in queue.

Example: Message published at 10:00 AM, consumed at 12:00 PM → 2-hour queue latency visible in trace.

Identify queue bottlenecks

5 Dead Letter Queue Tracing

- 👉 Maintain trace context when messages move to DLQ.

Example: Failed message → moved to DLQ → trace shows original request + all retry attempts.

Trace failure scenarios

6 Fan-Out Pattern Tracing

- 👉 Track one message spawning multiple child operations.

Example: Order created → triggers 5 parallel tasks (email, inventory, shipping, analytics, notification) → all linked to same trace.

Parallel operation visibility

Problem 3: Instrumenting a Slow Application

Question: INTERVIEWER: Your app is slow — but you don't know why. What do you instrument first?

Solution:

Systematic instrumentation reveals performance bottlenecks quickly.

1 Application Performance Monitoring (APM)

- 👉 Install APM agent to get automatic transaction tracing.

Example: New Relic/Datadog agent shows 80% of time spent in database queries.

New Relic Datadog Dynatrace

2 Database Query Logging

- 👉 Log slow queries and analyze patterns.

Example: Enable slow_query_log → discover 10 queries taking >1 second each.

PostgreSQL slow query log MySQL slow log

3 Distributed Tracing

- 👉 Trace requests across microservices.

Example: Jaeger shows checkout takes 8s: 5s waiting for payment service, 2s in database.

Jaeger Zipkin OpenTelemetry

4 Metrics Collection

- 👉 Track key metrics: response time, throughput, error rate.

Example: P95 latency jumped from 200ms to 3s starting at 2 PM.

Prometheus CloudWatch Grafana

5 Logging with Correlation IDs

- 👉 Add request IDs to all logs for request flow tracking.

Example: Follow request_id through 5 microservices to find where it slows down.

Structured logging correlation IDs

6 Real User Monitoring (RUM)

- 👉 Measure actual user experience in browser.

Example: RUM shows 90% of users experience 5s load time (server logs say 500ms).

Google Analytics SpeedCurve Browser timing API

7 Profiling

- 👉 Profile CPU and memory usage to find hotspots.

Example: Python profiler shows 60% CPU time in image resizing function.

Python cProfile Java Flight Recorder perf

💡 Problem 3: Instrumenting a Slow Application

Question: INTERVIEWER: Your app is slow — but you don't know why. What do you instrument first?

Solution:

Systematic instrumentation reveals performance bottlenecks quickly.

1 Application Performance Monitoring (APM)

- 👉 Install APM agent to get automatic transaction tracing.

Example: New Relic/Datadog agent shows 80% of time spent in database queries.

New Relic Datadog Dynatrace

2 Database Query Logging

- 👉 Log slow queries and analyze patterns.

Example: Enable `slow_query_log` → discover 10 queries taking >1 second each.

PostgreSQL slow query log

MySQL slow log

3 Distributed Tracing

- 👉 Trace requests across microservices.

Example: Jaeger shows checkout takes 8s: 5s waiting for payment service, 2s in database.

Jaeger

Zipkin

OpenTelemetry

4 Metrics Collection

- 👉 Track key metrics: response time, throughput, error rate.

Example: P95 latency jumped from 200ms to 3s starting at 2 PM.

Prometheus

CloudWatch

Grafana

5 Logging with Correlation IDs

- 👉 Add request IDs to all logs for request flow tracking.

Example: Follow `request_id` through 5 microservices to find where it slows down.

Structured logging

correlation IDs

6 Real User Monitoring (RUM)

- 👉 Measure actual user experience in browser.

Example: RUM shows 90% of users experience 5s load time (server logs say 500ms).

Google Analytics

SpeedCurve

Browser timing API

7 Profiling

- 👉 Profile CPU and memory usage to find hotspots.

Example: Python profiler shows 60% CPU time in image resizing function.

Python cProfile

Java Flight Recorder

perf



Retry Logic & Exponential Backoff

💡 Problem 1: Handling Transient Failures

Question: INTERVIEWER: Your API calls to payment gateway occasionally fail due to network blips. How do you handle retries?

Solution:

Smart retry logic with exponential backoff recovers from transient failures without overwhelming services.

1 Exponential Backoff

- 👉 Increase delay between retries exponentially.

Example: Retry after 1s, then 2s, then 4s, then 8s → gives service time to recover.

Prevents overwhelming recovering service

2 Jitter / Randomization

- 👉 Add random delay to prevent thundering herd.

Example: Instead of all clients retrying at exactly 1s, retry between 0.8s-1.2s.

Prevents synchronized retry storms

3 Maximum Retry Attempts

- 👉 Limit retries to prevent infinite loops.

Example: Max 3 retries → after that, fail permanently and alert.

Fail fast after reasonable attempts

4 Idempotency Requirement

- 👉 Ensure retries don't cause duplicate side effects.

Example: Payment processed with idempotency key → retries don't create duplicate charges.

Prevent duplicate side effects

5 Retry on Specific Errors Only

- 👉 Only retry transient errors, not permanent failures.

Example: Retry 503 Service Unavailable, 429 Rate Limit; Don't retry 400 Bad Request, 401 Unauthorized.

Smart error classification

6 Circuit Breaker Integration

- 👉 Stop retrying when circuit opens.

Example: After 5 failures, circuit opens → skip retry logic, fail fast for 30 seconds.

Prevent wasted retry attempts

7 Retry Budget

- 👉 Limit percentage of requests that can be retries.

Example: Allow max 10% retry traffic → prevents retry storms overwhelming backend.

Protect backend from retry storms

8 Timeout Configuration

- 👉 Set shorter timeouts for retry attempts.

Example: First attempt: 5s timeout; Retries: 2s timeout → faster failure detection.

Balance responsiveness and retry attempts

💡 Problem 2: Database Connection Retry Strategy

Question: INTERVIEWER: Your database restarts and all connections drop. How do you handle reconnection in your application?

Solution:

Database connection retry requires careful handling to avoid overwhelming recovering database.

1 Connection Pool Health Checks

- 👉 Validate connections before using them.

Example: Before executing query, test connection with `SELECT 1` → if fails, get new connection.

Prevent using dead connections

2 Exponential Backoff for Reconnection

- 👉 Retry database connection with increasing delays.

Example: DB connection lost → retry after 1s, 2s, 4s, 8s, 16s → max 5 attempts.

Give database time to fully restart

3 Connection Validation Query

- 👉 Periodically test idle connections.

Example: Every 30 seconds, run `SELECT 1` on idle connections → close if invalid.

Fast health verification

4 Maximum Lifetime & Idle Timeout

- 👉 Close long-lived or idle connections automatically.

Example: Max connection lifetime: 1 hour; Idle timeout: 10 minutes → prevents stale connections.

Prevent accumulating stale connections

5 Graceful Degradation During Outage

- 👉 Continue serving traffic with reduced capacity when systems are stressed.

Example: High load → disable non-essential features like recommendations → keep checkout working.

Maintain partial functionality

6 Read Replica Failover

- 👉 Automatically promote read replica to primary when primary fails.

Example: Primary DB crashes → replica promoted to primary within 60 seconds → minimal downtime.

High availability for reads

Topic 11: Database Indexing & Query Optimization

💡 Problem 31: When to Add an Index

Question: INTERVIEWER: Your query is slow on a users table with 10 million rows. When should you add an index?

Solution:

Indexes speed up reads but slow down writes — know when the tradeoff is worth it.

Here's when to add an index 👉

1 Frequently Queried Columns

- 👉 Add index on columns used in WHERE, JOIN, and ORDER BY clauses.

Example: Users table with frequent 'WHERE email = ?' queries → index on email column.

B-tree index PostgreSQL CREATE INDEX

2 High Selectivity Columns

- 👉 Index columns with many unique values (high cardinality).

Example: Email or user_id (unique) gets better index performance than gender (low cardinality).

Selectivity = unique_values / total_rows

3 Composite Indexes for Multi-Column Queries

- 👉 Create index on multiple columns when queries filter on those columns together.

Example: 'WHERE city = ? AND age > ?' → composite index on (city, age).

Leftmost prefix rule PostgreSQL multi-column indexes

4 Covering Indexes for SELECT Optimization

- 👉 Include all columns from SELECT clause in index to avoid table lookup.

Example: 'SELECT name, email WHERE user_id = ?' → index on (user_id, name, email).

Index-only scan INCLUDE clause in PostgreSQL

5 Avoid Over-Indexing Write-Heavy Tables

- 👉 Each index adds overhead to INSERT/UPDATE/DELETE operations.

Example: Table with 10 inserts/sec and 1 read/min → indexes hurt more than help.

Write amplification index maintenance cost

6 Use Partial Indexes for Subset Queries

- 👉 Index only rows matching a condition to save space and improve performance.

Example: 'WHERE status = 'active'' → partial index on only active rows.

PostgreSQL partial indexes

filtered indexes in SQL Server

💡 Problem 32: Optimizing Slow Queries

Question: INTERVIEWER: A query takes 5 seconds on production. How do you optimize it?

Solution:

Query optimization is like debugging — you need to see what's happening under the hood.

Here's how to optimize slow queries ↗

1 Use EXPLAIN to Analyze Query Plan

👉 Check how database executes the query — look for sequential scans, missing indexes.

Example: `EXPLAIN ANALYZE SELECT * FROM orders WHERE user_id = 123` → shows seq scan → add index.

EXPLAIN

EXPLAIN ANALYZE

execution plan

2 Rewrite N+1 Queries with JOINS

👉 Replace multiple single-row queries with one query using JOIN.

Example: 100 queries to fetch user for each order → 1 query with `JOIN users ON orders.user_id = users.id`.

Eager loading

LEFT JOIN

INNER JOIN

3 Add Indexes on Filter/Join Columns

👉 Index columns used in WHERE, JOIN, and ORDER BY clauses.

Example: Slow `WHERE created_at > ?` → add index on created_at.

B-tree index

composite index

4 Limit Result Set with LIMIT and Pagination

- 👉 Fetch only necessary rows instead of full table scan.

Example: `SELECT * FROM logs` (1M rows) → `SELECT * FROM logs LIMIT 100 OFFSET 0`.

[Cursor-based pagination](#) [keyset pagination](#)

5 Avoid SELECT * — Fetch Only Needed Columns

- 👉 Reduce data transfer by selecting specific columns.

Example: `SELECT *` transfers 20 columns → `SELECT id, name, email` transfers 3 columns.

[Projection](#) [covering index](#)

6 Denormalize for Read-Heavy Queries

- 👉 Duplicate data to avoid expensive JOINs in read-heavy workloads.

Example: Store user name in orders table instead of JOIN with users table on every query.

[Denormalization](#) [materialized views](#)

7 Use Database Query Cache

- 👉 Cache identical query results for repeated queries.

Example: Same dashboard query runs 100 times/min → cache result for 1 minute.

[Query result cache](#) [Redis](#)



Event-Driven Architecture

💡 Problem 33: Choosing the Right Index Type

Question: INTERVIEWER: You have a query searching for text patterns. What index type should you use?

Solution:

Different queries need different index types — B-tree isn't always the answer.

Here's how to choose the right index type 🤓

1 B-Tree Index for Equality and Range Queries

👉 Use for `=`, `<`, `>`, `BETWEEN`, `IN`, `ORDER BY` queries.

Example: `WHERE age > 30 AND age < 50` → B-tree index on age.

PostgreSQL default

MySQL InnoDB

2 Hash Index for Exact Match Queries

👉 Use for equality comparisons only — faster than B-tree but no range support.

Example: `WHERE user_id = 123` → hash index on user_id.

PostgreSQL HASH index

not range-searchable

3 Full-Text Index for Text Search

👉 Use for pattern matching and text search queries.

Example: `WHERE description LIKE '%backend%'` → full-text index on description.

PostgreSQL GIN/GIST

MySQL FULLTEXT

Elasticsearch

4 GiST/GIN Index for JSON and Array Queries

👉 Use for querying inside JSON fields or array contains operations.

Example: `WHERE tags @> ARRAY['python', 'backend']` → GIN index on tags.

PostgreSQL GIN for JSONB

array containment

5 Spatial Index for Geolocation Queries

👉 Use for location-based queries (nearby places, distance calculations).

Example: `WHERE ST_DWithin(location, point, 10km)` → spatial index on location.

PostGIS R-tree

MongoDB geospatial index

6 Bitmap Index for Low-Cardinality Columns

👉 Use for columns with few distinct values in data warehouse scenarios.

Example: `WHERE gender = 'M'` → bitmap index on gender (only 2-3 values).

Oracle Bitmap Index

not in PostgreSQL/MySQL



Problem 34: Index Maintenance and Monitoring

Question: INTERVIEWER: Your database performance degrades over time even though you have indexes. Why?

Solution:

Indexes aren't fire-and-forget — they need maintenance and monitoring.

Here's how to maintain indexes 👇

1 Monitor Index Usage

👉 Track which indexes are actually used vs unused indexes wasting space.

Example: Query `pg_stat_user_indexes` → find indexes with zero scans → drop them.

PostgreSQL pg_stat_user_indexes

MySQL INFORMATION_SCHEMA

2 Rebuild Fragmented Indexes

- 👉 Over time, indexes become fragmented and slower — rebuild to defragment.

Example: Like a fragmented hard drive slowing down → defrag improves performance.

PostgreSQL REINDEX MySQL OPTIMIZE TABLE

3 Update Table Statistics

- 👉 Database query planner uses statistics to choose optimal plan — keep them fresh.

Example: After bulk insert of 1M rows → run ANALYZE to update statistics.

PostgreSQL ANALYZE MySQL ANALYZE TABLE

4 Vacuum Dead Tuples in PostgreSQL

- 👉 DELETE/UPDATE leaves dead rows that bloat table and indexes — vacuum to reclaim space.

Example: Like emptying trash bin → reclaim disk space and improve performance.

PostgreSQL VACUUM autovacuum daemon

5 Check Index Bloat

- 👉 Indexes grow larger than necessary due to dead tuples and fragmentation.

Example: 500MB index should be 100MB → indicates bloat → rebuild index.

pg_stat_user_indexes pgstattuple extension

6 Monitor Query Plan Changes

- 👉 Database might stop using your index if statistics are outdated or query changes.

Example: EXPLAIN shows seq scan instead of index scan → update stats or fix query.

EXPLAIN ANALYZE query plan monitoring

Topic 12: Database Sharding

Problem 35: When to Shard Your Database

Question: INTERVIEWER: Your database has 500GB of data and queries are slowing down. Should you shard?

Solution:

Sharding is powerful but complex — make sure you actually need it first.

Here's when to shard your database 

1 Data Size Exceeds Single Server Capacity

 When your data grows beyond what one database server can handle.

Example: 5TB dataset → single PostgreSQL instance can't scale → shard across 10 servers (500GB each).

[Horizontal partitioning](#) [distributed database](#)

2 Write Traffic Exceeds Single Server Throughput

 When write operations bottleneck on a single database instance.

Example: 100K writes/sec → single database maxes at 10K writes/sec → shard to distribute writes.

[Write scaling](#) [distributed writes](#)

3 Read Replicas Can't Keep Up

 When adding more read replicas doesn't help because primary write bottleneck.

Example: Primary replication lag at 5 seconds even with 10 replicas → shard to reduce primary load.

[Replication lag](#) [write bottleneck](#)

4 Natural Data Boundaries Exist

- 👉 When your data has logical separation (multi-tenant, geographic regions).

Example: E-commerce platform → shard by region (US shard, EU shard, Asia shard).

Multi-tenancy geographic sharding

5 Before Sharding: Try These First

- 👉 Exhaust simpler options before adding sharding complexity.

Example: Vertical scaling, read replicas, caching, query optimization, table partitioning.

Avoid premature sharding operational complexity

💡 Problem 36: Zero-Downtime Database Migration

Question: INTERVIEWER: You need to add a NOT NULL column to a 500M-row table — how do you avoid downtime?

Solution:

Large database migrations require careful planning to avoid locking and downtime.

1 Add Column as Nullable First

- 👉 Add column without NOT NULL constraint initially.

Example: ALTER TABLE users ADD COLUMN status VARCHAR(20) → instant, no row updates.

Nullable first fast migration

2 Backfill in Batches

- 👉 Update existing rows in small batches.

Example: UPDATE users SET status='active' WHERE id BETWEEN 1 AND 100000 → repeat.

Batch updates gradual backfill

3 Add Default Value for New Rows

- 👉 Application sets value for all new inserts.

Example: App code adds status='active' for new users during backfill.

Application defaults dual writes

4 Add NOT NULL Constraint After Backfill

- 👉 Once all rows populated, add constraint.

Example: ALTER TABLE users MODIFY COLUMN status VARCHAR(20) NOT NULL → fast validation.

Constraint addition validation

5 Online Schema Change Tools

- 👉 Use tools that avoid locking.

Example: gh-ost (GitHub), pt-online-schema-change (Percona) → non-blocking migrations.

gh-ost pt-online-schema-change

6 Blue-Green Database Migration

- 👉 Migrate to new database, switch over when ready.

Example: Replicate to new DB → apply schema → catch up → switch traffic.

Blue-green DB replication

7 Rollback Plan

- 👉 Always have a way to undo the migration.

Example: Keep old column temporarily → if issues, revert code → remove new column later.

Rollback safety

💡 Problem 36: Zero-Downtime Database Migration

Question: INTERVIEWER: You need to add a NOT NULL column to a 500M-row table — how do you avoid downtime?

Solution:

Large database migrations require careful planning to avoid locking and downtime.

1 Add Column as Nullable First

- 👉 Add column without NOT NULL constraint initially.

Example: `ALTER TABLE users ADD COLUMN status VARCHAR(20)` → instant, no row updates.

Nullable first fast migration

2 Backfill in Batches

- 👉 Update existing rows in small batches.

Example: `UPDATE users SET status='active' WHERE id BETWEEN 1 AND 100000` → repeat.

Batch updates gradual backfill

3 Add Default Value for New Rows

- 👉 Application sets value for all new inserts.

Example: App code adds `status='active'` for new users during backfill.

Application defaults dual writes

4 Add NOT NULL Constraint After Backfill

- 👉 Once all rows populated, add constraint.

Example: `ALTER TABLE users MODIFY COLUMN status VARCHAR(20) NOT NULL` → fast validation.

Constraint addition validation

5 Online Schema Change Tools

- 👉 Use tools that avoid locking.

Example: *gh-ost* (GitHub), *pt-online-schema-change* (Percona) → non-blocking migrations.

`gh-ost` `pt-online-schema-change`

6 Blue-Green Database Migration

- 👉 Migrate to new database, switch over when ready.

Example: Replicate to new DB → apply schema → catch up → switch traffic.

`Blue-green DB` `replication`

7 Rollback Plan

- 👉 Always have a way to undo the migration.

Example: Keep old column temporarily → if issues, revert code → remove new column later.

`Rollback` `safety`



Microservices Architecture

💡 Problem 36: Sharding Strategies

Question: INTERVIEWER: You decided to shard. How do you choose which shard to send data to?

Solution:

The sharding key determines data distribution — choosing wrong can cause hotspots.

Here are common sharding strategies 🤞

1 Hash-Based Sharding

- 👉 Hash the sharding key (`user_id`, `order_id`) to determine shard.

Example: $\text{hash}(\text{user_id}) \% 10 \rightarrow \text{user 12345 always goes to shard 5.}$

Consistent hashing even distribution Redis Cluster

2 Range-Based Sharding

- 👉 Split data by value ranges of the sharding key.

Example: Users A-M → Shard 1, Users N-Z → Shard 2.

Sorted data range queries MongoDB range sharding

3 Geographic Sharding

- 👉 Shard by location to reduce latency and comply with data residency laws.

Example: US users → US shard, EU users → EU shard (GDPR compliance).

Multi-region data sovereignty latency optimization

4 Directory-Based Sharding

- 👉 Maintain lookup table mapping each key to its shard.

Example: `user_id 123` → Shard 2, `user_id 456` → Shard 5.

Flexible lookup overhead shard rebalancing

5 Entity Group Sharding

- 👉 Co-locate related data on same shard for transaction support.

Example: User and their orders on same shard → local transactions instead of distributed.

Data locality transaction support Google Spanner

6 Composite Sharding Key

- 👉 Use multiple columns for sharding to improve distribution.

Example: Shard by (`tenant_id`, `created_date`) → evenly distribute multi-tenant data.

Multi-tenancy time-series data

💡 Problem 37: Handling Shard Hotspots and Rebalancing

Question: INTERVIEWER: One of your shards is getting 80% of the traffic. How do you fix it?

Solution:

Uneven shard distribution causes hotspots — you need to rebalance or rethink sharding strategy.

Here's how to handle shard hotspots 🤞

1 Identify Hotspot Cause

👉 Monitor query patterns to find why one shard is overloaded.

Example: Celebrity user on Shard 3 → 1M followers querying their data → hotspot.

Query analysis shard metrics uneven distribution

2 Use Consistent Hashing to Minimize Rebalancing

👉 When adding/removing shards, only rebalance affected keys instead of all data.

Example: Adding Shard 11 with consistent hashing → only ~9% of data moves, not 100%.

Consistent hashing ring minimal data movement

3 Reshard Hot Keys to Separate Shard

👉 Move high-traffic keys to dedicated shard to isolate load.

Example: Celebrity user causing hotspot → move to dedicated shard.

Resharding data migration hotspot isolation

4 Use Finer-Grained Sharding

👉 Increase number of shards to distribute load more evenly.

Example: 4 shards with uneven load → 16 shards with better distribution.

More shards rebalancing complexity

5 Add Secondary Index Shard for Read-Heavy Keys

👉 Cache or replicate hot data separately for read queries.

Example: Popular product data → replicate to all shards or cache layer.

Read replicas caching denormalization

6 Implement Virtual Shards

👉 Create more logical shards than physical servers, then map multiple virtual shards to one physical shard.

Example: 1000 virtual shards mapped to 10 physical shards → easier rebalancing.

Virtual shards flexible mapping Vitess

Topic 13: Read Replicas & Replication

💡 Problem 38: Setting Up Read Replicas

Question: INTERVIEWER: Your application has heavy read traffic. How do you scale reads without affecting writes?

Solution:

Read replicas let you scale reads horizontally without overloading the primary database.

Here's how to set up read replicas 🤞

1 Asynchronous Replication to Replicas

- 👉 Primary writes data, then asynchronously replicates to read replicas.

Example: Like photocopying documents — original written first, copies made after.

PostgreSQL streaming replication MySQL replication

2 Route Reads to Replicas, Writes to Primary

- 👉 Application logic or database proxy routes read queries to replicas.

Example: 90% read traffic → 10 replicas handle reads, primary handles writes only.

Read/write splitting ProxySQL AWS RDS read replicas

3 Monitor Replication Lag

- 👉 Track delay between primary write and replica availability.

Example: Primary writes at 10:00:00, replica shows data at 10:00:05 → 5-second lag.

Replication lag monitoring pg_stat_replication

4 Use Multiple Replicas for Geographic Distribution

- 👉 Place replicas in different regions for lower latency reads.

Example: US primary → EU replica → EU users read from EU replica (lower latency).

Multi-region geographic distribution

5 Handle Eventual Consistency

- 👉 Accept that replicas may show stale data due to replication lag.

Example: User updates profile → immediately reads from replica → sees old data (read-after-write inconsistency).

Eventual consistency read-your-writes problem

💡 Problem 39: Handling Replication Lag

Question: INTERVIEWER: A user updates their profile but doesn't see the change. What's happening?

Solution:

Replication lag causes read-after-write inconsistency — you need strategies to handle it.

Here's how to handle replication lag 🤔

1 Read Your Own Writes from Primary

👉 Route reads to primary for data the user just modified.

Example: User updates profile → next profile read goes to primary, not replica.

Session affinity sticky routing

2 Use Synchronous Replication for Critical Data

👉 Primary waits for at least one replica to acknowledge write before returning success.

Example: Financial transaction → wait for replica confirmation → guarantees data available on replica.

Synchronous replication trade latency for consistency

3 Check Replication Position Before Read

👉 Query replica only if it has caught up to a specific replication position.

Example: Write returns position 1000 → read waits until replica reaches position 1000.

LSN tracking replication position

4 Use Monotonic Reads with Session Token

👉 Ensure a user always reads from replica at least as fresh as their last read.

Example: Session token tracks last seen timestamp → only query replicas ahead of that timestamp.

Monotonic reads session consistency

5 Add Cache Layer for Immediate Reads

👉 Cache recent writes so reads return cached data instead of stale replica data.

Example: Profile update → write to database and Redis → read from Redis (no lag).

Write-through cache Redis

6 Display Staleness to User

👉 Show users that data might be slightly out of date.

Example: "Data as of 5 seconds ago" message on dashboard.

User expectations eventual consistency UX



API Gateway

💡 Problem 40: Promoting Replica to Primary (Failover)

Question: INTERVIEWER: Your primary database crashes. How do you promote a replica to become the new primary?

Solution:

Database failover is critical for high availability — automate it or risk extended downtime.

Here's how to promote a replica to primary 🙌

1 Detect Primary Failure

👉 Use health checks and heartbeat monitoring to detect when primary is down.

Example: Primary doesn't respond to health check for 10 seconds → trigger failover.

Health checks Patroni AWS RDS automated failover

2 Stop Writes to Failed Primary

- 👉 Ensure no new writes go to the failed primary to prevent data conflicts.

Example: Update DNS or load balancer to stop routing traffic to old primary.

Fencing STONITH

3 Select Replica with Most Recent Data

- 👉 Choose replica with lowest replication lag as new primary.

Example: Replica A at position 1000, Replica B at position 995 → promote Replica A.

LSN comparison replication position

4 Promote Replica to Primary

- 👉 Reconfigure selected replica to accept writes and become new primary.

Example: `pg_promote()` in PostgreSQL → replica starts accepting writes.

Promote command configuration change

5 Reconfigure Other Replicas

- 👉 Point remaining replicas to replicate from the new primary.

Example: Replica B and C now replicate from Replica A (promoted to primary).

Replication topology change

6 Update Application Connection Strings

- 👉 Change application configuration or DNS to point to new primary.

Example: Update database endpoint from old-primary.com to new-primary.com.

DNS update connection failover service discovery

7 Use Automated Failover Tools

- 👉 Tools like Patroni, Orchestrator, or cloud-managed services automate entire failover process.

Example: AWS RDS Multi-AZ automatically fails over within 1-2 minutes.

Topic 14: Connection Pooling

💡 Problem 41: Why Use Connection Pooling

Question: INTERVIEWER: Your API creates a new database connection for every request. Why is this slow?

Solution:

Creating database connections is expensive — connection pooling reuses connections.

Here's why you need connection pooling 🤔

1 Connection Creation is Expensive

- 👉 Each new connection requires TCP handshake, authentication, and resource allocation.

Example: Creating connection takes 50ms → 1000 req/sec = 50,000ms wasted on connections.

2 Connection Pool Reuses Connections

- 👉 Pre-create connections and reuse them across requests.

Example: Like rental cars — instead of buying a car for each trip, rent from a shared pool.

3 Limit Max Connections to Database

- 👉 Database has maximum connection limit — pool prevents exhausting it.

Example: PostgreSQL max_connections = 100 → connection pool limits app to 100 connections.

Max connections resource limits

4 Reduce Database Connection Overhead

- 👉 Fewer connections means less memory and CPU usage on database server.

Example: 1000 idle connections consume GBs of memory → 20 pooled connections consume MBs.

Memory usage connection overhead

5 Improve Application Response Time

- 👉 Reusing connections removes connection setup latency from request path.

Example: 50ms connection setup → with pooling, 0ms connection setup.

Latency reduction connection reuse

💡 Problem 42: Configuring Connection Pool Size

Question: INTERVIEWER: How many connections should you have in your connection pool?

Solution:

Too few connections cause queuing, too many overwhelm the database — find the sweet spot.

Here's how to configure connection pool size 👉

1 Start with Formula: $\text{connections} = ((\text{core_count} * 2) + \text{effective_spindle_count})$

- 👉 For SSDs, effective_spindle_count ≈ 1. For 4-core DB server → pool size ≈ 9.

Example: Like a restaurant — too few servers (connections) = long wait, too many = chaos.

HikariCP recommendation empirical formula

2 Consider Database Max Connections

- 👉 Total connections across all app instances must stay below database limit.

Example: PostgreSQL max_connections = 100 → 10 app servers → 10 connections per server.

Max connections constraint distributed apps

3 Monitor Connection Usage

- 👉 Track active vs idle connections to see if pool is too small or too large.

Example: Pool size 20, always 19-20 active → increase pool size.

Connection metrics pool saturation

4 Use Connection Pool Per Service

- 👉 Each microservice or app instance gets its own connection pool.

Example: 5 app servers with 20 connections each = 100 total connections to database.

Distributed pools connection accounting

5 Set Connection Timeouts

- 👉 Configure how long app waits for available connection before timing out.

Example: Connection timeout = 10 seconds → if pool exhausted, request fails after 10s.

Connection timeout pool exhaustion handling

6 Use PgBouncer for Transaction-Level Pooling

- 👉 External connection pooler that multiplexes many app connections to fewer database connections.

Example: 1000 app connections → PgBouncer → 20 database connections.

PgBouncer transaction pooling mode connection multiplexing



Service Mesh

💡 Problem 43: Connection Pooling Issues

Question: INTERVIEWER: Your connection pool shows all connections idle, but requests are timing out. What's wrong?

Solution:

Connection pools have gotchas — idle connections might be dead, or pool might be misconfigured.

Here's how to debug connection pooling issues 🤞

1 Dead Connections in Pool

👉 Database closed idle connections but pool thinks they're still valid.

Example: Database closes connections after 5 minutes idle → app tries to use dead connection → query fails.

Connection validation testOnBorrow

2 Enable Connection Validation

👉 Test connections before using them from pool.

Example: Run 'SELECT 1' before giving connection to app → if fails, discard and create new one.

Connection health check HikariCP validationQuery

3 Set Max Connection Lifetime

👉 Close and recreate connections periodically to avoid stale connections.

Example: Max lifetime 30 minutes → connection used for 30 min → close and create new one.

Connection refresh maxLifetime

4 Connection Leak Detection

- 👉 Detect when app doesn't return connections to pool after use.

Example: App queries database but forgets to close connection → pool exhausted → timeouts.

Leak detection connection leak timeout

5 Firewall Killing Idle Connections

- 👉 Network firewall drops TCP connections after idle timeout.

Example: AWS NAT gateway drops connections idle >350 seconds → configure keepalive < 350s.

TCP keepalive firewall idle timeout

6 Pool Size Too Small for Workload

- 👉 More concurrent requests than available connections.

Example: 100 concurrent requests, pool size 10 → 90 requests wait for connection → timeouts.

Pool saturation increase pool size

Topic 15: ACID Transactions

💡 Problem 44: Explaining ACID Properties

Question: INTERVIEWER: What does ACID mean in databases?

Solution:

ACID guarantees ensure database transactions are reliable and consistent.

Here's what ACID means 👉

1 Atomicity

- 👉 Transaction either fully completes or fully fails — no partial updates.

Example: Bank transfer — debit \$100 from A and credit \$100 to B → both succeed or both fail.

All-or-nothing rollback on failure

2 Consistency

- 👉 Transaction takes database from one valid state to another valid state.

Example: Constraint "balance >= 0" → transaction that would make balance negative is rejected.

Database constraints invariants

3 Isolation

- 👉 Concurrent transactions don't interfere with each other.

Example: Two users buying last item — both see item available, but only one succeeds.

Concurrency control locks MVCC

4 Durability

- 👉 Committed transaction changes are permanent, even after crash.

Example: Transaction commits at 10:00:00, server crashes at 10:00:01 → data still saved.

Write-ahead log fsync persistent storage



Circuit Breaker Pattern

💡 Problem 45: Transaction Isolation Levels

Question: INTERVIEWER: Two transactions read and update the same row. What isolation level do you use?

Solution:

Isolation levels trade off consistency for performance — choose based on your requirements.

Here are the transaction isolation levels 🤞

1 Read Uncommitted (Lowest Isolation)

👉 Transaction can read uncommitted changes from other transactions.

Example: Transaction A updates row, Transaction B reads it before A commits → dirty read.

Dirty reads possible rarely used

2 Read Committed (Default in PostgreSQL/Oracle)

👉 Transaction only reads committed data from other transactions.

Example: Transaction A updates row, Transaction B can only read value after A commits.

No dirty reads prevents reading uncommitted data

3 Repeatable Read

👉 Transaction sees consistent snapshot of data throughout — no non-repeatable reads.

Example: Transaction A reads row twice → sees same value both times even if another transaction updates it.

No non-repeatable reads snapshot isolation

4 Serializable (Highest Isolation)

- 👉 Transactions execute as if they ran serially, one after another.

Example: Like a single-threaded program — no concurrency anomalies possible.

Full isolation performance cost deadlock risk

5 Trade-off: Isolation vs Performance

- 👉 Higher isolation = more locks/overhead = lower throughput.

Example: Serializable prevents all anomalies but can be 10x slower than Read Committed.

Lock contention MVCC choose based on requirements

💡 Problem 46: Handling Distributed Transactions

Question: INTERVIEWER: You need to update two different databases atomically. How do you ensure both succeed or both fail?

Solution:

Distributed transactions are hard — you need coordination protocols or eventual consistency.

Here's how to handle distributed transactions 👇

1 Two-Phase Commit (2PC)

- 👉 Coordinator asks all participants to prepare, then commits if all agree.

Example: Coordinator asks DB1 and DB2 "ready to commit?" → both say yes → coordinator says "commit now".

XA transactions blocking protocol coordinator failure risk

2 Saga Pattern with Compensating Transactions

- 👉 Break distributed transaction into local transactions with compensation logic for failures.

Example: Order service creates order → payment fails → order service cancels order (compensating action).

Choreography or orchestration eventual consistency

3 Use Message Queue for Reliable Delivery

- 👉 Write to database and publish message atomically using outbox pattern.

Example: Insert order in DB + insert event in outbox table → background job publishes event.

Outbox pattern transactional outbox Kafka

4 Avoid Distributed Transactions with Denormalization

- 👉 Duplicate data across services to avoid cross-service transactions.

Example: Store user email in orders table instead of JOIN with users service.

Denormalization data duplication eventual consistency

5 Use Idempotency for Retry Safety

- 👉 Design operations to be safely retried without side effects.

Example: Payment service receives same payment request twice → charges only once.

Idempotency keys retry safety

6 Google Spanner / CockroachDB for True Distributed ACID

- 👉 Use databases that provide ACID guarantees across multiple nodes.

Example: Like a magical database that coordinates transactions globally — but complex and expensive.

TrueTime distributed consensus Spanner CockroachDB

💡 Problem 47: Optimistic vs Pessimistic Locking

Question: INTERVIEWER: Two users try to book the last hotel room at the same time — how do you handle it?

Solution:

Locking strategies balance data consistency with system performance.

1 Pessimistic Locking

- 👉 Lock row immediately, preventing concurrent access.

Example: `SELECT * FROM rooms WHERE id=123 FOR UPDATE` → blocks other transactions.

Row locking `FOR UPDATE`

2 Pessimistic Lock Downsides

- 👉 Blocks other users, reduces concurrency, risk of deadlocks.

Example: 100 users trying to book → 99 wait → slow user experience.

Low concurrency deadlocks

3 Optimistic Locking

- 👉 Don't lock, detect conflicts at commit time.

Example: Add version column → `UPDATE WHERE version=old_version` → fails if changed.

Version column conflict detection

4 Optimistic Lock Failure Handling

- 👉 Retry with updated data when conflict detected.

Example: Version mismatch → reload data → retry update → or show error to user.

Retry logic user notification

5 When to Use Pessimistic

- 👉 High contention, critical consistency needs.

Example: Financial transactions, inventory with very limited stock.

High contention strict consistency

6 When to Use Optimistic

- 👉 Low contention, performance-critical reads.

Example: Social media likes, view counts → rare conflicts → optimistic faster.

Low contention performance

7 Hybrid Approach

- 👉 Use optimistic by default, pessimistic for critical operations.

Example: Product catalog updates → optimistic; Final checkout → pessimistic.

Hybrid strategy balanced approach



Problem 47: Optimistic vs Pessimistic Locking

Question: INTERVIEWER: Two users try to book the last hotel room at the same time — how do you handle it?

Solution:

Locking strategies balance data consistency with system performance.

1 Pessimistic Locking

- 👉 Lock row immediately, preventing concurrent access.

Example: `SELECT * FROM rooms WHERE id=123 FOR UPDATE` → blocks other transactions.

Row locking FOR UPDATE

2 Pessimistic Lock Downsides

- 👉 Blocks other users, reduces concurrency, risk of deadlocks.

Example: 100 users trying to book → 99 wait → slow user experience.

Low concurrency deadlocks

3 Optimistic Locking

- 👉 Don't lock, detect conflicts at commit time.

Example: Add version column → UPDATE WHERE version=old_version → fails if changed.

Version column conflict detection

4 Optimistic Lock Failure Handling

- 👉 Retry with updated data when conflict detected.

Example: Version mismatch → reload data → retry update → or show error to user.

Retry logic user notification

5 When to Use Pessimistic

- 👉 High contention, critical consistency needs.

Example: Financial transactions, inventory with very limited stock.

High contention strict consistency

6 When to Use Optimistic

- 👉 Low contention, performance-critical reads.

Example: Social media likes, view counts → rare conflicts → optimistic faster.

Low contention performance

7 Hybrid Approach

- 👉 Use optimistic by default, pessimistic for critical operations.

Example: Product catalog updates → optimistic; Final checkout → pessimistic.

Hybrid strategy

balanced approach

Topic 16: NoSQL Data Modeling



Rate Limiting



Problem 47: When to Use NoSQL vs SQL

Question: INTERVIEWER: Your app needs to store user profiles with flexible attributes. Should you use SQL or NoSQL?

Solution:

SQL and NoSQL excel at different use cases — choose based on data structure and access patterns.

Here's when to use NoSQL vs SQL ↗

1 Use SQL for Structured Data with Relationships

👉 Data has well-defined schema and relationships between entities.

Example: E-commerce orders, users, products with foreign keys and JOINs.

PostgreSQL

MySQL

ACID transactions

referential integrity

2 Use NoSQL for Flexible Schema

👉 Each document/row can have different fields without schema changes.

Example: User profiles where some users have 5 fields, others have 20 different fields.

MongoDB DynamoDB schema flexibility

3 Use NoSQL for High Write Throughput

- 👉 Distributed NoSQL databases scale writes horizontally better than SQL.

Example: Logging system with 1M writes/sec → Cassandra distributes across 100 nodes.

Cassandra DynamoDB eventual consistency

4 Use SQL for Complex Queries and Aggregations

- 👉 Need JOINs, GROUP BY, complex filtering, and ad-hoc queries.

Example: Analytics dashboard querying across multiple tables with aggregations.

SQL JOINs GROUP BY HAVING complex WHERE clauses

5 Use NoSQL for Key-Value or Document Access Patterns

- 👉 Primarily access data by primary key without complex joins.

Example: User session storage — fetch session by session_id.

Redis DynamoDB MongoDB

6 Use SQL for Strong Consistency Requirements

- 👉 Need ACID guarantees and immediate consistency.

Example: Banking system where balance must be accurate immediately.

PostgreSQL MySQL ACID

7 Use NoSQL for Horizontal Scalability

- 👉 Data grows beyond single server capacity and needs distributed architecture.

Example: Social media platform with billions of posts across thousands of servers.

Cassandra MongoDB sharding DynamoDB

💡 Problem 48: Modeling Data in Document Databases (MongoDB)

Question: INTERVIEWER: How do you model a blog with users, posts, and comments in MongoDB?

Solution:

Document databases let you embed or reference related data — choose based on access patterns.

Here's how to model data in document databases 🤓

1 Embed Related Data in Same Document

👉 Store related data together when accessed together.

Example: Blog post with comments → embed comments array inside post document.

Atomic updates

single query retrieval

2 Reference with Document IDs

👉 Store reference to another document's ID when data is large or accessed separately.

Example: Post has author_id → separate users collection → JOIN in application code.

Normalization

avoid document size limits

3 Embed for One-to-Few Relationships

👉 Embed when child documents are small and limited in number.

Example: User has 3 addresses → embed addresses array in user document.

Simplicity

atomic updates

4 Reference for One-to-Many Relationships

👉 Use references when child documents are numerous or accessed independently.

Example: Author has 1000 blog posts → posts collection with author_id field.

Avoid unbounded arrays

separate access patterns

5 Denormalize for Read Performance

👉 Duplicate data to avoid lookups when querying.

Example: Store author name in post document even though it's in users collection.

Faster reads data duplication eventual consistency

6 Use Two-Way Referencing for Many-to-Many

👉 Store references in both documents for efficient bidirectional queries.

Example: Students and courses → student has course_ids, course has student_ids.

Query from either direction data duplication

💡 Problem 49: Modeling Data in Wide-Column Stores (Cassandra)

Question: INTERVIEWER: How do you design a time-series data model in Cassandra for sensor readings?

Solution:

Wide-column stores require modeling based on query patterns — denormalize and duplicate data.

Here's how to model data in wide-column stores 👉

1 Design Tables Based on Queries, Not Entities

👉 Create separate table for each query pattern, duplicating data across tables.

Example: Query by sensor_id → table with sensor_id as partition key. Query by location → different table with location as partition key.

Query-driven design data duplication

2 Choose Partition Key for Even Distribution

- 👉 Partition key determines which node stores data — avoid hotspots.

Example: sensor_id as partition key distributes 1000 sensors across cluster. Using country as partition key creates hotspots.

Even distribution avoid hotspots

3 Use Clustering Key for Sorting Within Partition

- 👉 Clustering key determines sort order within partition.

Example: Partition key = sensor_id, clustering key = timestamp → all readings for sensor sorted by time.

Range queries efficient time-series queries

4 Keep Partition Size Manageable

- 👉 Avoid unbounded partitions that grow infinitely.

Example: Partition by (sensor_id, date) instead of just sensor_id → creates new partition each day.

Partition size limits time bucketing

5 Denormalize for Read Efficiency

- 👉 Duplicate data to serve different query patterns without JOINs.

Example: Sensor readings table AND sensor readings by location table — same data, different partition keys.

No JOINs write amplification

6 Use Materialized Views for Multiple Query Patterns

- 👉 Cassandra automatically maintains additional tables for different queries.

Example: Base table partitioned by sensor_id → materialized view partitioned by location.

Automatic denormalization Cassandra feature

Topic 17: Eventually Consistent Systems



Authentication & Authorization



Problem 50: Understanding Eventual Consistency

Question: INTERVIEWER: Your distributed system shows different data on different servers. Is this a bug?

Solution:

Eventual consistency is a trade-off — accept temporary inconsistency for availability and performance.

Here's how eventual consistency works 🤓

1 What is Eventual Consistency?

👉 Replicas may temporarily show different data, but will eventually converge.

Example: Update post on Server A → Server B shows old data for 2 seconds → eventually both show same data.

CAP theorem

AP systems

2 Why Accept Eventual Consistency?

👉 Enables high availability and low latency by not waiting for all replicas to update.

Example: Amazon cart allows adding items even if other data centers are unreachable.

Availability over consistency

partition tolerance

3 Read-After-Write Inconsistency

👉 User writes data but immediately reads old data from replica.

Example: User posts comment → refresh page → doesn't see comment (replica lag).

Replication lag user experience issue

4 Conflict Resolution with Last-Write-Wins

👉 When conflicting updates occur, use timestamp to decide winner.

Example: User edits profile on two devices → both writes succeed → later timestamp wins.

LWW vector clocks Cassandra

5 Use Cases for Eventual Consistency

👉 Non-critical data where temporary inconsistency is acceptable.

Example: Social media likes counter, view counts, caching, DNS.

Best-effort consistency high availability

6 When NOT to Use Eventual Consistency

👉 Financial transactions, inventory management, or any data requiring immediate accuracy.

Example: Bank balance must be accurate immediately — can't accept stale data.

Strong consistency required ACID transactions

💡 Problem 51: Handling Conflicts in Eventually Consistent Systems

Question: INTERVIEWER: Two users edit the same document at the same time in different data centers.
How do you resolve the conflict?

Solution:

Conflicts are inevitable in eventually consistent systems — you need resolution strategies.

Here's how to handle conflicts 🤞

1 Last-Write-Wins (LWW)

- 👉 Use timestamp to determine which update wins.

Example: User A updates at 10:00:00, User B updates at 10:00:05 → B's update wins.

Simple but data loss Cassandra DynamoDB

2 Version Vectors / Vector Clocks

- 👉 Track causality of updates to detect concurrent modifications.

Example: [A:1, B:0] and [A:0, B:1] → both concurrent → conflict detected.

Riak Cassandra causality tracking

3 Merge Conflicts Automatically with CRDTs

- 👉 Use Conflict-free Replicated Data Types that merge updates deterministically.

Example: Two users add items to shopping cart → both additions preserved in merged cart.

CRDTs Riak Redis

4 Application-Level Conflict Resolution

- 👉 Store all conflicting versions and let application decide which to keep.

Example: Google Docs tracks all changes and merges text edits operationally.

Custom merge logic operational transformation

5 Prevent Conflicts with Distributed Locks

- 👉 Use locking to ensure only one writer at a time.

Example: User acquires lock on document → edits → releases lock → other user can edit.

Redlock ZooKeeper trade availability for consistency

6 Use Strong Consistency for Critical Operations

- 👉 Switch to synchronous replication for operations that can't have conflicts.

Example: Decrement inventory with strong consistency → prevent overselling.

Quorum writes synchronous replication

Problem 52: Implementing Distributed Locks

Question: INTERVIEWER: Two servers try to process the same job simultaneously — how do you prevent that?

Solution:

Distributed locks coordinate access to shared resources across multiple processes.

1 Redis Distributed Lock (Redlock)

👉 Use Redis SET NX EX for simple distributed locking.

Example: `SET lock:job123 server1 NX EX 30` → only one server acquires lock.

Redlock Redis locks

2 Lock Expiration

👉 Auto-release locks after timeout to prevent deadlocks.

Example: Lock expires after 30 seconds → if server crashes, lock auto-released.

TTL auto-release

3 Lock Renewal (Heartbeat)

👉 Extend lock if job still running.

Example: Job takes 60s → renew lock every 20s → prevents premature expiration.

Lock renewal heartbeat

4 Zookeeper Locks

- 👉 Use Zookeeper ephemeral nodes for locks.

Example: Create /locks/job123 → if client dies, node auto-deleted → lock released.

Zookeeper ephemeral nodes

5 Database-Based Locks

- 👉 Use database row with unique constraint for lock.

Example: INSERT INTO locks (resource, owner) → unique constraint prevents duplicate.

Database locks row locking

6 Handling Split-Brain

- 👉 Prevent two nodes thinking they both hold lock.

Example: Use fencing tokens → each lock gets incrementing ID → reject stale locks.

Fencing tokens split-brain

7 When to Avoid Locks

- 👉 Consider lock-free algorithms or idempotency instead.

Example: Process job idempotently → duplicate processing is safe → no lock needed.

Lock-free idempotency



Problem 52: Implementing Distributed Locks

Question: INTERVIEWER: Two servers try to process the same job simultaneously — how do you prevent that?

Solution:

Distributed locks coordinate access to shared resources across multiple processes.

1 Redis Distributed Lock (Redlock)

- 👉 Use Redis SET NX EX for simple distributed locking.

Example: `SET lock:job123 server1 NX EX 30` → only one server acquires lock.

Redlock

Redis locks

2 Lock Expiration

- 👉 Auto-release locks after timeout to prevent deadlocks.

Example: Lock expires after 30 seconds → if server crashes, lock auto-released.

TTL

auto-release

3 Lock Renewal (Heartbeat)

- 👉 Extend lock if job still running.

Example: Job takes 60s → renew lock every 20s → prevents premature expiration.

Lock renewal

heartbeat

4 Zookeeper Locks

- 👉 Use Zookeeper ephemeral nodes for locks.

Example: Create `/locks/job123` → if client dies, node auto-deleted → lock released.

Zookeeper

ephemeral nodes

5 Database-Based Locks

- 👉 Use database row with unique constraint for lock.

Example: `INSERT INTO locks (resource, owner)` → unique constraint prevents duplicate.

Database locks

row locking

6 Handling Split-Brain

- 👉 Prevent two nodes thinking they both hold lock.

Example: Use fencing tokens → each lock gets incrementing ID → reject stale locks.

Fencing tokens split-brain

7 When to Avoid Locks

👉 Consider lock-free algorithms or idempotency instead.

Example: Process job idempotently → duplicate processing is safe → no lock needed.

Lock-free idempotency

Topic 18: N+1 Query Problem

💡 Problem 52: Identifying N+1 Query Problem

Question: INTERVIEWER: Your API endpoint is slow and making 1000 database queries. What's wrong?

Solution:

N+1 problem happens when you query in a loop — one query for list + N queries for details.

Here's how to identify N+1 queries 👇

1 What is N+1 Query Problem?

👉 Fetch N items in one query, then loop through items making N additional queries.

Example: Fetch 100 users → loop through users → fetch posts for each user → 101 queries total.

Loop queries performance killer

2 Signs of N+1 Problem

👉 Number of queries grows with number of records, slow response times.

Example: API takes 10ms with 1 user, 1000ms with 100 users → linear growth = $N+1$ issue.

Query count scales with data database CPU spike

3 Enable Query Logging to Detect

- 👉 Log all SQL queries to see if queries are running in a loop.

Example: Enable slow query log → see same query pattern repeated 100 times with different IDs.

PostgreSQL log_statement MySQL slow query log

4 Use Database Query Profiler

- 👉 Monitor query patterns in production to identify repeated queries.

Example: APM tool shows 'SELECT * FROM posts WHERE user_id = ?' executed 100 times.

New Relic Datadog query monitoring

5 Check ORM Query Behavior

- 👉 ORMs can hide $N+1$ queries behind lazy loading.

Example: Django ORM 'users = User.objects.all()' then 'user.posts.all()' in template → $N+1$ queries.

Lazy loading ORM gotcha

💡 Problem 3: Lazy vs Eager Loading Trade-offs

Question: INTERVIEWER: When should you use lazy loading vs eager loading for related data?

Solution:

Loading strategies balance memory usage with query efficiency.

1 Lazy Loading

- 👉 Load related data only when accessed.

Example: Fetch user → access user.posts → separate query to fetch posts.

On-demand loading deferred queries

2 Eager Loading

- 👉 Load related data upfront with JOIN.

*Example: SELECT * FROM users JOIN posts → all data in one query.*

Upfront loading JOIN queries

3 Lazy Loading Benefits

- 👉 Don't load unnecessary data, saves memory and query time.

Example: User profile page doesn't need posts → lazy loading avoids fetching them.

Memory efficient selective loading

4 Eager Loading Benefits

- 👉 Avoid N+1 queries when you know you'll need related data.

Example: User list with post counts → eager load to avoid 100 separate queries.

Prevents N+1 fewer queries

5 N+1 Query Problem

- 👉 Lazy loading can cause N+1 queries if not careful.

Example: Loop through 100 users → access user.posts each time → 101 queries total.

Performance pitfall query explosion

6 Selective Eager Loading

- 👉 Eager load only specific relations you need.

Example: User.findAll({ include: ['posts'] }) → loads posts, not comments.

Selective includes controlled loading

7 Pagination Considerations

- 👉 Eager loading with pagination can be complex.

Example: Paginate users with eager-loaded posts → need careful query design.

Pagination complexity query optimization

💡 Problem 53: Solving N+1 Query Problem

Question: INTERVIEWER: You identified N+1 queries fetching users and their posts. How do you fix it?

Solution:

Fix N+1 by fetching all related data in fewer queries using JOINs or batch loading.

Here's how to solve N+1 queries 🤞

1 Use JOIN to Fetch Related Data

- 👉 Fetch parent and child data in single query with JOIN.

Example: `SELECT users. , posts.* FROM users LEFT JOIN posts ON users.id = posts.user_id` → 1 query instead of N+1.*

SQL JOIN eager loading

2 Use ORM Eager Loading

- 👉 Tell ORM to prefetch related data upfront.

Example: Django `User.objects.prefetch_related('posts')` → 2 queries (users + posts) instead of N+1.

prefetch_related select_related in Django/Rails

3 Batch Load with IN Clause

👉 Fetch all related records in one query using IN clause.

*Example: Fetch 100 users → collect user IDs → 'SELECT * FROM posts WHERE user_id IN (1,2,3...100)' → 2 queries total.*

[Batch loading](#) [dataloader pattern](#)

4 Use DataLoader Library

👉 Automatically batch and cache database requests within single request.

Example: GraphQL dataloader batches all user queries → fetches in single query with IN clause.

[Facebook DataLoader](#) [automatic batching](#)

5 Cache Related Data

👉 Cache frequently accessed related data to avoid queries entirely.

Example: Cache user's posts in Redis → fetch from cache instead of database.

[Redis](#) [Memcached](#) [query caching](#)

6 Denormalize Data

👉 Duplicate data to avoid JOINs entirely.

Example: Store author name in posts table → no need to query users table.

[Denormalization](#) [data duplication](#)



SQL vs NoSQL Trade-offs

Topic 19: SQL vs NoSQL Trade-offs



OAuth 2.0 & OpenID Connect

💡 Problem 54: Migrating from SQL to NoSQL

Question: INTERVIEWER: Your PostgreSQL database is hitting scaling limits. Should you migrate to NoSQL?

Solution:

Migration is costly and risky — make sure NoSQL actually solves your problem.

Here's how to decide on migrating to NoSQL 🤔

1 Identify Why SQL is Hitting Limits

👉 Is it query complexity, write throughput, data size, or scaling cost?

Example: If slow queries → optimize indexes first. If write throughput → maybe NoSQL.

Diagnose root cause exhaust SQL optimizations first

2 Consider Hybrid Approach

👉 Keep SQL for transactional data, add NoSQL for specific use cases.

Example: PostgreSQL for orders/payments + Redis for caching + Cassandra for logs.

Polyglot persistence best tool for each job

3 Understand What You're Giving Up

👉 NoSQL means losing JOINs, transactions, foreign keys, and mature tooling.

Example: Multi-table reports requiring JOINs become application-side joins (slow and complex).

Trade-offs application complexity

4 Evaluate Schema Flexibility Needs

- 👉 Do you actually need flexible schema or is it premature optimization?

Example: User profiles with varying fields → MongoDB. Fixed schema → stick with SQL.

JSONB in PostgreSQL can provide flexibility without full migration

5 Consider NewSQL Databases

- 👉 Get SQL semantics with horizontal scalability.

Example: CockroachDB, Google Spanner, TiDB → SQL interface with distributed architecture.

NewSQL best of both worlds higher cost

6 Plan Incremental Migration

- 👉 Don't do big-bang migration — move one service/table at a time.

Example: Migrate analytics data to Cassandra first, keep transactional data in SQL.

Risk mitigation incremental approach dual-write pattern

💡 Problem 55: Handling Transactions Without ACID

Question: INTERVIEWER: You moved to DynamoDB but now need to update multiple items atomically.
How?

Solution:

NoSQL databases lack multi-item transactions — you need workarounds.

Here's how to handle transactions in NoSQL 👉

1 Use Database-Specific Transaction Features

- 👉 Some NoSQL databases support limited transactions.

Example: DynamoDB transactions support up to 25 items, MongoDB supports multi-document transactions.

Native transaction support check database capabilities

2 Implement Two-Phase Commit in Application

- 👉 Manually coordinate writes across multiple items with prepare/commit phases.

Example: Mark items as "pending" → verify all ready → commit all → mark as "complete".

Application-level 2PC complex error handling

3 Use Saga Pattern with Compensating Actions

- 👉 Break transaction into steps with rollback logic for each step.

Example: Reserve inventory → charge payment → if payment fails → unreserve inventory.

Compensating transactions eventual consistency

4 Design Data Model to Avoid Multi-Item Transactions

- 👉 Embed related data in same document/item for atomic updates.

Example: Store order items inside order document → update entire document atomically.

Denormalization data locality

5 Use Conditional Writes for Optimistic Locking

- 👉 Update only if item hasn't changed since you read it.

Example: Read item with version=5 → update with condition "version=5" → fails if someone else updated.

Optimistic concurrency DynamoDB conditional writes MongoDB update operators

6 Accept Eventual Consistency

- 👉 Design system to tolerate temporary inconsistencies.

Example: Shopping cart shows approximate inventory — occasional oversells handled by customer service.

Eventual consistency business logic accommodation



Session Management



Problem 56: Querying Patterns in NoSQL

Question: INTERVIEWER: In SQL you'd write complex queries with JOINs and WHERE clauses. How do you do this in NoSQL?

Solution:

NoSQL databases optimize for specific query patterns — you need to design schema around queries.

Here's how to handle querying in NoSQL

1 Denormalize Data for Query Patterns

👉 Duplicate data so each query hits one table/collection.

Example: Store user name in posts collection even though it's in users → no JOIN needed.

Data duplication write amplification faster reads

2 Create Multiple Tables for Different Access Patterns

👉 In wide-column stores, create separate table per query.

Example: Posts by author table + posts by date table → same data, different partition keys.

Cassandra DynamoDB query-driven design

3 Use Secondary Indexes Cautiously

👉 NoSQL secondary indexes are slower than SQL indexes — use sparingly.

Example: DynamoDB GSI for querying by email → slower than primary key access.

Global secondary index additional cost eventual consistency

4 Perform JOINs in Application Code

- 👉 Fetch related data in multiple queries and join in application.

Example: Fetch posts → collect author IDs → fetch authors with IN query → merge in code.

Application-side joins N+1 problem dataloader pattern

5 Use Aggregation Pipelines

- 👉 Document databases support aggregation frameworks for complex queries.

Example: MongoDB aggregation → \$match, \$group, \$sort → similar to SQL GROUP BY.

MongoDB aggregation limited compared to SQL

6 Pre-Compute Analytics Data

- 👉 Run background jobs to aggregate data for dashboards.

Example: Count posts per user overnight → store in separate collection → dashboard reads pre-computed data.

Materialized views batch processing Spark/Flink

Topic 20: Consistent Hashing

💡 Problem 57: How Consistent Hashing Works

Question: INTERVIEWER: You have 10 cache servers. When you add an 11th server, all cache keys get invalidated. Why?

Solution:

Simple modulo hashing causes massive cache invalidation when nodes change — consistent hashing solves this.

Here's how consistent hashing works 🤓

1 Problem with Simple Modulo Hashing

👉 Hash(key) % N maps keys to servers, but changing N remaps all keys.

Example: $\text{Hash("user123") \% 10 = server 3}$. Add 11th server $\rightarrow \text{Hash("user123") \% 11 = server 8} \rightarrow \text{cache miss}$.

Massive reshuffling cache invalidation storm

2 Consistent Hashing Ring

👉 Map servers and keys onto circular hash space (0 to 2^{32}).

Example: Like a clock — servers at positions 3, 7, 11 o'clock \rightarrow keys map to next clockwise server.

Hash ring minimal disruption

3 Key Assignment to Servers

👉 Hash key \rightarrow place on ring \rightarrow clockwise to next server.

Example: $\text{Hash("user123") = 5} \rightarrow \text{server at position 7 handles it.}$

Deterministic assignment consistent across clients

4 Adding Server Only Affects Neighbors

👉 New server takes some keys from next clockwise server only.

Example: Add server at position 5 → takes keys between 3 and 5 from server at 7 → other 9 servers unaffected.

Minimal data movement ~1/N keys move

5 Removing Server Redistributes to Next Server

👉 When server fails, its keys move to next clockwise server.

Example: Server at 7 fails → its keys go to server at 11.

Automatic failover ~1/N keys move

6 Virtual Nodes for Even Distribution

👉 Each physical server gets multiple positions on ring for balanced load.

Example: 10 servers × 150 virtual nodes each → 1500 points on ring → more even distribution.

Virtual nodes avoid hotspots DynamoDB uses 128 virtual nodes

💡 Problem 58: Implementing Consistent Hashing

Question: INTERVIEWER: How would you implement consistent hashing for a distributed cache?

Solution:

Implementation requires hash function, ring data structure, and virtual nodes for balance.

Here's how to implement consistent hashing 🤓

1 Choose Hash Function

👉 Use uniform distribution hash like MD5, SHA-1, or MurmurHash.

Example: `MurmurHash("cache-1-vnode-5")` → 32-bit integer → position on ring.

Uniform distribution fast computation

2 Represent Ring with Sorted Data Structure

- 👉 Store virtual node positions in sorted list or tree.

Example: TreeMap in Java, bisect in Python → $O(\log N)$ lookup for next server.

Binary search efficient lookups

3 Create Virtual Nodes for Each Server

- 👉 Hash (server + vnodeNumber) to create multiple ring positions per server.

Example: For cache-1 with 150 vnodes → Hash("cache-1-vnode-0") through Hash("cache-1-vnode-149").

Load balancing 100-200 vnodes per server typical

4 Lookup Function: Find Next Server

- 👉 Hash key → binary search for next virtual node on ring → return physical server.

Example: Hash("user123") = 1234567 → binary search ring → find first vnode ≥ 1234567 → return server.

$O(\log N)$ complexity ceiling search

5 Handle Ring Wrap-Around

- 👉 If key hash > largest ring position, wrap to beginning.

Example: Hash("user456") = max+100 → wrap to first vnode on ring.

Circular ring modulo arithmetic

6 Add/Remove Servers Dynamically

- 👉 Insert/delete virtual nodes from ring → only affected keys need remapping.

Example: Remove cache-5 → delete its 150 vnodes from ring → clients automatically use next server.

Membership changes minimal disruption



CORS (Cross-Origin Resource Sharing)

💡 Problem 59: Consistent Hashing Use Cases

Question: INTERVIEWER: When would you use consistent hashing in your system?

Solution:

Consistent hashing excels when you need dynamic cluster membership with minimal disruption.

Here are use cases for consistent hashing 🤝

1 Distributed Caching

👉 Cache servers can scale up/down without invalidating entire cache.

Example: Memcached cluster with 20 servers → add 5 more → only 20% of cache invalidated instead of 100%.

Redis Cluster Memcached CDN caching

2 Load Balancing Across Servers

👉 Distribute requests consistently to same server for session affinity.

Example: User sessions stick to same backend server → consistent hashing on session_id.

Stateful services session persistence

3 Distributed Databases

👉 Partition data across nodes with minimal rebalancing when nodes change.

Example: Cassandra, DynamoDB, Riak use consistent hashing for data distribution.

Data sharding node membership changes

4 CDN Request Routing

- 👉 Route requests to same edge server for cache locality.

Example: Video chunks for movie_id=123 always go to same CDN edge → better cache hit rate.

Akamai Cloudflare edge caching

5 Service Discovery

- 👉 Route service requests consistently to same instance for caching/state.

Example: API requests for user_id always routed to same service instance.

Microservices consul etcd

6 When NOT to Use Consistent Hashing

- 👉 Avoid if you need range queries or perfect load balance with few nodes.

Example: Range query "users WHERE age > 30" doesn't work with consistent hashing.

Range queries small clusters prefer range partitioning

Topic 21: Caching Strategies (Cache Invalidation)

💡 Problem 60: Cache Invalidation Strategies

Question: INTERVIEWER: Your cached data is getting stale. How do you keep cache in sync with database?

Solution:

Cache invalidation is hard — choose strategy based on consistency requirements and access patterns.

Here are cache invalidation strategies 🤔

1 Time-Based Expiration (TTL)

👉 Set expiration time on cached entries — simple but may serve stale data.

Example: Cache user profile for 5 minutes → stale for up to 5 min after update.

TTL Redis EXPIRE simple but eventual consistency

2 Write-Through Cache

👉 Write to cache and database simultaneously — cache always fresh.

Example: Update user profile → write to database AND cache → reads always see latest data.

Synchronous writes slower writes strong consistency

3 Write-Behind (Write-Back) Cache

👉 Write to cache immediately, asynchronously write to database later.

Example: Update user profile → write to cache → background job writes to database.

Fast writes risk of data loss eventual consistency

4 Cache Invalidation on Write

👉 Delete cache entry when database is updated — next read repopulates cache.

Example: Update user profile → delete from cache → next read fetches from database and caches.

Cache-aside pattern read-after-write fetches from database

5 Event-Based Invalidation

👉 Publish event when data changes — cache subscribers invalidate their entries.

Example: User service updates profile → publishes event → cache service receives event → deletes cache entry.

Pub/sub Kafka RabbitMQ decoupled architecture

6 Version-Based Invalidation

👉 Include version number in cache key — new version creates new cache entry.

Example: Cache key "user:123:v5" → update user → cache key becomes "user:123:v6" → old cache ignored.

Versioned keys no explicit invalidation cache bloat risk



WebSockets

💡 Problem 61: Cache Warming Strategies

Question: INTERVIEWER: After deploying new cache servers, cache hit rate is 0% and database is overloaded. How do you fix this?

Solution:

Cold cache causes thundering herd on database — pre-populate cache before taking traffic.

Here's how to warm up cache ↗

1 Pre-Populate Critical Data

👉 Before taking traffic, load frequently accessed data into cache.

Example: Load top 1000 products before launching sale → cache ready for traffic spike.

Batch pre-loading startup script

2 Gradual Traffic Ramp-Up

- 👉 Send small percentage of traffic to new cache server, gradually increase.

Example: 1% traffic for 10 min → 10% for 10 min → 100% → cache warms gradually.

Traffic shaping canary deployment

3 Copy from Existing Cache

- 👉 If adding cache servers, copy hot keys from existing servers.

Example: Redis SCAN existing server → copy top 10K hot keys to new server.

Cache replication hot key migration

4 Background Refresh for Predictable Access

- 👉 Periodically refresh cache entries before they expire.

Example: Refresh homepage data every 4 minutes (TTL 5 min) → always warm.

Proactive refresh scheduled jobs

5 Probabilistic Early Expiration

- 👉 Refresh cache entries probabilistically before expiration to smooth load.

Example: TTL 60s → when TTL < 10s, 20% chance to refresh → avoids synchronized expiration.

Stampede prevention jittered expiration

6 Read-Through with Lazy Loading

- 👉 Accept initial cache misses and let organic traffic warm cache.

Example: First request hits database → populates cache → subsequent requests hit cache.

Lazy loading simpler but initial spike

💡 Problem 62: Multi-Layer Caching

Question: INTERVIEWER: How would you design caching for a high-traffic website with multiple layers?

Solution:

Multi-layer caching reduces latency and load at each tier — place caches close to access points.

Here's how to design multi-layer caching 

1 Browser/Client Cache

👉 Cache static assets and API responses in browser.

Example: Cache CSS/JS with Cache-Control: max-age=31536000 → served from browser cache.

HTTP caching ETags localStorage

2 CDN Edge Cache

👉 Cache content at CDN edge locations close to users.

Example: Product images cached at CDN edge → 10ms latency instead of 100ms.

Cloudflare Akamai CloudFront

3 Application-Level Cache (In-Memory)

👉 Cache frequently accessed data in application memory.

Example: Cache user session in application RAM → no network call needed.

Caffeine Guava Cache local cache

4 Distributed Cache (Redis/Memcached)

👉 Shared cache across multiple application servers.

Example: User profile cached in Redis → all app servers share same cache.

Redis Memcached shared state

5 Database Query Cache

- 👉 Database caches query results internally.

Example: PostgreSQL caches frequently executed queries in shared buffers.

PostgreSQL shared_buffers MySQL query cache

6 Cache Hierarchy Strategy

- 👉 Check caches from closest to farthest — L1, L2, L3 → database.

Example: Check local cache → check Redis → check database → populate caches on way back.

Cascade caching return-and-populate pattern

7 Cache Consistency Across Layers

- 👉 Invalidate all cache layers when data changes to avoid stale data.

Example: Update user profile → invalidate local cache, Redis, CDN.

Multi-layer invalidation cache coherence

Topic 22: Cache-Aside Pattern



💡 Problem 63: Implementing Cache-Aside Pattern

Question: INTERVIEWER: How do you implement caching in your application to reduce database load?

Solution:

Cache-aside (lazy loading) is the most common pattern — application manages cache explicitly.

Here's how cache-aside pattern works 🤝

1 Read Path: Check Cache First

👉 Application checks cache before querying database.

Example: Fetch user → check Redis → if found, return → if miss, query database.

Cache lookup Redis GET

2 Cache Miss: Fetch from Database

👉 On cache miss, query database for data.

*Example: User not in cache → query 'SELECT * FROM users WHERE id = 123'.*

Database query fallback to source of truth

3 Populate Cache After Fetch

👉 Store database result in cache for future requests.

Example: Fetched user from database → store in Redis with 'SET user:123 {data} EX 3600'.

Cache population set TTL

4 Write Path: Update Database First

- 👉 On write, update database then invalidate cache.

Example: Update user → write to database → delete cache key 'DEL user:123'.

Write-through invalidation next read repopulates cache

5 Handle Race Conditions

- 👉 Multiple threads may fetch same data simultaneously on cache miss.

Example: 100 requests for user:123 → all miss cache → all query database → use lock or accept duplicate queries.

Thundering herd use setnx lock

6 Cache-Aside Benefits and Trade-offs

- 👉 Simple to implement, cache only what's accessed, but cache misses impact latency.

Example: Cache hit = 1ms, cache miss = 100ms database query.

Lazy loading eventual consistency read-heavy workloads

💡 Problem 64: Cache-Aside vs Write-Through vs Write-Behind

Question: INTERVIEWER: What's the difference between cache-aside, write-through, and write-behind caching?

Solution:

Different patterns trade off consistency, latency, and complexity.

Here are the caching pattern trade-offs 🤞

1 Cache-Aside (Lazy Loading)

- 👉 Application explicitly manages cache — read from cache, on miss fetch from database and populate.

Example: Most common pattern — application checks Redis, then database.

Application-managed eventual consistency simple

2 Write-Through Cache

- 👉 Writes go through cache to database — cache always consistent with database.

Example: Update user → cache writes to database synchronously → returns success.

Synchronous writes slower writes strong consistency

3 Write-Behind (Write-Back) Cache

- 👉 Writes go to cache immediately, asynchronously written to database.

Example: Update user → write to cache → return → background job flushes to database.

Fast writes data loss risk eventual consistency

4 Read-Through Cache

- 👉 Cache sits in front of database — cache automatically fetches from database on miss.

Example: Application queries cache → cache fetches from database if miss → returns data.

Cache-managed transparent to application

5 When to Use Each Pattern

- 👉 Cache-aside for most use cases, write-through for critical consistency, write-behind for high write throughput.

Example: User profiles → cache-aside. Financial transactions → write-through. Analytics events → write-behind.

Choose based on requirements



Topic 23: Distributed Caching (Redis)

💡 Problem 65: Redis vs Memcached

INTERVIEWER: When would you use Redis vs Memcached?

Question:

Solution:

Both are distributed caches, but Redis offers richer data structures and persistence.

Here's when to use Redis vs Memcached 🤗

1 Use Memcached for Simple Key-Value Caching

👉 Memcached is simpler, faster for basic get/set operations.

Example: Cache API responses, HTML fragments → Memcached sufficient.

Simple caching multi-threaded less memory overhead

2 Use Redis for Rich Data Structures

👉 Redis supports lists, sets, sorted sets, hashes beyond simple strings.

Example: Leaderboard with sorted set, shopping cart with hash, job queue with list.

Data structures ZADD HSET LPUSH

3 Use Redis for Persistence

👉 Redis can persist cache to disk for recovery after restart.

Example: Session data → Redis with AOF persistence → survives server restart.

RDB snapshots AOF durability

4 Use Redis for Pub/Sub

- 👉 Redis supports publish/subscribe messaging pattern.

Example: Real-time notifications → publish to Redis channel → subscribers receive instantly.

PUBLISH SUBSCRIBE message broker

5 Use Redis for Transactions and Lua Scripts

- 👉 Redis supports atomic operations with MULTI/EXEC and Lua scripting.

Example: Decrement inventory atomically with Lua script → prevent race conditions.

MULTI/EXEC EVAL atomicity

6 Use Memcached for Multi-Threaded Performance

- 👉 Memcached uses multi-threading, better CPU utilization on multi-core systems.

Example: 16-core server → Memcached uses all cores, Redis single-threaded (one core per instance).

Multi-threading horizontal scaling



💡 Problem 66: Redis Clustering and High Availability

INTERVIEWER: Your single Redis instance is a single point of failure. How do you make Redis highly available?

Question:

Solution:

Redis supports replication, Sentinel for failover, and clustering for horizontal scaling.

Here's how to make Redis highly available 🤝

1 Redis Replication (Master-Slave)

👉 Replicate data from master to one or more replicas for read scaling and failover.

Example: 1 master for writes, 3 replicas for reads → reads scale horizontally.

REPLICAOF asynchronous replication

2 Redis Sentinel for Automatic Failover

👉 Sentinel monitors master, automatically promotes replica if master fails.

Example: Master crashes → Sentinel detects → promotes replica to master → updates clients.

Sentinel quorum automatic failover monitoring

3 Redis Cluster for Horizontal Scaling

👉 Shard data across multiple master nodes using hash slots.

Example: 3 master nodes → 16384 hash slots divided evenly → keys distributed across masters.

Redis Cluster 16384 hash slots no single point of failure

4 Read Replicas for Read Scaling

👉 Route read queries to replicas to reduce load on master.

Example: 90% read traffic → replicas handle reads → master handles writes only.

Read/write splitting replica lag

5 Persistence for Data Durability

👉 Enable RDB snapshots or AOF to recover data after crash.

Example: Server crashes → restart → load RDB snapshot → data restored.

RDB AOF fsync policy

6 Connection Pooling and Client-Side Failover

👉 Use connection pool with multiple Redis endpoints for client-side failover.

Example: Client connects to master → master fails → client switches to replica.

Jedis redis-py client libraries

💡 Problem 67: Redis Persistence (RDB vs AOF)

INTERVIEWER: How does Redis persist data to disk, and what are the trade-offs?

Question:

Solution:

Redis offers two persistence mechanisms with different durability and performance trade-offs.

Here's how Redis persistence works 🤓

1 RDB (Redis Database Snapshot)

👉 Periodically save full snapshot of data to disk.

Example: Save snapshot every 5 minutes → server crashes → restore from last snapshot → lose up to 5 min of data.

BGSAVE point-in-time snapshots compact files

2 AOF (Append-Only File)

👉 Log every write operation to disk — can replay to recover exact state.

Example: Every write appended to AOF → server crashes → replay AOF → restore all data.

Append-only log better durability larger files

3 RDB Performance

👉 RDB is faster and generates smaller files, but higher data loss risk.

Example: 1GB RDB snapshot takes 10 seconds to load → faster recovery than AOF.

Faster restarts background fork data loss risk

4 AOF Durability

- 👉 AOF provides better durability with configurable fsync policy.

Example: fsync every second → lose max 1 second of data on crash.

appendfsync everysec appendfsync always appendfsync no

5 Use Both RDB and AOF

- 👉 Enable both for best durability and fast recovery.

Example: RDB every 5 min + AOF fsync every second → recover from RDB → replay recent AOF for latest state.

Hybrid approach Redis loads RDB then AOF

6 AOF Rewrite for Compaction

- 👉 AOF grows large over time — rewrite to compact.

Example: 1000 SET operations on same key → rewrite as single SET → smaller file.

BGREWRITEAOF automatic rewrite

💡 Problem 68: Choosing Redis Persistence Strategy

INTERVIEWER: Your Redis has 10GB of session data — how do you ensure it survives a restart?

Question:

Solution:

Redis persistence balances durability, performance, and recovery time.

1 RDB (Point-in-time Snapshots)

- 👉 Periodic snapshots saved to disk.

Example: Save snapshot every 5 minutes → fast restart, but lose up to 5 minutes of data on crash.

RDB snapshots backup

2 AOF (Append-Only File)

- 👉 Log every write operation to file.

Example: Every SET, INCR command logged → replay on restart → minimal data loss but slower.

AOF durability

3 AOF fsync Policies

- 👉 Control how often AOF is flushed to disk.

Example: fsync always (slow, durable), everysec (balanced), or no (fast, risky).

fsync policies trade-offs

4 RDB + AOF Hybrid

- 👉 Use both for best durability and fast restarts.

Example: RDB snapshot every hour + AOF for recent writes → fast recovery with minimal loss.

Hybrid persistence best of both

5 AOF Rewrite

- 👉 Compact AOF file by removing redundant operations.

Example: SET key val1, SET key val2, SET key val3 → rewrites to single SET key val3.

AOF compaction disk space

6 No Persistence (Cache-Only)

- 👉 For ephemeral data that can be regenerated.

Example: Temporary rate-limit counters → no persistence needed → faster.

Cache-only ephemeral data

7 Backup Strategy

- 👉 Regular backups of RDB files to external storage.

Example: Cron job copies RDB to S3 every 6 hours → disaster recovery.

Backup S3 disaster recovery



Problem 68: Choosing Redis Persistence Strategy

INTERVIEWER: Your Redis has 10GB of session data — how do you ensure it survives a restart?

Question:

Solution:

Redis persistence balances durability, performance, and recovery time.

1 RDB (Point-in-time Snapshots)

- 👉 Periodic snapshots saved to disk.

Example: Save snapshot every 5 minutes → fast restart, but lose up to 5 minutes of data on crash.

RDB snapshots backup

2 AOF (Append-Only File)

- 👉 Log every write operation to file.

Example: Every SET, INCR command logged → replay on restart → minimal data loss but slower.

AOF durability

3 AOF fsync Policies

- 👉 Control how often AOF is flushed to disk.

Example: fsync always (slow, durable), everysec (balanced), or no (fast, risky).

[fsync policies](#) [trade-offs](#)

4 RDB + AOF Hybrid

- 👉 Use both for best durability and fast restarts.

Example: RDB snapshot every hour + AOF for recent writes → fast recovery with minimal loss.

[Hybrid persistence](#) [best of both](#)

5 AOF Rewrite

- 👉 Compact AOF file by removing redundant operations.

Example: SET key val1, SET key val2, SET key val3 → rewrites to single SET key val3.

[AOF compaction](#) [disk space](#)

6 No Persistence (Cache-Only)

- 👉 For ephemeral data that can be regenerated.

Example: Temporary rate-limit counters → no persistence needed → faster.

[Cache-only](#) [ephemeral data](#)

7 Backup Strategy

- 👉 Regular backups of RDB files to external storage.

Example: Cron job copies RDB to S3 every 6 hours → disaster recovery.

[Backup](#) [S3](#) [disaster recovery](#)

Topic 24: Redis Data Structures

REST API Best Practices

💡 Problem 68: Using Redis Strings and Hashes

INTERVIEWER: How would you store user session data in Redis?

Question:

Solution:

Redis strings and hashes are perfect for key-value storage with expiration.

Here's how to use Redis strings and hashes ↗

1 Simple String Storage

👉 Store serialized data (JSON, MessagePack) as string.

Example: `SET session:abc123 "{user_id:5,name:Alice}" EX 3600` → expires in 1 hour.

SET GET EXPIRE serialization

2 Hash for Structured Data

👉 Store individual fields in hash for partial access.

Example: `HSET user:123 name Alice age 30 email alice@example.com` → access fields individually.

HSET HGET HGETALL field-level access

3 Atomic Increment/Decrement

👉 Use INCR/DECR for counters with atomic operations.

Example: `INCR page:views:123` → increment view counter atomically → prevent race conditions.

INCR DECR INCRBY atomic operations

4 Set Expiration on Keys

👉 Auto-delete keys after TTL for session management and cache invalidation.

Example: Session expires after 30 minutes → `EXPIRE session:abc123 1800`.

EXPIRE TTL automatic cleanup

5 Conditional Operations with NX/XX

👉 SETNX sets only if key doesn't exist → useful for distributed locks.

Example: `SET lock:resource:123 1 NX EX 10` → acquire lock if not held.

SETNX distributed lock Redlock

6 Batch Operations with MGET/MSET

👉 Fetch/set multiple keys in single round trip.

Example: `MGET user:1 user:2 user:3` → 1 network call instead of 3.

MGET MSET pipelining performance

💡 Problem 69: Using Redis Lists and Sets

INTERVIEWER: How would you implement a job queue and a tagging system in Redis?

Question:

Solution:

Redis lists are perfect for queues, sets for unique collections and membership tests.

Here's how to use Redis lists and sets 

1 List as FIFO Queue

 LPUSH to enqueue, RPOP to dequeue — simple job queue.

Example: `LPUSH jobs:pending {job_id:1,task:send_email}` → `RPOP jobs:pending` → process job.

LPUSH RPOP FIFO queue

2 Blocking Pop for Worker Pools

 BRPOP blocks until item available — efficient worker pattern.

Example: Worker calls `BRPOP jobs:pending 30` → blocks up to 30 seconds → processes job when available.

BRPOP blocking operation worker pool

3 List as Activity Feed

 Store recent activity, trim to keep limited history.

Example: `LPUSH feed:user:123 {action:like,post:456}` → `LTRIM feed:user:123 0 99` → keep last 100 items.

LPUSH LTRIM LRANGE recent activity

4 Set for Unique Collections

 Store unique items, automatic deduplication.

Example: `SADD tags:post:456 python backend redis` → no duplicates.

SADD SMEMBERS unique values

5 Set Operations (Union, Intersection, Difference)

 Combine sets to find common or unique elements.

Example: `SINTER tags:post:1 tags:post:2` → find common tags between posts.

SINTER SUNION SDIFF set algebra

6 Set Membership Testing

👉 Check if element exists in set — O(1) operation.

Example: `SISMEMBER followers:user:123 user:456` → check if user:456 follows user:123.

SISMEMBER fast membership test

💡 Problem 70: Using Redis Sorted Sets

INTERVIEWER: How would you implement a real-time leaderboard that ranks users by score?

Question:

Solution:

Redis sorted sets maintain sorted order automatically — perfect for leaderboards and priority queues.

Here's how to use Redis sorted sets 🌟

1 Add Elements with Scores

👉 ZADD stores members with scores, automatically sorted.

Example: `ZADD leaderboard 1500 user:123 2000 user:456` → users ranked by score.

ZADD score-based sorting

2 Get Top N Elements

👉 ZREVRANGE returns highest-scoring elements.

Example: `ZREVRANGE leaderboard 0 9 WITHSCORES` → top 10 users with scores.

ZREVRANGE leaderboard queries

3 Get User Rank

👉 ZREVRANK returns user's position in leaderboard.

Example: 'ZREVRANK leaderboard user:123' → returns rank 42 → user is 42nd place.

ZREVRANK position lookup O(log N)

4 Increment Score Atomically

👉 ZINCRBY atomically increases user's score.

Example: User earns 50 points → 'ZINCRBY leaderboard 50 user:123' → atomic update.

ZINCRBY atomic increment race-free

5 Range Queries by Score

👉 ZRANGEBYSCORE returns elements within score range.

Example: 'ZRANGEBYSCORE leaderboard 1000 2000' → users with scores between 1000 and 2000.

ZRANGEBYSCORE score filtering

6 Time-Based Leaderboards with Timestamps

👉 Use timestamp as score for time-ordered data.

Example: 'ZADD trending:posts <unix_timestamp> post:456' → get recent posts with ZREVRANGEBYSCORE.

Timestamp as score time-series data

Topic 25: Cache Stampede Problem

Problem 71: Understanding Cache Stampede

INTERVIEWER: A popular cache entry expires, and suddenly 10,000 requests hit your database at once. What happened?

Question:

Solution:

Cache stampede (thundering herd) happens when many requests try to regenerate expired cache simultaneously.

Here's how cache stampede happens 

1 What is Cache Stampede?

👉 Multiple requests simultaneously detect cache miss and query database.

Example: Popular product page cached → cache expires → 10K concurrent requests → all hit database → database overload.

Thundering herd synchronized cache miss

2 Why It's Problematic

👉 Database handles 1K req/sec normally → suddenly receives 10K concurrent queries → crashes or slows down.

Example: Homepage cache expires → millions of users hit database → database CPU at 100%.

Database overload cascading failure

3 When Does It Occur?

👉 Popular cached items with synchronized expiration times.

Example: All cache entries created at deploy time with TTL 3600 → all expire at same time → stampede.

Synchronized expiration hot keys

4 Effect on User Experience

- 👉 Users experience slow responses during stampede.

Example: Page loads in 50ms from cache → during stampede takes 5 seconds from database.

Latency spike timeout errors

5 Risk of Cascading Failure

- 👉 Database overload causes timeouts → retries increase load → more timeouts → death spiral.

Example: Database slow → app retries → more queries → database crashes → entire system down.

Cascading failure retry storm

💡 Problem 3: Where to Add Caching for Slow Homepage

INTERVIEWER: Your homepage takes 2 seconds to load — where exactly do you add caching?

Question:

Solution:

Strategic caching placement at multiple layers dramatically reduces page load time.

1 Browser Caching

- 👉 Cache static assets (CSS, JS, images) in user's browser.

Example: Set Cache-Control: max-age=31536000 for images → users download once, use for a year.

Browser cache headers CDN integration

2 CDN Caching

- 👉 Serve static content from edge servers close to users.

Example: Homepage HTML cached at CDN → 50ms latency instead of 500ms from origin server.

CloudFront Cloudflare Akamai

3 Application-Level Caching

👉 Cache database query results in Redis/Memcached.

Example: Product listings, user profiles cached → avoid 10 DB queries on every page load.

Redis Memcached in-memory cache

4 Database Query Caching

👉 Cache expensive aggregation queries at database level.

Example: `SELECT COUNT(*)` from million-row table → cache result for 5 minutes.

MySQL query cache materialized views

5 Full Page Caching

👉 Cache entire rendered HTML for anonymous users.

Example: Homepage identical for all logged-out users → generate once, serve millions.

Varnish Nginx proxy cache static site generation

6 Fragment/Partial Caching

👉 Cache reusable page components separately.

Example: Product recommendations widget cached independently from rest of page.

Edge Side Includes React Server Components

7 Cache Warming

👉 Pre-populate cache before traffic surge.

Example: Before Black Friday, pre-cache top 1000 products → no cold start delays.

Scheduled jobs background workers

Problem 4: Reducing API Latency from 450ms to <100ms

INTERVIEWER: Your search API must respond in <100ms — but DB takes 450ms. What's your move?

Question:

Solution:

Sub-100ms API responses require eliminating database calls from the critical path.

1 Add Search Engine Layer

👉 Use Elasticsearch/Algolia for fast full-text search instead of DB.

Example: Elasticsearch returns search results in 10-30ms vs 450ms from PostgreSQL.

Elasticsearch Algolia Typesense

2 Aggressive Caching

👉 Cache search results for popular queries.

Example: "iPhone 15" searched 10K times/hour → cache results for 60 seconds.

Redis cache CDN caching stale-while-revalidate

3 Read Replicas for Search

👉 Route search queries to optimized read replicas.

Example: Dedicated read replica with search-optimized indexes → 200ms → 80ms.

PostgreSQL replicas read-only nodes

4 Database Query Optimization

👉 Add proper indexes and optimize search queries.

Example: Add GIN index for full-text search → 450ms → 150ms (still not enough).

[Index optimization](#) [query planning](#)

5 Denormalization

- 👉 Pre-compute search-friendly data structures.

Example: Denormalized search_index table with all searchable fields → simpler, faster queries.

[Materialized views](#) [denormalized tables](#)

6 Asynchronous Search

- 👉 Return instant results from cache, update in background if stale.

Example: Return cached results in 20ms, trigger background refresh if >5 minutes old.

[Background jobs](#) [eventual consistency](#)

7 Search-Specific Database

- 👉 Use database optimized for search workloads.

Example: Migrate from PostgreSQL to Elasticsearch for product search only.

[Specialized databases](#) [polyglot persistence](#)

💡 Problem 3: Where to Add Caching for Slow Homepage

INTERVIEWER: Your homepage takes 2 seconds to load — where exactly do you add caching?

Question:

Solution:

Strategic caching placement at multiple layers dramatically reduces page load time.

1 Browser Caching

- 👉 Cache static assets (CSS, JS, images) in user's browser.

Example: Set Cache-Control: max-age=31536000 for images → users download once, use for a year.

Browser cache headers CDN integration

2 CDN Caching

- 👉 Serve static content from edge servers close to users.

Example: Homepage HTML cached at CDN → 50ms latency instead of 500ms from origin server.

CloudFront Cloudflare Akamai

3 Application-Level Caching

- 👉 Cache database query results in Redis/Memcached.

Example: Product listings, user profiles cached → avoid 10 DB queries on every page load.

Redis Memcached in-memory cache

4 Database Query Caching

- 👉 Cache expensive aggregation queries at database level.

Example: SELECT COUNT() from million-row table → cache result for 5 minutes.*

MySQL query cache materialized views

5 Full Page Caching

- 👉 Cache entire rendered HTML for anonymous users.

Example: Homepage identical for all logged-out users → generate once, serve millions.

Varnish Nginx proxy cache static site generation

6 Fragment/Partial Caching

- 👉 Cache reusable page components separately.

Example: Product recommendations widget cached independently from rest of page.

Edge Side Includes

React Server Components

7 Cache Warming

- 👉 Pre-populate cache before traffic surge.

Example: Before Black Friday, pre-cache top 1000 products → no cold start delays.

Scheduled jobs

background workers

💡 Problem 4: Reducing API Latency from 450ms to <100ms

INTERVIEWER: Your search API must respond in <100ms — but DB takes 450ms. What's your move?

Question:

Solution:

Sub-100ms API responses require eliminating database calls from the critical path.

1 Add Search Engine Layer

- 👉 Use Elasticsearch/Algolia for fast full-text search instead of DB.

Example: Elasticsearch returns search results in 10-30ms vs 450ms from PostgreSQL.

Elasticsearch

Algolia

Typesense

2 Aggressive Caching

- 👉 Cache search results for popular queries.

Example: "iPhone 15" searched 10K times/hour → cache results for 60 seconds.

Redis cache

CDN caching

stale-while-revalidate

3 Read Replicas for Search

- 👉 Route search queries to optimized read replicas.

Example: Dedicated read replica with search-optimized indexes → 200ms → 80ms.

PostgreSQL replicas read-only nodes

4 Database Query Optimization

- 👉 Add proper indexes and optimize search queries.

Example: Add GIN index for full-text search → 450ms → 150ms (still not enough).

Index optimization query planning

5 Denormalization

- 👉 Pre-compute search-friendly data structures.

Example: Denormalized search_index table with all searchable fields → simpler, faster queries.

Materialized views denormalized tables

6 Asynchronous Search

- 👉 Return instant results from cache, update in background if stale.

Example: Return cached results in 20ms, trigger background refresh if >5 minutes old.

Background jobs eventual consistency

7 Search-Specific Database

- 👉 Use database optimized for search workloads.

Example: Migrate from PostgreSQL to Elasticsearch for product search only.

Specialized databases polyglot persistence

Webhooks

Problem 72: Preventing Cache Stampede

INTERVIEWER: How do you prevent cache stampede from bringing down your database?

Question:

Solution:

Use locking, probabilistic expiration, or early cache refresh to prevent thundering herd.

Here's how to prevent cache stampede 

1 Lock-Based Cache Regeneration

 First request acquires lock, regenerates cache → others wait for lock to be released.

Example: Request 1 gets lock → queries database → populates cache → releases lock → other requests read from cache.

Distributed lock

Redis SETNX

Redlock

2 Probabilistic Early Expiration

 Randomly refresh cache before expiration based on probability.

Example: TTL 60s → when TTL < 10s, 10% chance to refresh → avoids synchronized expiration.

XFetch algorithm

jittered expiration

3 Serve Stale While Revalidating

 Return expired cache immediately, refresh asynchronously in background.

Example: Cache expired → return stale data → background job refreshes cache → next request gets fresh data.

Stale-while-revalidate eventual consistency

4 Background Refresh Before Expiration

- 👉 Proactively refresh cache before it expires.

Example: Cache expires in 60s → background job refreshes at 55s → cache never expires.

Proactive refresh scheduled jobs cron

5 Add Jitter to TTLs

- 👉 Randomize expiration times to avoid synchronized cache misses.

Example: Base TTL 3600s → add random 0-300s → TTL between 3600-3900s → expirations spread out.

TTL jitter desynchronization

6 Use Semaphore/Rate Limiting

- 👉 Limit concurrent database queries during cache miss.

Example: Max 10 concurrent regenerations → others wait or return stale data.

Semaphore rate limiting backpressure

Topic 26: Kafka Architecture

- 💡 Problem 73: Understanding Kafka Topics, Partitions, and Consumer Groups

INTERVIEWER: Explain how Kafka distributes messages across multiple consumers.

Question:

Solution:

Kafka partitions topics for parallelism — consumer groups enable load balancing across consumers. Here's how Kafka topics and partitions work 

1 Topics as Message Categories

 Topic is a logical grouping of related messages.

Example: 'user-events' topic for user actions, 'payment-events' topic for payments.

Logical separation message routing

2 Partitions for Parallelism

 Topic divided into partitions — each partition is ordered log of messages.

Example: 'user-events' topic with 10 partitions → messages distributed across 10 partitions.

Horizontal scaling parallelism ordering per partition

3 Partition Assignment by Key

 Messages with same key go to same partition for ordering guarantee.

Example: All events for user_id=123 go to same partition → ordered per user.

Hash(key) % num_partitions ordering guarantee

4 Consumer Groups for Load Balancing

 Consumers in same group split partitions among themselves.

Example: 10 partitions, 5 consumers → each consumer reads from 2 partitions.

Load balancing horizontal scaling

5 One Partition → One Consumer Per Group

- 👉 Each partition consumed by exactly one consumer in a group at a time.

Example: 10 partitions, 15 consumers → 10 consumers active, 5 idle (can't have more consumers than partitions).

Partition-to-consumer mapping scaling limit

6 Multiple Consumer Groups for Different Use Cases

- 👉 Each consumer group independently consumes all messages.

Example: Group A for analytics, Group B for notifications → both process all messages.

Fan-out pattern independent processing

📚 Idempotency

💡 Problem 74: Kafka Offsets and Message Delivery Guarantees

INTERVIEWER: How does Kafka ensure messages aren't lost or processed twice?

Question:

Solution:

Kafka uses offsets for tracking and offers three delivery semantics.

Here's how Kafka offsets and delivery guarantees work 🌟

1 Offset as Message Position

- 👉 Offset is sequential ID of message within partition.

Example: Partition has messages at offsets 0, 1, 2, 3... → consumer tracks "I've read up to offset 100".

`Sequential position` `per-partition offset`

2 Consumer Commits Offset

- 👉 Consumer periodically saves its current offset to Kafka.

Example: Consumer reads offset 100 → commits offset 100 → crashes → restarts from offset 100.

`Offset commit` `resume from last commit`

3 At-Most-Once Delivery

- 👉 Commit offset before processing message — message may be lost on crash.

Example: Read offset 100 → commit 100 → crash before processing → message lost.

`Commit-then-process` `data loss risk` `lowest latency`

4 At-Least-Once Delivery

- 👉 Process message then commit offset — message may be reprocessed on crash.

Example: Read offset 100 → process → crash before commit → restart → process offset 100 again.

`Process-then-commit` `duplicates possible` `most common`

5 Exactly-Once Semantics

- 👉 Kafka transactions ensure message processed and offset committed atomically.

Example: Process message and commit offset in single transaction → no loss or duplication.

`Kafka transactions` `idempotent producer` `highest consistency`

6 Manual vs Automatic Offset Commit

- 👉 Auto-commit commits offsets periodically, manual commit gives fine-grained control.

Example: Auto-commit every 5 seconds → simple but may duplicate/lose messages. Manual commit after processing → more control.

`enable.auto.commit` `commitSync` `commitAsync`

💡 Problem 75: Kafka Replication and High Availability

INTERVIEWER: How does Kafka ensure messages aren't lost when a broker crashes?

Question:

Solution:

Kafka replicates partitions across brokers for fault tolerance.

Here's how Kafka replication works 🤓

1 Replication Factor

👉 Each partition replicated to N brokers.

Example: Replication factor 3 → partition stored on 3 different brokers.

Fault tolerance data redundancy

2 Leader and Follower Replicas

👉 One replica is leader (handles reads/writes), others are followers (replicate).

Example: Broker 1 is leader for partition 0 → Brokers 2 and 3 are followers.

Leader election read/write routing

3 In-Sync Replicas (ISR)

👉 Replicas caught up with leader are "in-sync" — eligible for leadership.

Example: Leader has offset 1000 → followers at 999 and 1000 are ISR → follower at 900 is not ISR.

Replica lag ISR set

4 Producer Acknowledgments (acks)

👉 Control how many replicas must acknowledge write before success.

Example: `acks=1` → leader ack only (fast, data loss risk). `acks=all` → wait for all ISR (slow, durable).

[Durability vs latency](#) [acks config](#)

5 Leader Election on Failure

- 👉 When leader fails, Kafka elects new leader from ISR.

Example: Leader broker crashes → Kafka controller picks ISR follower → promotes to leader.

[Automatic failover](#) [ZooKeeper/KRaft coordination](#)

6 Min In-Sync Replicas

- 👉 Minimum number of ISRs required for write to succeed.

Example: Replication factor 3, `min.insync.replicas=2` → write succeeds if 2+ replicas available.

[Availability vs durability](#) [min.insync.replicas](#)

Topic 27: RabbitMQ Patterns

Pagination

💡 Problem 76: RabbitMQ Exchanges and Routing

INTERVIEWER: How does RabbitMQ route messages from producers to queues?

Question:

Solution:

RabbitMQ uses exchanges with routing rules to deliver messages to queues.

Here's how RabbitMQ exchanges work 

1 Producer → Exchange → Queue → Consumer

👉 Producers send to exchanges, exchanges route to queues based on routing key.

Example: Producer sends to "orders" exchange → exchange routes to specific queue.

Decoupling flexible routing

2 Direct Exchange (Routing Key Match)

👉 Route message to queue whose binding key exactly matches routing key.

Example: Message with routing key "error" → routed to queue bound with "error".

Exact match unicast

3 Fanout Exchange (Broadcast)

👉 Broadcast message to all bound queues, ignore routing key.

Example: Log message → fanout exchange → delivered to all logging queues.

Broadcast pub/sub pattern

4 Topic Exchange (Pattern Matching)

👉 Route based on wildcard pattern matching.

Example: Routing key "user.signup.email" → matches queue bound to "user..email" or "user.signup.*".*

Wildcard routing * = one word # = zero or more words

5 Headers Exchange (Header Matching)

👉 Route based on message header attributes instead of routing key.

Example: Message with headers {type: email, priority: high} → routed to matching queue.

Attribute-based routing complex matching

6 Default Exchange (Empty String)

- 👉 Automatically routes to queue with name matching routing key.

Example: Send to default exchange with routing key "my-queue" → delivered to queue named "my-queue".

Implicit routing simple use cases

💡 Problem 77: RabbitMQ Message Acknowledgments and Delivery Guarantees

INTERVIEWER: How does RabbitMQ ensure messages aren't lost if a consumer crashes?

Question:

Solution:

RabbitMQ uses acknowledgments and persistence for message durability.

Here's how RabbitMQ delivery guarantees work 👉

1 Manual Acknowledgment (Ack)

- 👉 Consumer explicitly acks message after successful processing.

Example: Receive message → process → send ack → RabbitMQ deletes message.

Manual ack reliable delivery

2 Auto-Acknowledgment

- 👉 RabbitMQ marks message delivered as soon as sent to consumer.

Example: Send message → auto-ack immediately → consumer crashes → message lost.

Auto-ack fire-and-forget data loss risk

3 Negative Acknowledgment (Nack/Reject)

- 👉 Consumer rejects message → RabbitMQ requeues or discards it.

Example: Processing fails → send nack with requeue=true → message goes back to queue.

Error handling retry mechanism

4 Publisher Confirms

- 👉 RabbitMQ confirms to producer that message was received and persisted.

Example: Producer publishes → waits for confirmation → ensures message not lost.

Producer reliability at-least-once guarantee

5 Message Persistence

- 👉 Mark messages and queues as durable → survives broker restart.

Example: Durable queue + persistent message → RabbitMQ writes to disk → survives crash.

Durability disk storage performance cost

6 Prefetch Count for Flow Control

- 👉 Limit unacknowledged messages per consumer to prevent overload.

Example: Prefetch=10 → consumer gets max 10 messages before acking → prevents overwhelming slow consumer.

Flow control backpressure QoS

📚 Search & Indexing

- 💡 Problem 78: RabbitMQ Work Queues and Competing Consumers

INTERVIEWER: You have 10,000 jobs to process. How do you distribute them across multiple workers using RabbitMQ?

Question:



Solution:

Work queues distribute tasks across competing consumers for parallel processing.

Here's how RabbitMQ work queues work 

1 Multiple Consumers on Same Queue

👉 Multiple consumers read from same queue — RabbitMQ round-robin messages.

Example: Queue with 1000 jobs, 10 workers → each worker gets ~100 jobs.

Competing consumers load balancing

2 Round-Robin Dispatching

👉 RabbitMQ distributes messages evenly across consumers.

Example: Message 1 → Consumer A, Message 2 → Consumer B, Message 3 → Consumer C.

Fair distribution default behavior

3 Prefetch for Fair Dispatch

👉 Set prefetch count so workers get messages based on processing speed.

Example: Fast worker finishes and gets new message → slow worker still processing → avoids slow worker bottleneck.

Fair dispatch prefetch count QoS

4 Message Acknowledgment for Reliability

👉 If worker crashes, unacked messages redelivered to other workers.

Example: Worker A crashes while processing → RabbitMQ redelivers to Worker B.

Fault tolerance no message loss

5 Priority Queues

👉 Assign priorities to messages → higher priority processed first.

Example: Urgent email → priority 10, regular email → priority 1 → urgent processed first.

Priority-based processing x-max-priority

6 Dead Letter Exchange for Failed Jobs

👉 Failed jobs routed to dead letter queue after max retries.

Example: Job fails 3 times → moved to DLX → manual investigation or retry later.

Error handling DLX x-dead-letter-exchange

Topic 28: Event-Driven Architecture

💡 Problem 79: Designing Event-Driven Systems

INTERVIEWER: How would you design a system where services communicate via events instead of direct API calls?

Question:

Solution:

Event-driven architecture decouples services — producers emit events, consumers react asynchronously.

Here's how to design event-driven systems 👉

1 Events as State Changes

- 👉 Event represents something that happened in the past.

Example: "OrderCreated", "PaymentProcessed", "UserSignedUp" → immutable facts.

Event naming past tense immutable

2 Event Bus/Broker as Central Backbone

- 👉 Services publish events to central broker, other services subscribe.

Example: Order service publishes OrderCreated → inventory, notification, analytics services subscribe.

Kafka RabbitMQ AWS EventBridge

3 Loose Coupling Between Services

- 👉 Producers don't know about consumers — add new consumers without changing producers.

Example: Add fraud detection service → subscribes to PaymentProcessed → no changes to payment service.

Decoupling independent deployment

4 Choreography vs Orchestration

- 👉 Choreography: services react to events independently. Orchestration: central controller directs workflow.

Example: Choreography → OrderCreated → inventory decrements, email sends, analytics logs. Orchestration → Order orchestrator calls each service.

Distributed coordination saga pattern

5 Event Sourcing for Audit Trail

- 👉 Store all events as source of truth, rebuild state by replaying events.

Example: Account balance = replay all deposit/withdrawal events.

Event sourcing append-only log full audit trail

6 CQRS (Command Query Responsibility Segregation)

- 👉 Separate read and write models — events update read models asynchronously.

Example: Write to event store → events update materialized views for queries.

CQRS read model write model separation

7 Handle Eventual Consistency

- 👉 Events processed asynchronously → temporary inconsistency between services.

Example: Order created → inventory updated 2 seconds later → eventual consistency.

Eventual consistency compensating actions

💡 Problem 80: Event Schema Design and Versioning

INTERVIEWER: Your event schema needs to change. How do you handle event schema evolution without breaking consumers?

Question:

Solution:

Backward and forward compatibility are critical — use versioning and schema registries.

Here's how to handle event schema evolution 🤞

1 Use Schema Registry

- 👉 Centralized registry stores and validates event schemas.

Example: Confluent Schema Registry with Avro → enforces schema compatibility.

Schema registry Avro Protobuf JSON Schema

2 Add Fields Only (Backward Compatibility)

- 👉 Add optional fields → old consumers ignore new fields → no breaking changes.

Example: Add "email_verified" field to UserCreated → old consumers still work.

[Backward compatibility](#) [optional fields](#) [default values](#)

3 Never Remove Fields (Forward Compatibility)

👉 Mark fields as deprecated instead of removing → old events still valid.

Example: Deprecate "phone" field but keep it → new consumers handle its absence.

[Forward compatibility](#) [deprecation](#)

4 Version Events Explicitly

👉 Include version in event name or metadata.

Example: UserCreatedV1, UserCreatedV2 → consumers handle specific versions.

[Explicit versioning](#) [parallel versions](#)

5 Schema Evolution Rules

👉 Follow strict compatibility rules to prevent breaking changes.

Example: Only add fields with defaults, never change field types, never reorder fields.

[Compatibility modes](#) [schema registry enforcement](#)

6 Dual Publishing During Migration

👉 Temporarily publish both old and new event versions during transition.

Example: Publish UserCreatedV1 and UserCreatedV2 → consumers migrate gradually.

[Gradual migration](#) [dual publishing](#)

Logging & Monitoring

Problem 81: Event Ordering and Causality

INTERVIEWER: Two events for the same user arrive out of order. How do you handle this?

Question:

Solution:

Ordering is only guaranteed within partition — handle out-of-order events explicitly.

Here's how to handle event ordering 

1 Partition by Entity ID for Ordering

 All events for same entity go to same partition → ordering guaranteed.

Example: All events for `user_id=123` in same Kafka partition → processed in order.

`Kafka partitioning` `ordering per key`

2 Include Timestamp and Sequence Number

 Add timestamp and sequence number to detect out-of-order events.

Example: Event 1 at time T with sequence 5, Event 2 at time T-1 with sequence 4 → detect out-of-order.

`Timestamp` `sequence number` `causality`

3 Reorder Buffer for Late Events

 Buffer events temporarily to reorder before processing.

Example: Event seq=4 arrives → buffer → Event seq=3 arrives → process in order 3, 4.

`Reorder buffer` `windowing` `complexity`

4 Idempotent Processing

 Design consumers to handle duplicate and out-of-order events safely.

Example: "SetEmail" event is idempotent → processing twice or out-of-order doesn't cause issues.

Idempotency state-based vs delta-based

5 Use Vector Clocks for Causality

- 👉 Track causal relationships between events across services.

Example: Vector clock detects that Event B happened after Event A → preserve causality.

Vector clocks Lamport timestamps distributed systems

6 Accept Out-of-Order for Non-Critical Events

- 👉 Some events don't require strict ordering — trade-off consistency for simplicity.

Example: Page view events → order doesn't matter → process as they arrive.

Eventual consistency acceptable for analytics

💡 Problem 82: Implementing CQRS Pattern

INTERVIEWER: Your app has complex reads but simple writes — how does CQRS help?

Question:

Solution:

CQRS separates read and write models for independent optimization and scaling.

1 Separate Read and Write Models

- 👉 Different data structures for queries vs commands.

Example: Write to normalized SQL; Read from denormalized Elasticsearch.

Model separation optimization

2 Command Side (Write)

- 👉 Focus on business logic validation and consistency.

Example: CreateOrder command → validate inventory → save to SQL → publish event.

Commands write model

3 Query Side (Read)

- 👉 Optimized for fast, complex queries.

Example: Product search → query Elasticsearch with filters, facets, full-text search.

Queries read model

4 Eventual Consistency

- 👉 Read model updated asynchronously after write.

Example: Order created → event published → read model updated 100ms later.

Eventual consistency async updates

5 Event Sourcing Integration

- 👉 CQRS often paired with event sourcing.

Example: Store all OrderCreated, OrderShipped events → rebuild read model from events.

Event sourcing event replay

6 Multiple Read Models

- 👉 Different read models for different use cases.

Example: SQL for reporting, Elasticsearch for search, Redis for real-time dashboard.

Multiple models specialized views

7 When NOT to Use CQRS

- 👉 Adds complexity — only worth it for complex domains.

Example: Simple CRUD app → CQRS overkill → stick with single model.

Complexity trade-offs

Problem 82: Implementing CQRS Pattern

INTERVIEWER: Your app has complex reads but simple writes — how does CQRS help?

Question:

[Redacted]

Solution:

CQRS separates read and write models for independent optimization and scaling.

1 Separate Read and Write Models

👉 Different data structures for queries vs commands.

Example: Write to normalized SQL; Read from denormalized Elasticsearch.

Model separation optimization

2 Command Side (Write)

👉 Focus on business logic validation and consistency.

Example: CreateOrder command → validate inventory → save to SQL → publish event.

Commands write model

3 Query Side (Read)

👉 Optimized for fast, complex queries.

Example: Product search → query Elasticsearch with filters, facets, full-text search.

Queries read model

4 Eventual Consistency

- 👉 Read model updated asynchronously after write.

Example: Order created → event published → read model updated 100ms later.

Eventual consistency async updates

5 Event Sourcing Integration

- 👉 CQRS often paired with event sourcing.

Example: Store all OrderCreated, OrderShipped events → rebuild read model from events.

Event sourcing event replay

6 Multiple Read Models

- 👉 Different read models for different use cases.

Example: SQL for reporting, Elasticsearch for search, Redis for real-time dashboard.

Multiple models specialized views

7 When NOT to Use CQRS

- 👉 Adds complexity — only worth it for complex domains.

Example: Simple CRUD app → CQRS overkill → stick with single model.

Complexity trade-offs

Topic 29: Dead Letter Queues

💡 Problem 82: Implementing Dead Letter Queues

INTERVIEWER: Some messages repeatedly fail processing. How do you prevent them from blocking the queue?

Question:

Solution:

Dead letter queues capture failed messages for investigation and retry.

Here's how to implement dead letter queues 🤝

1 What is a Dead Letter Queue?

👉 Separate queue for messages that failed processing after max retries.

Example: Message fails 3 times → moved to DLQ → manual investigation.

Error handling poison messages retry exhaustion

2 Move Messages After Max Retries

👉 Configure max retry count → exceeded → send to DLQ.

Example: RabbitMQ message rejected 5 times → moved to DLX → sent to DLQ.

x-max-retries retry policy

3 Include Failure Metadata

👉 Add error details, retry count, timestamps to DLQ message.

Example: DLQ message includes original message + error reason + retry history.

Debugging info failure context

4 Monitor DLQ Size

👉 Alert when DLQ grows beyond threshold → indicates systemic issues.

Example: DLQ has 10K messages → alarm → investigate root cause.

Monitoring alerting CloudWatch

5 Replay from DLQ After Fix

👉 Fix bug → replay DLQ messages to reprocess.

Example: Bug fixed → batch replay DLQ messages to main queue → reprocess successfully.

Replay mechanism

manual intervention

6 TTL on DLQ Messages

👉 Auto-expire old DLQ messages to prevent unbounded growth.

Example: DLQ messages older than 30 days → automatically deleted.

TTL

message expiration

cleanup

💡 Problem 83: Handling Poison Messages

INTERVIEWER: A malformed message keeps getting retried and fails every time. How do you handle it?

Question:

Solution:

Poison messages can deadlock consumers — detect and route to DLQ quickly.

Here's how to handle poison messages 👇

1 Detect Poison Messages Early

👉 Validate message format before processing — fail fast on malformed messages.

Example: JSON parsing fails → immediately send to DLQ without retries.

Validation

fail fast

schema validation

2 Set Max Retries to Prevent Infinite Loops

👉 Limit retries to prevent poison message from blocking queue forever.

Example: Max 3 retries → after 3 failures → send to DLQ.

Retry budget bounded retries

3 Use Exponential Backoff for Retries

👉 Increase retry delay exponentially to avoid overwhelming system.

Example: Retry after 1s → 2s → 4s → 8s → max retries → DLQ.

Exponential backoff jitter

4 Separate Queue for Known Issues

👉 Route specific types of failures to dedicated queues for investigation.

Example: Validation errors → validation-errors queue, transient errors → retry queue.

Error classification multiple DLQs

5 Add Circuit Breaker

👉 Stop processing after consecutive failures to prevent cascading failures.

Example: 10 consecutive failures → pause consumer → alert → investigate.

Circuit breaker backpressure

6 Manual Inspection and Fix

👉 Have process for manually inspecting and fixing poison messages.

Example: DevOps reviews DLQ → fixes message format → replays to main queue.

Manual intervention runbooks

Kafka vs RabbitMQ

Topic 30: Kafka vs RabbitMQ

Distributed Tracing

💡 Problem 84: When to Use Kafka vs RabbitMQ

INTERVIEWER: Should you use Kafka or RabbitMQ for your messaging needs?

Question:

Solution:

Kafka excels at high-throughput event streaming, RabbitMQ at flexible routing and task queues.

Here's when to use Kafka vs RabbitMQ ↗

1 Use Kafka for High-Throughput Event Streams

👉 Kafka handles millions of messages/sec with horizontal scaling.

Example: Log aggregation, clickstream analytics, activity tracking.

High throughput append-only log partitioned

2 Use RabbitMQ for Complex Routing

👉 RabbitMQ exchanges support topic, fanout, headers routing.

Example: Route notifications based on user preferences, delivery method, priority.

Flexible routing exchanges routing keys

3 Use Kafka for Event Replay and Reprocessing

- 👉 Kafka stores messages with configurable retention — consumers can rewind.

Example: Deploy new analytics service → replay last 7 days of events.

Message retention replay offset management

4 Use RabbitMQ for Task Queues and Work Distribution

- 👉 RabbitMQ competing consumers pattern ideal for distributing jobs.

Example: Process 10K background jobs across worker pool.

Work queues competing consumers acks

5 Use Kafka for Event Sourcing and CQRS

- 👉 Kafka's append-only log perfect for storing event history.

Example: Store all account transactions as events → rebuild state by replaying.

Event sourcing immutable log audit trail

6 Use RabbitMQ for Request/Reply Patterns

- 👉 RabbitMQ supports RPC pattern with reply queues.

Example: Send request to worker → worker responds to reply queue → synchronous feel.

RPC request/reply correlation ID

7 Use Kafka for Multi-Consumer Scenarios

- 👉 Multiple consumer groups independently consume all messages.

Example: Same event stream consumed by analytics, notifications, audit systems.

Fan-out consumer groups independent processing

💡 Problem 85: Kafka vs RabbitMQ Performance and Scalability

INTERVIEWER: Which is faster and more scalable — Kafka or RabbitMQ?

Question:

Solution:

Kafka optimizes for throughput and horizontal scaling, RabbitMQ for low latency and routing flexibility.

Here's how Kafka and RabbitMQ compare ↗

1 Kafka Throughput

👉 Kafka handles millions of messages/sec with batch writes and zero-copy.

Example: Kafka can sustain 100K-1M msgs/sec per broker.

High throughput batch processing sequential disk I/O

2 RabbitMQ Latency

👉 RabbitMQ has lower per-message latency than Kafka.

Example: RabbitMQ <1ms latency, Kafka 5-10ms due to batching.

Lower latency real-time delivery

3 Kafka Horizontal Scaling

👉 Add partitions and brokers to scale linearly.

Example: 10 partitions → add 10 more → double throughput.

Linear scaling partition-based

4 RabbitMQ Vertical Scaling

👉 RabbitMQ scales primarily by adding nodes to cluster or using federation.

Example: Add more RabbitMQ nodes → shared queues → limited horizontal scaling.

Cluster federation mirrored queues

5 Kafka Message Retention

👉 Kafka stores messages for configured time regardless of consumption.

Example: Retain for 7 days → 10TB storage → high disk usage.

Storage cost retention policy

6 RabbitMQ Memory Usage

👉 RabbitMQ stores messages in memory when possible → faster but memory-constrained.

Example: 1M queued messages → high RAM usage → potential memory pressure.

Memory-based disk fallback flow control

💡 Problem 86: Building Reliable ETL Pipelines

INTERVIEWER: You need to sync 100M user records from MySQL to Elasticsearch nightly — how do you build that?

Question:

Solution:

ETL pipelines extract, transform, and load data reliably at scale.

1 Batch vs Streaming

👉 Batch for scheduled bulk loads, streaming for real-time sync.

Example: Nightly reports → batch; Real-time search → streaming (CDC).

Batch ETL streaming ETL

2 Change Data Capture (CDC)

- 👉 Capture database changes in real-time.

Example: Debezium reads MySQL binlog → publishes changes to Kafka → Elasticsearch consumes.

CDC Debezium binlog

3 Idempotent Processing

- 👉 Handle duplicate records gracefully.

Example: Use primary key for upsert → re-processing same record doesn't break.

Idempotency upsert

4 Checkpointing

- 👉 Track progress to resume from failure.

Example: Process 10M records → checkpoint every 100K → crash at 5.2M → resume from 5.2M.

Checkpointing resumable

5 Data Validation

- 👉 Verify data quality before loading.

Example: Check for null emails, invalid dates → reject bad records → alert team.

Validation data quality

6 Error Handling & Dead Letter Queue

- 👉 Route failed records to DLQ for manual review.

Example: Record transformation fails → sent to DLQ → team fixes manually.

DLQ error handling

7 Monitoring & Alerting

- 👉 Track pipeline health, lag, and errors.

Example: Kafka lag > 1 hour → PagerDuty alert → investigate bottleneck.

Monitoring pipeline lag

💡 Problem 86: Building Reliable ETL Pipelines

INTERVIEWER: You need to sync 100M user records from MySQL to Elasticsearch nightly — how do you build that?

Question:

Solution:

ETL pipelines extract, transform, and load data reliably at scale.

1 Batch vs Streaming

👉 Batch for scheduled bulk loads, streaming for real-time sync.

Example: Nightly reports → batch; Real-time search → streaming (CDC).

Batch ETL streaming ETL

2 Change Data Capture (CDC)

👉 Capture database changes in real-time.

Example: Debezium reads MySQL binlog → publishes changes to Kafka → Elasticsearch consumes.

CDC Debezium binlog

3 Idempotent Processing

👉 Handle duplicate records gracefully.

Example: Use primary key for upsert → re-processing same record doesn't break.

Idempotency upsert

4 Checkpointing

- 👉 Track progress to resume from failure.

Example: Process 10M records → checkpoint every 100K → crash at 5.2M → resume from 5.2M.

Checkpointing resumable

5 Data Validation

- 👉 Verify data quality before loading.

Example: Check for null emails, invalid dates → reject bad records → alert team.

Validation data quality

6 Error Handling & Dead Letter Queue

- 👉 Route failed records to DLQ for manual review.

Example: Record transformation fails → sent to DLQ → team fixes manually.

DLQ error handling

7 Monitoring & Alerting

- 👉 Track pipeline health, lag, and errors.

Example: Kafka lag > 1 hour → PagerDuty alert → investigate bottleneck.

Monitoring pipeline lag

Topic 31: Message Deduplication

Alerting & Incident Management

💡 Problem 86: Detecting Duplicate Messages

INTERVIEWER: Your message queue delivers the same message twice. How do you detect duplicates?

Question:

Solution:

Use message IDs and deduplication windows to identify duplicate messages.

Here's how to detect duplicate messages 🤝

1 Unique Message ID

👉 Assign unique ID to each message — check if already processed.

Example: Generate UUID for message → store in processed set → reject if seen before.

Message ID UUID idempotency key

2 Content-Based Hashing

👉 Hash message content → use hash as dedup key.

Example: Hash(message payload) → check if hash exists in cache → skip if duplicate.

Content hash MD5 SHA-256

3 Sliding Window Deduplication

👉 Track message IDs in time window to detect recent duplicates.

Example: Keep last 5 minutes of message IDs in Redis → check new messages against set.

Time window Redis SET with TTL bounded memory

4 Database Unique Constraint

👉 Use database unique index to enforce deduplication.

Example: INSERT message_id → unique constraint violation → duplicate detected.

Unique index database enforcement

5 Sequence Numbers

- 👉 Producer assigns increasing sequence numbers — detect gaps and duplicates.

Example: Receive seq 1, 2, 3, 2 → detect duplicate seq 2.

Sequence number gap detection

6 Bloom Filter for Probabilistic Dedup

- 👉 Use Bloom filter for memory-efficient duplicate detection.

Example: Check Bloom filter → "definitely new" or "probably seen" → handle accordingly.

Bloom filter false positives memory efficient

💡 Problem 87: Implementing Idempotent Message Processing

INTERVIEWER: How do you ensure processing the same message twice doesn't cause problems?

Question:

Solution:

Design idempotent operations that produce same result regardless of how many times they're executed.

Here's how to implement idempotent processing 👉

1 Idempotency Key from Message ID

- 👉 Use message ID as idempotency key — store in database to prevent re-execution.

Example: Process payment with idempotency key → check if already processed → skip if yes.

Idempotency key dedupe table

2 Set-Based Operations (Naturally Idempotent)

- 👉 Operations like SET are inherently idempotent.

Example: SET user.email = "alice@example.com" → executing twice has same result.

Set operations state-based

3 Conditional Updates with Version/Timestamp

- 👉 Update only if record hasn't changed since read.

Example: UPDATE balance WHERE version=5 → fails if someone else updated → prevents duplicate updates.

Optimistic locking compare-and-swap

4 Database Transactions with Dedup Check

- 👉 Within transaction, check if message processed, then process and mark done.

Example: BEGIN TRANSACTION → check processed → insert record → mark processed → COMMIT.

Atomic dedup transactional processing

5 External System Idempotency

- 👉 Use external system's idempotency features.

Example: Stripe payment API with idempotency key → duplicate requests return same result.

Third-party idempotency API support

6 Avoid Non-Idempotent Operations

- 👉 Replace delta operations with set operations.

Example: Instead of "balance += 10", use "balance = old_balance + 10" with version check.

State-based vs delta-based design for idempotency

Performance Optimization

Problem 88: Handling Exactly-Once Delivery

INTERVIEWER: Can you achieve true exactly-once message delivery in distributed systems?

Question:

Solution:

Exactly-once is hard — combine at-least-once delivery with idempotent processing.

Here's how to achieve exactly-once semantics 

1 At-Least-Once Delivery + Idempotency

 Accept duplicates from message system, handle with idempotent processing.

Example: Kafka at-least-once + idempotency keys → effective exactly-once.

Practical approach **most common**

2 Kafka Exactly-Once Semantics

 Kafka transactions ensure message produced and consumed exactly once.

Example: Kafka producer → transaction → consume from input topic + produce to output topic → commit → all-or-nothing.

Kafka transactions **enable.idempotence=true** **transactional.id**

3 Database Transaction with Message Commit

 Process message and commit offset in same database transaction.

Example: PostgreSQL transaction → process message → insert result → save offset → COMMIT.

Transactional inbox/outbox **atomicity**

4 Two-Phase Commit Across Systems

- 👉 Coordinate commit across message system and database.

Example: 2PC → prepare database + message ack → commit both → distributed transaction.

Complex blocking rarely used in practice

5 Idempotent Producer

- 👉 Kafka producer deduplicates messages on broker side.

Example: Network retry sends duplicate → broker detects via sequence number → ignores duplicate.

Producer idempotence Kafka feature

6 Accept "Effectively Once"

- 👉 In practice, at-least-once + idempotency is "exactly-once enough".

Example: Payment processed at most once due to idempotency key → effectively exactly-once.

Pragmatic approach good enough for most use cases

Topic 32: REST API Best Practices

💡 Problem 89: Designing RESTful API Endpoints

INTERVIEWER: How do you design clean, intuitive REST API endpoints?

Question:

Solution:

RESTful APIs use HTTP methods and resource-based URLs for predictable, intuitive interfaces.

Here's how to design RESTful API endpoints 

1 Use Nouns for Resources, Not Verbs

 Endpoints represent resources — use HTTP methods for actions.

Example: `/users` instead of `/getUsers`, `/users/123` instead of `/getUserById/123`.

Resource-based URLs RESTful convention

2 Use HTTP Methods Correctly

 GET for reading, POST for creating, PUT/PATCH for updating, DELETE for deleting.

Example: GET `/users/123`, POST `/users`, PUT `/users/123`, DELETE `/users/123`.

HTTP verb semantics idempotent methods

3 Use Hierarchical URLs for Relationships

 Nest resources to show relationships.

Example: `/users/123/orders` for user's orders, `/users/123/orders/456` for specific order.

Nested resources relationship clarity

4 Use Plural Nouns for Collections

 Collection endpoints use plural form.

Example: `/users` for user collection, `/orders` for order collection.

Plural convention consistency

5 Use Query Parameters for Filtering and Pagination

 Keep URLs clean, use query params for optional filters.

Example: `/users?status=active&page=2&limit=20` instead of `/users/active/page/2/limit/20`.

Query parameters optional filters

6 Version Your API

👉 Include version in URL or header to allow breaking changes.

Example: '/v1/users' or '/v2/users' with different response formats.

API versioning

backward compatibility



Problem 90: REST API Status Codes and Error Handling

INTERVIEWER: What HTTP status codes should you return for different scenarios?

Question:

Solution:

HTTP status codes communicate what happened — use them correctly for clear API semantics.

Here's how to use HTTP status codes 🤝

1 2xx Success Codes

👉 200 OK for successful GET/PUT, 201 Created for POST, 204 No Content for DELETE.

Example: POST '/users' → 201 with created user in response body.

Success codes

semantic meaning

2 4xx Client Error Codes

👉 400 Bad Request for invalid input, 401 Unauthorized for missing auth, 403 Forbidden for insufficient permissions, 404 Not Found for missing resource.

Example: POST '/users' with invalid email → 400 Bad Request with error details.

Client errors

validation failures

3 5xx Server Error Codes

- 👉 500 Internal Server Error for unexpected errors, 503 Service Unavailable for temporary unavailability.

Example: Database connection fails → 500 with generic error message (don't expose internals).

Server errors don't leak implementation details

4 Return Error Details in Response Body

- 👉 Include error code, message, and field-level details for debugging.

Example: `{"error": "VALIDATION_ERROR", "message": "Email is invalid", "fields": {"email": "Invalid format"}}`.

Error structure debugging info

5 Use 422 Unprocessable Entity for Validation Errors

- 👉 Distinguish between malformed request (400) and validation failure (422).

Example: Valid JSON but email already exists → 422 Unprocessable Entity.

Semantic validation errors GitHub convention

6 Use 429 Too Many Requests for Rate Limiting

- 👉 Return 429 with Retry-After header when rate limit exceeded.

Example: Client makes 1000 req/min → 429 with 'Retry-After: 60'.

Rate limiting backpressure

💡 Problem 91: REST API Pagination, Filtering, and Sorting

INTERVIEWER: Your API returns 10,000 users. How do you handle large result sets?

Question:

Solution:

Pagination, filtering, and sorting make large datasets manageable for clients.

Here's how to handle large result sets ↗

1 Offset-Based Pagination

👉 Use `page` and `limit` or `offset` and `limit` query parameters.

Example: `/users?page=2&limit=20` returns users 21-40.

Simple stateless deep pagination slow

2 Cursor-Based Pagination

👉 Use cursor (encoded ID) to fetch next page — better for real-time data.

Example: `/users?cursor=abc123&limit=20` → response includes `next_cursor`.

Consistent results handles insertions/deletions

3 Return Pagination Metadata

👉 Include total count, page info, and next/prev links in response.

Example: `{"data": [...], "total": 1000, "page": 2, "limit": 20, "next": "/users?page=3&limit=20"}`.

Discoverability client convenience

4 Filtering with Query Parameters

👉 Allow filtering by resource fields.

Example: `/users?status=active&role=admin&created_after=2024-01-01`.

Flexible filtering query parameters

5 Sorting with Query Parameters

👉 Use `sort` or `order_by` parameter with direction.

Example: `/users?sort=created_at:desc` or `/users?order_by=name&direction=asc`.

Sorting ascending/descending

6 Limit Maximum Page Size

👉 Cap `limit` parameter to prevent expensive queries.

Example: Client requests `limit=10000` → cap at 100 → return max 100 results.

Resource protection DoS prevention

Database Indexing

💡 Problem 92: Optimizing API Response Time from 2 Seconds to 200ms

INTERVIEWER: Our API endpoint takes 2 seconds to respond. How would you improve it to 200ms?

Question:

Solution:

API performance optimization requires identifying bottlenecks and applying targeted improvements.

Here's how to optimize API response time 👉

1 Identify the Bottleneck First

👉 Profile the request to find where time is spent — database, external APIs, or computation.

Example: Use APM tools (New Relic, Datadog) or simple logging to measure each operation's duration.

Measure first optimize later avoid premature optimization

2 Database Query Optimization

👉 Add indexes on frequently queried columns, optimize N+1 queries with eager loading, use EXPLAIN to analyze query plans.

Example: Query takes 1.8s → add index on user_id → query takes 50ms.

[Indexing](#) [query optimization](#) [database performance](#)

3 Implement Caching

👉 Cache frequently accessed data in Redis/Memcached to avoid repeated database queries.

Example: Product catalog queried 1000x/sec → cache in Redis with 5-min TTL → 99% of requests served from cache (5ms instead of 500ms).

[Redis caching](#) [TTL strategy](#) [cache-aside pattern](#)

4 Use Database Connection Pooling

👉 Reuse existing database connections instead of creating new ones for each request.

Example: Connection creation takes 100ms → connection pool provides instant connection → save 100ms per request.

[Connection pooling](#) [reduce latency](#) [resource reuse](#)

5 Optimize External API Calls

👉 Run external API calls in parallel, implement timeouts, use circuit breakers, or cache responses.

Example: 3 sequential API calls (500ms each = 1.5s total) → run in parallel → 500ms total.

[Parallel execution](#) [async/await](#) [Promise.all](#)

6 Reduce Payload Size

👉 Return only necessary fields, compress responses with gzip/brotli, paginate large result sets.

Example: 5MB JSON response → remove unused fields → 500KB → enable gzip → 50KB → network transfer time reduced.

[Payload optimization](#) [compression](#) [selective fields](#)

7 Asynchronous Processing

👉 Move heavy operations to background jobs and return response immediately.

Example: Order placement with email notification → accept order (200ms) → send email asynchronously via queue.

Background jobs async processing immediate response

8 Use Read Replicas

👉 Route read-heavy operations to read replicas to reduce load on primary database.

Example: Dashboard query hits primary DB → move to read replica → reduce primary load → faster writes.

Read replicas database scaling load distribution

9 Implement CDN for Static Assets

👉 Serve images, CSS, JS from CDN edge locations closer to users.

Example: Loading 10 images from origin server (200ms each) → CDN edge cache (10ms each) → save 1.9s.

CDN edge caching geographic distribution

👉 Remove unnecessary computations, optimize algorithms, use efficient data structures.

Example: $O(n^2)$ loop processing 10k records → optimize to $O(n)$ → reduce from 1s to 10ms.

Algorithm optimization computational efficiency

1 1 Monitor and Measure Continuously

👉 Set up performance monitoring, track P50/P95/P99 latencies, set up alerts.

Example: Track API response time → set alert for P95 > 300ms → proactively fix degradations.

APM observability performance metrics

Topic 33: API Versioning

💡 Problem 93: API Versioning Strategies

INTERVIEWER: You need to make breaking changes to your API. How do you version it?

Question:

Solution:

API versioning allows backward-incompatible changes without breaking existing clients.

Here are API versioning strategies 🤔

1 URL Path Versioning

👉 Include version in URL path.

Example: `/v1/users` vs `/v2/users` — clear and visible.

Most common

easy to route

URL pollution

2 Header Versioning

👉 Specify version in custom header or Accept header.

Example: `X-API-Version: 2` or `Accept: application/vnd.myapi.v2+json`.

Clean URLs

less discoverable

harder to test

3 Query Parameter Versioning

👉 Pass version as query parameter.

Example: `/users?version=2` — simple but not RESTful.

Easy to implement

non-standard

easily forgotten

4 Maintain Multiple Versions Simultaneously

- 👉 Support old versions for transition period, deprecate gradually.

Example: Support v1 and v2 for 6 months → notify clients → sunset v1.

Backward compatibility maintenance burden

5 Deprecation Headers

- 👉 Warn clients about deprecated endpoints with headers.

Example: 'Deprecation: true', 'Sunset: 2024-12-31', 'Link: <https://api.example.com/v2/users>; rel="successor"'.

RFC 8594 graceful migration

6 Semantic Versioning for Major Changes

- 👉 Only increment major version for breaking changes.

Example: v1 → v2 when response structure changes, v1.1 for new optional fields.

Semantic versioning clear expectations

💡 Problem 94: Managing API Breaking Changes

INTERVIEWER: What constitutes a breaking change in an API?

Question:

Solution:

Breaking changes break existing clients — distinguish from backward-compatible changes.

Here's what counts as a breaking change 👉

1 Breaking Changes Requiring New Version

- 👉 Removing fields, changing field types, changing URL structure, changing required parameters.

Example: Remove 'phone' field from response → breaks clients expecting it.

Major version bump client impact

2 Non-Breaking Changes

- 👉 Adding optional fields, adding new endpoints, adding optional query parameters.

Example: Add new 'email_verified' field → old clients ignore it.

Backward compatible minor version

3 Additive Changes Are Safe

- 👉 Clients should ignore unknown fields — adding fields is safe.

Example: Response adds 'created_at' field → clients using 'name' and 'email' still work.

Robustness principle tolerant readers

4 Deprecate Before Removing

- 👉 Mark field/endpoint as deprecated, give clients time to migrate.

Example: Mark 'phone' deprecated → wait 6 months → remove in v2.

Gradual deprecation migration window

5 Support Old and New Formats Temporarily

- 👉 Accept both formats during transition, return new format.

Example: Accept 'user_id' and 'userId' during migration → eventually only 'userId'.

Dual support transition period



Authentication (OAuth 2.0, JWT)

Topic 34: Authentication (OAuth 2.0, JWT)



N+1 Query Problem

 **Problem 95: Understanding OAuth 2.0 Flow**

INTERVIEWER: Explain how OAuth 2.0 works for third-party authentication.

Question:

Solution:

OAuth 2.0 delegates authentication to authorization server — users grant limited access without sharing passwords.

Here's how OAuth 2.0 works 

1 Authorization Code Flow (Most Secure)

- 👉 User redirected to auth server → grants permissions → auth code returned → exchange code for access token.

Example: "Login with Google" → redirect to Google → user approves → code returned → backend exchanges for token.

Server-side apps

PKCE extension for mobile

2 Access Token for API Requests

- 👉 Client includes access token in Authorization header for API calls.

Example: `Authorization: Bearer eyJhbGc...` → API validates token → grants access.

Bearer token short-lived

3 Refresh Token for Token Renewal

- 👉 Long-lived refresh token exchanges for new access token without re-authentication.

Example: Access token expires in 1 hour → use refresh token to get new access token.

Refresh token long-lived secure storage

4 Scopes for Permission Granularity

- 👉 Limit access to specific resources/actions.

Example: Request `scope=read:user,write:repos` → token only allows reading user info and writing repos.

Least privilege scope-based access

5 Client Credentials Flow for Service-to-Service

- 👉 Backend service authenticates with client ID and secret to get token.

Example: Microservice A calls microservice B → uses client credentials → no user involved.

Machine-to-machine service accounts

6 Implicit Flow (Deprecated)

- 👉 Access token returned directly in redirect URL — less secure, avoid for new apps.

Example: Single-page app → token in URL fragment → token exposure risk.

Legacy use Authorization Code with PKCE instead

💡 Problem 96: JWT (JSON Web Tokens) Structure and Validation

INTERVIEWER: How do you implement stateless authentication with JWTs?

Question:



Solution:

JWTs encode claims in signed tokens — servers validate without database lookups.

Here's how JWTs work 🤝

1 JWT Structure: Header.Payload.Signature

👉 Header has algorithm, payload has claims, signature proves authenticity.

Example: `eyJhbGciOiJIUzI1NiJ9.eyJc2VyX2lkIjoxMjN9.signature` → base64-encoded JSON.

Three parts base64url encoded

2 Claims in Payload

👉 Standard claims: `sub` (subject/user ID), `exp` (expiration), `iat` (issued at), `iss` (issuer).

Example: `{"sub": "user123", "exp": 1735689600, "role": "admin"}`.

JWT claims application-specific data

3 Signature Verification

👉 Server verifies signature using secret key (HMAC) or public key (RSA).

Example: Verify signature with secret → if valid, trust claims → if invalid, reject token.

HMAC-SHA256 RSA signature validation

4 Stateless Authentication

👉 No database lookup needed — all info in token.

Example: JWT contains user_id and role → server reads from token → validates signature → grants access.

Stateless horizontal scaling no session storage

5 Short Expiration Times

- 👉 Set short expiration to limit damage if token stolen.

Example: Access token expires in 15 minutes, refresh token in 7 days.

exp claim security vs usability

6 Can't Revoke JWTs Easily

- 👉 JWTs valid until expiration — use short expiration and refresh tokens for revocation.

Example: User logs out → can't invalidate JWT → short expiration limits exposure.

Revocation challenge blacklist or short TTL

Connection Pooling

💡 Problem 97: JWT Security Best Practices

INTERVIEWER: What are the security risks of using JWTs and how do you mitigate them?

Question:

Solution:

JWTs have specific security considerations — follow best practices to avoid vulnerabilities.

Here are JWT security best practices 🌟

1 Always Verify Signature

- 👉 Never trust JWT without signature verification.

Example: Attacker modifies payload → changes 'role: user' to 'role: admin' → signature invalid → reject.

Signature verification cryptographic integrity

2 Use Strong Signing Algorithm

👉 Use RS256 (RSA) or HS256 (HMAC-SHA256), avoid 'none' algorithm.

Example: Attacker sets 'alg: none' → no signature → server must reject 'none' algorithm.

Algorithm whitelist reject 'none'

3 Validate Claims (exp, iat, iss, aud)

👉 Check expiration, issuer, audience to prevent misuse.

Example: Token expired 1 hour ago → check 'exp' claim → reject expired token.

Claim validation expiration check

4 Store Secrets Securely

👉 Use environment variables or secret management for signing keys.

Example: Store JWT secret in AWS Secrets Manager, not hardcoded in code.

Secret management KMS Vault

5 Use HTTPS Only

👉 JWTs sent over HTTP can be intercepted.

Example: Man-in-the-middle attack intercepts JWT → attacker impersonates user.

TLS/SSL encrypted transport

6 Don't Store Sensitive Data in JWT

👉 Payload is base64-encoded, not encrypted — anyone can read it.

Example: Don't put SSN, credit card in JWT → only store user_id, role.

JWT is signed not encrypted public data only

💡 Problem 98: Session-Based vs Token-Based Authentication

INTERVIEWER: Should you use session-based or token-based authentication?

Question:

Solution:

Sessions store state on server, tokens are stateless — each has trade-offs.

Here's when to use sessions vs tokens 🤔

1 Session-Based Authentication

👉 Server stores session data, client gets session ID cookie.

Example: User logs in → server creates session in Redis → returns session ID cookie → client sends cookie on each request.

Server-side state easy revocation

2 Token-Based Authentication (JWT)

👉 Client stores token, server validates signature — stateless.

Example: User logs in → server issues JWT → client stores in localStorage → sends in Authorization header.

Stateless scales horizontally no session storage

3 Use Sessions for Traditional Web Apps

👉 Server-rendered apps with cookies.

Example: Server-side rendered pages with Express sessions.

HttpOnly cookies CSRF protection needed

4 Use Tokens for APIs and SPAs

- 👉 Mobile apps, single-page apps, microservices.

Example: React app calls API → sends JWT in Authorization header.

CORS-friendly mobile-friendly stateless

5 Session Revocation is Easier

- 👉 Delete session from storage → user logged out immediately.

Example: User logs out → delete Redis session → subsequent requests denied.

Immediate revocation centralized control

6 Tokens Scale Better

- 👉 No shared session storage needed across servers.

Example: 10 API servers → no session synchronization → validate JWT locally.

Horizontal scaling no sticky sessions



Topic 35: Authorization (RBAC)



Batch Processing



💡 Problem 99: Implementing Role-Based Access Control (RBAC)

INTERVIEWER: How do you implement authorization where different users have different permissions?

Question:

Solution:

RBAC assigns roles to users, roles have permissions — check permissions before granting access.

Here's how to implement RBAC ↗

1 Define Roles

👉 Create roles representing job functions.

Example: Roles: 'admin', 'editor', 'viewer' → different permission levels.

Role definition job-based

2 Assign Permissions to Roles

👉 Roles have specific permissions to perform actions.

Example: 'admin' can 'create:users', 'delete:users' → 'editor' can 'edit:posts' → 'viewer' can 'read:posts'.

Permission mapping role-permission relationship

3 Assign Roles to Users

👉 Users get one or more roles.

Example: User Alice has role 'admin' → Bob has role 'editor'.

User-role assignment database relationship

4 Check Permissions in Code

👉 Before action, verify user has required permission.

Example: User requests DELETE '/users/123' → check if user has 'delete:users' permission → allow or deny.

Permission check authorization middleware

5 Store Roles in JWT or Session

- 👉 Include user's roles in authentication token.

Example: JWT payload `{"user_id": 123, "roles": ["admin", "editor"]}`.

Stateless authorization claims-based

6 Database Schema for RBAC

- 👉 Tables: users, roles, permissions, user_roles, role_permissions.

Example: User → user_roles → Role → role_permissions → Permission.

Many-to-many relationships normalized schema

💡 Problem 100: Attribute-Based Access Control (ABAC) and Fine-Grained Authorization

INTERVIEWER: RBAC is too coarse-grained. How do you implement fine-grained authorization?

Question:

Solution:

ABAC uses attributes (user, resource, environment) for dynamic access decisions.

Here's how to implement fine-grained authorization 👉

1 ABAC with Policies

- 👉 Define policies based on user attributes, resource attributes, and context.

Example: User can edit post if `user.id == post.author_id OR user.role == admin`.

Policy-based dynamic rules

2 Resource-Level Permissions

- 👉 Check ownership or relationship to resource.

Example: User can DELETE '/posts/123' only if they own the post.

Ownership check resource-specific

3 Use Authorization Library

- 👉 Libraries like Casbin, OPA (Open Policy Agent) for policy enforcement.

Example: Define policies in Rego (OPA) → query OPA for access decision.

Policy engine centralized authorization

4 Row-Level Security in Database

- 👉 Database enforces access control at row level.

Example: PostgreSQL RLS → 'ALTER TABLE posts ENABLE ROW LEVEL SECURITY' → users only see their own posts.

Database-level security PostgreSQL RLS

5 Context-Based Authorization

- 👉 Consider time, location, device for access decisions.

Example: Admin can delete users only from office IP range.

Environmental attributes contextual policies

6 Combine RBAC and ABAC

- 👉 Use roles for general permissions, attributes for fine-grained control.

Example: User has 'editor' role → can edit posts → but only their own posts (attribute check).

Hybrid approach layered authorization

Topic 36: API Security

Stream Processing

Problem 101: Protecting APIs from Common Attacks

INTERVIEWER: What are the most common API security vulnerabilities and how do you prevent them?

Question:

Solution:

APIs face threats like injection, broken auth, excessive data exposure — follow OWASP API Security Top 10.

Here's how to protect APIs 

1 SQL Injection Prevention

 Use parameterized queries, never concatenate user input into SQL.

Example: `SELECT * FROM users WHERE id = ?` with parameter, not `SELECT * FROM users WHERE id = \${userId}`.

[Prepared statements](#) [ORM](#) [input validation](#)

2 Broken Authentication and Authorization

 Validate JWT signature, check expiration, enforce permissions.

Example: Every API request → verify JWT → check user has required permission.

[Token validation](#) [RBAC](#) [principle of least privilege](#)

3 Excessive Data Exposure

 Return only necessary fields, don't expose internal IDs or sensitive data.

Example: `/users/me` returns only `name`, `email` → not `password_hash`, `ssn`.

[Data minimization](#) [API response filtering](#)

4 Rate Limiting

👉 Limit requests per user/IP to prevent abuse.

Example: Max 100 requests per minute per API key → return 429 if exceeded.

[Rate limiting](#) [DDoS protection](#) [API gateway](#)

5 Input Validation

👉 Validate all inputs against expected format, type, range.

Example: Email field → validate regex → reject if invalid.

[Schema validation](#) [sanitization](#) [joi/yup libraries](#)

6 HTTPS Only

👉 Enforce TLS for all API traffic to prevent eavesdropping.

Example: Redirect HTTP to HTTPS, use HSTS header.

[TLS/SSL](#) [encrypted transport](#) [HSTS](#)

7 CORS Configuration

👉 Configure CORS to allow only trusted origins.

Example: `Access-Control-Allow-Origin: https://trusted-app.com` → not `*`.

[CORS](#) [origin whitelist](#) [preflight requests](#)

💡 Problem 102: API Rate Limiting and Throttling

INTERVIEWER: How do you prevent API abuse and ensure fair usage?

Question:

Solution:

Rate limiting controls request frequency per client — protects infrastructure and ensures fair access.

Here's how to implement API rate limiting 

1 Fixed Window Rate Limiting

👉 Allow N requests per time window.

Example: 100 requests per minute → counter resets every minute at :00.

Simple burst at window boundary

2 Sliding Window Rate Limiting

👉 Track requests in rolling time window.

Example: 100 requests per 60-second window → smooth enforcement.

Accurate more complex Redis sorted set

3 Token Bucket Algorithm

👉 Bucket refills at fixed rate, consume token per request.

Example: Bucket capacity 100, refill 10/sec → allows bursts up to 100, then 10 req/sec.

Burst handling smooth rate

4 Return Rate Limit Headers

👉 Tell clients their rate limit status.

Example: 'X-RateLimit-Limit: 100', 'X-RateLimit-Remaining: 42', 'X-RateLimit-Reset: 1735689600'.

Transparency client awareness

5 Different Limits for Different Tiers

- 👉 Free tier gets 100 req/min, paid tier gets 10K req/min.

Example: Check API key → determine tier → apply corresponding limit.

Tiered pricing API key-based

6 Return 429 with Retry-After

- 👉 Return 429 Too Many Requests with retry delay.

Example: '429 Too Many Requests', 'Retry-After: 60' → client waits 60 seconds.

Backpressure client retry guidance

 **Data Serialization**
 **Problem 103: API Keys and Secret Management**

INTERVIEWER: How do you manage API keys and secrets securely?

Question:

Solution:

API keys authenticate clients — store securely, rotate regularly, limit scope.

Here's how to manage API keys securely 👉

1 Hash API Keys Before Storing

- 👉 Store hashed version in database like passwords.

Example: Generate API key → show user once → store SHA-256 hash → validate by hashing incoming key.

Hash storage prevent database leak exposure

2 Use Secret Management Services

- 👉 Store secrets in dedicated secret managers, not environment variables or code.

Example: AWS Secrets Manager, HashiCorp Vault, Azure Key Vault.

Centralized secrets rotation access control

3 Rotate API Keys Regularly

- 👉 Force rotation or allow users to rotate keys.

Example: API keys expire after 90 days → user must generate new key.

Key rotation limit exposure window

4 Scope API Keys to Specific Permissions

- 👉 Each key has limited permissions, not full access.

Example: Read-only API key for analytics, write key for data ingestion.

Least privilege scoped keys

5 Rate Limit by API Key

- 👉 Track usage per API key to enforce quotas.

Example: API key 'abc123' → 100K requests/day → deny after limit.

Per-key rate limiting quota enforcement

6 Monitor API Key Usage

- 👉 Log and alert on suspicious activity.

Example: API key suddenly used from different country → alert security team.

Anomaly detection audit logs

💡 Problem 104: Managing Secrets Securely

INTERVIEWER: Where do you store database passwords and API keys in production?

Question:

Solution:

Secrets must be encrypted, rotated, and never hardcoded or committed to version control.

1 Never Commit Secrets to Git

👉 Use .env files or secret managers, not hardcoded values.

Example: GitHub secret scanning detects committed AWS keys → auto-revoked → service breaks.

Git secrets security breach

2 Secret Management Services

👉 Use dedicated services for secret storage.

Example: AWS Secrets Manager, HashiCorp Vault, Azure Key Vault → encrypted storage.

Secrets Manager Vault Key Vault

3 Encryption at Rest & In Transit

👉 Secrets encrypted in storage and during transmission.

Example: Secrets Manager encrypts with KMS → accessed via TLS → end-to-end encryption.

Encryption KMS

4 Secret Rotation

👉 Automatically rotate secrets periodically.

Example: Database password rotated every 90 days → Secrets Manager updates app automatically.

Rotation automated updates

5 Least Privilege Access

- 👉 Only authorized services can access specific secrets.

Example: Payment service can access Stripe key, not database password.

IAM least privilege

6 Audit Logging

- 👉 Track all secret access for compliance.

Example: CloudTrail logs every Secrets Manager access → audit who accessed what when.

Audit logs compliance

7 Environment-Specific Secrets

- 👉 Use different secrets for dev, staging, production.

Example: Dev uses dev database credentials → limits blast radius of leaked secrets.

Environment separation blast radius

💡 Problem 104: Managing Secrets Securely

INTERVIEWER: Where do you store database passwords and API keys in production?

Question:

Solution:

Secrets must be encrypted, rotated, and never hardcoded or committed to version control.

1 Never Commit Secrets to Git

- 👉 Use .env files or secret managers, not hardcoded values.

Example: GitHub secret scanning detects committed AWS keys → auto-revoked → service breaks.

Git secrets security breach

2 Secret Management Services

- 👉 Use dedicated services for secret storage.

Example: AWS Secrets Manager, HashiCorp Vault, Azure Key Vault → encrypted storage.

Secrets Manager Vault Key Vault

3 Encryption at Rest & In Transit

- 👉 Secrets encrypted in storage and during transmission.

Example: Secrets Manager encrypts with KMS → accessed via TLS → end-to-end encryption.

Encryption KMS

4 Secret Rotation

- 👉 Automatically rotate secrets periodically.

Example: Database password rotated every 90 days → Secrets Manager updates app automatically.

Rotation automated updates

5 Least Privilege Access

- 👉 Only authorized services can access specific secrets.

Example: Payment service can access Stripe key, not database password.

IAM least privilege

6 Audit Logging

- 👉 Track all secret access for compliance.

Example: CloudTrail logs every Secrets Manager access → audit who accessed what when.

[Audit logs](#) [compliance](#)

7 Environment-Specific Secrets

- 👉 Use different secrets for dev, staging, production.

Example: Dev uses dev database credentials → limits blast radius of leaked secrets.

[Environment separation](#) [blast radius](#)



Problem 104: Preventing API Injection and Data Leakage

INTERVIEWER: How do you prevent attackers from exploiting your API to access unauthorized data?

Question:

Solution:

Injection attacks and data leakage exploit poor input validation and authorization — validate, sanitize, and enforce access control.

Here's how to prevent API injection and leakage 👇

1 Parameterized Queries for SQL Injection

- 👉 Never concatenate user input into queries.

Example: `db.query('SELECT * FROM users WHERE id = ?', [userId])` → not `SELECT * FROM users WHERE id = ' + userId`.

[Prepared statements](#) [ORM](#)

2 Validate Input Against Schema

- 👉 Reject requests with unexpected fields or types.

Example: Expect `{name: string, age: number}` → reject `{name: string, age: "drop table"}`.

JSON schema validation joi yup

3 NoSQL Injection Prevention

- 👉 Validate types and sanitize operators in NoSQL queries.

Example: MongoDB '{username: req.body.username}' → attacker sends '{username: {\$ne: null}}' → bypass auth.

Type checking sanitize operators

4 Check Resource Ownership

- 👉 Verify user owns resource before allowing access.

Example: User requests '/orders/123' → check 'orders.user_id == current_user.id' → deny if mismatch.

Authorization check ownership validation

5 Filter Response Data

- 👉 Don't return entire database objects — select specific fields.

Example: Return '{id, name, email}' → not entire user object with 'password_hash', 'internal_notes'.

Data minimization response filtering

6 Prevent Mass Assignment

- 👉 Don't blindly assign request body to database model.

Example: User sends '{name: "Alice", role: "admin"}' → only update allowed fields → ignore 'role'.

Whitelist allowed fields mass assignment protection

💡 Problem 105: Preventing SQL Injection Attacks

INTERVIEWER: A user enters in login field — how do you protect your database?

Question: ' OR '1'='1

Solution:

SQL injection is a top security risk that can expose or destroy your entire database.

1 Parameterized Queries (Prepared Statements)

- 👉 Use placeholders instead of string concatenation.

*Example: SELECT * FROM users WHERE email = ? → database treats input as data, not code.*

Prepared statements safe queries

2 ORM/Query Builders

- 👉 Use libraries that automatically parameterize queries.

Example: User.find({email: input}) in Sequelize → auto-escapes input.

ORM Sequelize TypeORM

3 Input Validation

- 👉 Validate and sanitize all user inputs.

Example: Email field → regex validation → reject if contains SQL keywords.

Input validation sanitization

4 Least Privilege Database User

- 👉 App database user should have minimal permissions.

Example: App user can only SELECT/INSERT users table, not DROP or ALTER.

Least privilege database permissions

5 Stored Procedures

- 👉 Pre-defined database procedures with fixed logic.

Example: CALL get_user(email) → prevents dynamic SQL injection.

Stored procedures encapsulation

6 WAF (Web Application Firewall)

- 👉 Filter malicious SQL patterns at network level.

Example: AWS WAF blocks requests with SQL keywords → prevents attacks before reaching app.

WAF CloudFlare ModSecurity

7 Database Activity Monitoring

- 👉 Alert on suspicious database queries.

Example: Query with 'DROP TABLE' → instant PagerDuty alert → block IP.

Monitoring anomaly detection

💡 Problem 105: Preventing SQL Injection Attacks

INTERVIEWER: A user enters in login field — how do you protect your database?

Question: ' OR '1'='1

Solution:

SQL injection is a top security risk that can expose or destroy your entire database.

1 Parameterized Queries (Prepared Statements)

- 👉 Use placeholders instead of string concatenation.

*Example: SELECT * FROM users WHERE email = ? → database treats input as data, not code.*

[Prepared statements](#) [safe queries](#)

2 ORM/Query Builders

- 👉 Use libraries that automatically parameterize queries.

Example: `User.find({email: input})` in Sequelize → auto-escapes input.

[ORM](#) [Sequelize](#) [TypeORM](#)

3 Input Validation

- 👉 Validate and sanitize all user inputs.

Example: Email field → regex validation → reject if contains SQL keywords.

[Input validation](#) [sanitization](#)

4 Least Privilege Database User

- 👉 App database user should have minimal permissions.

Example: App user can only SELECT/INSERT users table, not DROP or ALTER.

[Least privilege](#) [database permissions](#)

5 Stored Procedures

- 👉 Pre-defined database procedures with fixed logic.

Example: `CALL get_user(email)` → prevents dynamic SQL injection.

[Stored procedures](#) [encapsulation](#)

6 WAF (Web Application Firewall)

- 👉 Filter malicious SQL patterns at network level.

Example: AWS WAF blocks requests with SQL keywords → prevents attacks before reaching app.

[WAF](#) [CloudFlare](#) [ModSecurity](#)

7 Database Activity Monitoring

👉 Alert on suspicious database queries.

Example: Query with 'DROP TABLE' → instant PagerDuty alert → block IP.

Monitoring anomaly detection

Topic 37: Microservices Communication

💡 Problem 105: Synchronous vs Asynchronous Communication

INTERVIEWER: When should microservices communicate synchronously vs asynchronously?

Question:

Solution:

Synchronous communication couples services tightly, asynchronous decouples — choose based on requirements.

Here's when to use sync vs async communication 🤔

1 Synchronous Communication (REST/gRPC)

👉 Direct request/response between services — immediate response needed.

Example: User service calls payment service → waits for payment confirmation → returns order status.

REST API gRPC tight coupling latency dependency

2 Asynchronous Communication (Message Queues/Events)

- 👉 Services publish events, others consume asynchronously — fire-and-forget.

Example: Order created → publish OrderCreated event → inventory, notification, analytics services react independently.

Kafka RabbitMQ loose coupling eventual consistency

3 Use Sync for Critical Path Operations

- 👉 When user waits for immediate response and consistency required.

Example: Login API → validate credentials synchronously → return success/failure.

Real-time response user-facing operations

4 Use Async for Non-Critical Background Tasks

- 👉 When operation can complete later without blocking user.

Example: User uploads photo → return success immediately → resize image asynchronously in background.

Background jobs worker queues improved UX

5 Sync Creates Cascade Failures

- 👉 If downstream service fails, entire request chain fails.

Example: Service A calls B calls C → C is down → entire request fails.

Cascading failures circuit breaker needed

6 Async Provides Better Fault Tolerance

- 👉 Services continue operating even if consumers are down.

Example: Order service publishes events → notification service down → events queued → processed when service recovers.

Resilience eventual consistency retry mechanisms

Problem 4: Cross-Service Joins Without Shared Database

INTERVIEWER: Each microservice uses its own DB — so how do you handle cross-service joins?

Question:

Solution:

Distributed data requires denormalization, API composition, or event-driven synchronization.

1 API Composition Pattern

👉 Application layer queries multiple services and combines results.

Example: *Order details = call Order service + call User service + call Product service → merge in app.*

API Gateway composition BFF pattern

2 Denormalization / Data Duplication

👉 Store frequently joined data in same service database.

Example: *Order service stores {order_id, user_name, user_email} even though Users service owns user data.*

Controlled duplication eventual consistency

3 CQRS with Read Models

👉 Build specialized read databases optimized for specific queries.

Example: *OrderHistory service subscribes to events → builds denormalized view combining orders + users + products.*

CQRS read models materialized views

4 Event-Driven Synchronization

👉 Services publish events when data changes, others subscribe and update local copies.

Example: User service publishes *UserUpdated* event → Order service updates cached user data.

Event sourcing domain events Kafka

5 Backend for Frontend (BFF)

- 👉 Create service that aggregates data from multiple services for specific UI needs.

Example: Mobile BFF queries 5 services → returns single optimized response to mobile app.

BFF pattern aggregation layer

6 GraphQL Federation

- 👉 Use GraphQL to query across multiple services transparently.

Example: Single GraphQL query fetches order (Order service) + user (User service) + products (Product service).

Apollo Federation GraphQL stitching

7 Avoid Cross-Service Joins

- 👉 Redesign bounded contexts to minimize need for joins.

Example: Instead of joining Orders + Users, make Orders service store all needed user data at order time.

Domain-driven design bounded contexts



Problem 4: Cross-Service Joins Without Shared Database

INTERVIEWER: Each microservice uses its own DB — so how do you handle cross-service joins?

Question:

Solution:

Distributed data requires denormalization, API composition, or event-driven synchronization.

1 API Composition Pattern

- 👉 Application layer queries multiple services and combines results.

Example: Order details = call Order service + call User service + call Product service → merge in app.

API Gateway composition BFF pattern

2 Denormalization / Data Duplication

- 👉 Store frequently joined data in same service database.

Example: Order service stores {order_id, user_name, user_email} even though Users service owns user data.

Controlled duplication eventual consistency

3 CQRS with Read Models

- 👉 Build specialized read databases optimized for specific queries.

Example: OrderHistory service subscribes to events → builds denormalized view combining orders + users + products.

CQRS read models materialized views

4 Event-Driven Synchronization

- 👉 Services publish events when data changes, others subscribe and update local copies.

Example: User service publishes UserUpdated event → Order service updates cached user data.

Event sourcing domain events Kafka

5 Backend for Frontend (BFF)

- 👉 Create service that aggregates data from multiple services for specific UI needs.

Example: Mobile BFF queries 5 services → returns single optimized response to mobile app.

BFF pattern aggregation layer

6 GraphQL Federation

- 👉 Use GraphQL to query across multiple services transparently.

Example: Single GraphQL query fetches order (Order service) + user (User service) + products (Product service).

Apollo Federation

GraphQL stitching

7 Avoid Cross-Service Joins

- 👉 Redesign bounded contexts to minimize need for joins.

Example: Instead of joining Orders + Users, make Orders service store all needed user data at order time.

Domain-driven design

bounded contexts

Consensus Algorithms

💡 Problem 106: Service Discovery and Load Balancing

INTERVIEWER: How do microservices find and connect to each other?

Question:

Solution:

Service discovery enables dynamic routing — services register and discover each other at runtime.

Here's how service discovery works ↗

1 Client-Side Service Discovery

- 👉 Client queries service registry to find available instances.

Example: Service A queries Consul → gets list of Service B instances → picks one → makes request.

Eureka Consul client-side load balancing

2 Server-Side Service Discovery

- 👉 Client makes request to load balancer, which queries registry and routes.

Example: Service A calls load balancer → load balancer queries registry → routes to healthy Service B instance.

AWS ALB Nginx simpler clients

3 Service Registry (Consul, etcd, Eureka)

- 👉 Central database of service instances and their health status.

Example: Service B instances register with Consul → Service A discovers from Consul.

Service registry health checks heartbeats

4 DNS-Based Discovery

- 👉 Use DNS to resolve service names to IP addresses.

Example: Service A queries 'service-b.internal' → DNS returns IP addresses of healthy instances.

Kubernetes DNS AWS Cloud Map simple but limited

5 Client-Side Load Balancing

- 👉 Client library chooses which instance to call.

Example: Service A has list of Service B instances → uses round-robin to pick one.

Ribbon client libraries flexible algorithms

6 Health Checks for Instance Removal

- 👉 Registry removes unhealthy instances from available pool.

Example: Service B instance fails health check → removed from registry → traffic stops routing to it.

Health endpoints TTL automatic deregistration

💡 Problem 107: API Gateway Pattern

INTERVIEWER: Should clients call microservices directly or through an API gateway?

Question:

Solution:

API gateway provides single entry point — handles routing, authentication, rate limiting for all services.

Here's how API gateway pattern works 🤝

1 Single Entry Point for Clients

👉 Gateway exposes unified API, routes requests to appropriate microservices.

Example: Mobile app calls '/api/orders' → gateway routes to order service.

Kong AWS API Gateway Apigee single endpoint

2 Authentication and Authorization

👉 Gateway validates auth tokens before routing to services.

Example: Client sends JWT → gateway validates → extracts user_id → passes to service.

Centralized auth services trust gateway

3 Rate Limiting and Throttling

👉 Gateway enforces rate limits per client or API key.

Example: Free tier limited to 100 req/min → gateway enforces → returns 429 if exceeded.

Centralized rate limiting tiered access

4 Request/Response Transformation

👉 Gateway transforms external API format to internal service format.

Example: External API uses camelCase → gateway converts to snake_case for services.

Protocol translation format adaptation

5 Aggregation of Multiple Services

👉 Gateway calls multiple services and combines responses.

Example: User dashboard → gateway calls user service, order service, notification service → combines into single response.

BFF pattern reduce client round trips

6 Gateway as Single Point of Failure

👉 If gateway goes down, entire API unavailable — need high availability.

Example: Deploy multiple gateway instances behind load balancer.

HA redundancy health checks

Topic 38: Saga Pattern

Service Discovery

💡 Problem 108: Understanding Saga Pattern for Distributed Transactions

INTERVIEWER: How do you handle transactions across multiple microservices?

Question:

Solution:

Sagas break distributed transactions into local transactions with compensation — eventual consistency instead of ACID.

Here's how saga pattern works 

1 What is a Saga?

👉 Sequence of local transactions, each with compensating transaction for rollback.

Example: Order → Reserve Inventory → Charge Payment → if payment fails → Unreserve Inventory.

Distributed transaction eventual consistency

2 Choreography-Based Saga

👉 Each service publishes events, other services react independently.

Example: Order service publishes OrderCreated → Inventory service reserves → publishes InventoryReserved → Payment service charges.

Event-driven decentralized complex to trace

3 Orchestration-Based Saga

👉 Central orchestrator coordinates saga steps.

Example: Saga orchestrator tells Order service → then Inventory → then Payment → handles failures.

Centralized control easier to understand single point of failure

4 Compensating Transactions

👉 Undo previous steps if later step fails.

Example: Payment fails → run compensation → unreserve inventory → cancel order.

Rollback logic compensation handlers

5 Saga State Machine

- 👉 Track saga progress through states.

Example: States: *OrderCreated* → *InventoryReserved* → *PaymentProcessed* → *Completed* or *OrderCreated* → *InventoryReserved* → *PaymentFailed* → *Compensating* → *Cancelled*.

[State tracking](#) [recovery](#)

6 Idempotency for Retry Safety

- 👉 Each saga step must be idempotent for safe retries.

Example: Reserve inventory request sent twice → only reserves once.

[Idempotency keys](#) [safe retries](#)

💡 Problem 109: Saga Failure Handling and Compensation

INTERVIEWER: What happens if a step in the middle of a saga fails?

Question:

[Answer](#)

Solution:

Sagas must compensate completed steps when failure occurs — undo in reverse order.

Here's how to handle saga failures 🤞

1 Backward Recovery (Compensation)

- 👉 Execute compensating transactions in reverse order.

Example: *Order* → *Reserve Inventory* → *Charge Payment fails* → *Unreserve Inventory* → *Cancel Order*.

[Rollback](#) [compensating actions](#)

2 Forward Recovery (Retry)

- 👉 Retry failed step until success.

Example: Payment service temporarily down → retry payment → eventually succeeds.

Retry with exponential backoff eventual success

3 Semantic Lock for Reservations

- 👉 Reserve resources without committing until saga completes.

Example: Inventory reserved but not decremented → if saga fails → release reservation.

Pessimistic locking reservation pattern

4 Saga Log for Recovery

- 👉 Store saga state for crash recovery.

Example: Orchestrator crashes mid-saga → restart → read log → resume from last step.

Persistence recovery database log

5 Timeout Handling

- 👉 If step doesn't complete in time, trigger compensation.

Example: Payment service doesn't respond after 30 seconds → assume failure → compensate.

Timeout policies failure detection

6 Manual Intervention for Non-Compensable Steps

- 👉 Some actions can't be undone — require manual handling.

Example: Email sent to user → can't unsend → send apology email or mark as error.

Human intervention support queue

Container Orchestration

Problem 110: Saga vs Two-Phase Commit (2PC)

INTERVIEWER: Why not use two-phase commit for distributed transactions?

Question:

Solution:

2PC provides strong consistency but blocks on failures — sagas trade consistency for availability.

Here's saga vs 2PC trade-offs 

1 Two-Phase Commit (2PC)

 Coordinator asks all participants to prepare → all agree → coordinator commits all.

Example: Coordinator: "Prepare?" → All: "Yes" → Coordinator: "Commit now" → All commit.

XA transactions strong consistency blocking

2 2PC Blocks on Coordinator Failure

 If coordinator crashes after prepare, participants locked indefinitely.

Example: Coordinator crashes → participants in prepared state → can't commit or abort → resources locked.

Blocking protocol single point of failure

3 Sagas Are Non-Blocking

 Each step commits immediately, compensation undoes if needed.

Example: Payment commits → inventory commits → if later step fails → run compensations.

Non-blocking eventual consistency

4 2PC Requires All Services Up

- 👉 If one service down, entire transaction blocks.

Example: Payment service down → 2PC can't complete → blocks forever or times out.

High availability requirement coupled services

5 Sagas Provide Better Availability

- 👉 Services continue operating, eventual consistency achieved.

Example: Payment service down → saga retries → eventually succeeds → saga completes.

Resilience retry mechanisms

6 When to Use Each

- 👉 Use 2PC for tightly coupled systems requiring strong consistency (rare in microservices). Use sagas for loosely coupled services with eventual consistency.

Example: Banking core system → 2PC. E-commerce order flow → Saga.

Trade-offs choose based on requirements

Topic 39: API Gateway Pattern

💡 Problem 111: Implementing Backend for Frontend (BFF)

INTERVIEWER: Your mobile and web apps need different API responses. How do you handle this?

Question:

Solution:

BFF creates separate backend API for each frontend — optimized for specific client needs.

Here's how BFF pattern works ↗

1 Separate API Gateway per Client Type

👉 Mobile BFF, web BFF, desktop BFF — each optimized for that client.

Example: Mobile BFF returns minimal data for bandwidth, web BFF returns full data.

Client-specific API optimized responses

2 Aggregate Multiple Service Calls

👉 BFF calls multiple microservices, combines responses.

Example: Dashboard request → BFF calls user service, order service, notification service → returns combined JSON.

Reduce client round trips N+1 problem avoided

3 Transform Data for Client Format

👉 BFF adapts internal service format to client-friendly format.

Example: Services return snake_case → BFF converts to camelCase for JavaScript clients.

Data transformation protocol adaptation

4 Client-Specific Business Logic

👉 BFF implements logic specific to client needs.

Example: Mobile BFF returns paginated results with limit=10, web BFF returns limit=50.

Client-optimized logic experience customization

5 BFF Owned by Frontend Team

👉 Frontend team controls BFF, can iterate independently.

Example: Mobile team updates mobile BFF without affecting web team.

Team autonomy independent deployment

6 Avoid Duplication Across BFFs

- 👉 Extract shared logic to libraries or shared services.

Example: Authentication logic shared across all BFFs via library.

Code reuse shared utilities

📘 Blue-Green Deployment

💡 Problem 112: GraphQL as API Gateway Alternative

INTERVIEWER: Should you use GraphQL instead of REST API gateway?

Question:

Solution:

GraphQL lets clients request exactly what they need — reduces over-fetching and under-fetching.

Here's when to use GraphQL 🤝

1 Client Specifies Required Fields

- 👉 Client queries only needed fields, avoids over-fetching.

Example: Query '{ user { name, email } }' → returns only name and email, not full user object.

Flexible queries bandwidth optimization

2 Fetch Related Data in Single Request

- 👉 GraphQL resolves nested relationships without N+1 requests.

Example: Query '{ user { orders { items } } }' → single request gets user, orders, and items.

Nested queries eliminates multiple round trips

3 Schema-Driven API

- 👉 GraphQL schema defines available types and fields.

Example: Type definitions → clients discover API capabilities → strongly typed.

Type safety introspection schema validation

4 Use GraphQL Gateway Over Microservices

- 👉 GraphQL server aggregates data from multiple microservices.

Example: GraphQL resolvers call order service, user service, inventory service → combine data.

Federation Apollo Gateway schema stitching

5 Trade-offs: Complexity and Caching

- 👉 GraphQL adds complexity, harder to cache than REST.

Example: REST '/users/123' cached by URL → GraphQL POST with query body harder to cache.

HTTP caching CDN complexity trade-off

6 When to Use REST vs GraphQL

- 👉 Use REST for simple CRUD, public APIs, caching. Use GraphQL for complex data fetching, mobile apps.

Example: Public API with rate limiting → REST. Mobile app with varied data needs → GraphQL.

Choose based on requirements not hype

Topic 40: Serverless Architecture



Problem 113: When to Use Serverless (AWS Lambda, Cloud Functions)

INTERVIEWER: Should you use serverless functions or traditional servers?

Question:

Solution:

Serverless is great for event-driven, variable workloads — but has cold start and execution limits.

Here's when to use serverless 📈

1 Event-Driven Workloads

👉 Functions triggered by events like file uploads, HTTP requests, queue messages.

Example: S3 upload triggers Lambda to resize image.

AWS Lambda Google Cloud Functions event triggers

2 Variable or Unpredictable Traffic

👉 Auto-scales from zero to thousands of concurrent executions.

Example: Black Friday sale → traffic spikes 100x → Lambda scales automatically.

Auto-scaling pay per execution

3 Short-Lived Tasks

👉 Functions run for seconds to minutes, not long-running processes.

Example: API endpoint, data transformation, webhook handler.

Max 15 minutes on Lambda short executions

4 Avoid Serverless for Persistent Connections

- 👉 WebSockets, database connection pools don't work well.

Example: WebSocket server needs persistent connection → use container instead.

Connection limits cold starts

5 Cold Start Latency

- 👉 First request after idle takes longer to start function.

Example: Function idle for 5 minutes → next request waits 1-2 seconds for cold start.

Cold start provisioned concurrency option

6 Cost Effective for Low Volume

- 👉 Pay only for execution time, no idle server costs.

Example: Function runs 1M times/month at 100ms each → cheaper than running server 24/7.

Pay-per-use cost optimization

Bloom Filters

💡 Problem 114: Serverless Best Practices

INTERVIEWER: How do you optimize serverless functions for performance and cost?

Question:

[Redacted]

Solution:

Minimize cold starts, optimize bundle size, use async patterns — serverless has unique constraints.

Here are serverless best practices 👉

1 Minimize Cold Start Time

- 👉 Keep deployment package small, minimize dependencies.

Example: Lambda with 50MB package → slow cold start. 5MB package → fast cold start.

Bundle optimization tree shaking reduce dependencies

2 Reuse Connections Outside Handler

- 👉 Initialize database connections, HTTP clients outside handler function.

Example: Database pool created once → reused across invocations in same container.

Connection reuse warm containers

3 Use Provisioned Concurrency for Critical Paths

- 👉 Keep functions warm to eliminate cold starts.

Example: API endpoint with strict latency SLA → provision 10 warm instances.

Provisioned concurrency cost vs performance

4 Avoid Heavy Synchronous Chains

- 👉 Don't chain many Lambda functions synchronously.

Example: Lambda 1 → Lambda 2 → Lambda 3 → Lambda 4 → slow and expensive. Use async events instead.

Step Functions async patterns

5 Right-Size Memory Allocation

- 👉 More memory = more CPU → sometimes faster and cheaper.

Example: Function with 128MB takes 10 seconds, 1024MB takes 2 seconds → 1024MB might cost less.

Memory optimization CPU allocation

6 Use Environment Variables for Config

- 👉 Avoid hardcoded values, use env vars for config.

Example: Database URL in environment variable → change without redeploying code.

Configuration management flexibility

Problem 115: Serverless Data Processing Patterns

INTERVIEWER: How do you process large amounts of data with serverless?

Question:

Solution:

Serverless excels at parallel processing — trigger multiple functions to process batches concurrently.

Here are serverless data processing patterns 

1 Fan-Out Pattern for Parallel Processing

👉 Split work into chunks, process each chunk in separate function.

Example: Process 1M records → split into 100 batches of 10K → 100 Lambdas process in parallel.

Parallel processing SQS fanout high throughput

2 S3 Event Triggers for File Processing

👉 File uploaded to S3 triggers Lambda to process.

Example: CSV uploaded → Lambda parses and inserts into database.

Event-driven S3 notifications automatic trigger

3 Stream Processing with Kinesis/Kafka

👉 Lambda consumes events from stream, processes in batches.

Example: Clickstream data in Kinesis → Lambda aggregates and writes to database.

Real-time processing stream batching

4 Step Functions for Orchestration

- 👉 Coordinate multiple Lambda functions in workflow.

Example: ETL pipeline → Step 1: Extract → Step 2: Transform → Step 3: Load → Step Functions orchestrates.

Workflow state machine error handling

5 Use SQS for Rate Limiting and Retry

- 👉 Queue messages, Lambda pulls at controlled rate.

Example: 1M messages in SQS → Lambda processes 100 concurrently → rate-limited processing.

Queue-based backpressure retry on failure

\n

6 Avoid Timeouts with Partial Processing

- 👉 Process batch partially, requeue remaining work.

Example: Lambda has 10 minutes to process 100K records → processes 50K → requeues rest.

Chunking checkpoint pattern progress tracking

💡 Problem 116: Handling Large File Uploads at Scale

INTERVIEWER: Your users upload 50GB of videos every hour — where do you store them?

Question:

Solution:

Scalable file storage requires object storage, not local disks or databases.

1 Object Storage (S3, GCS, Azure Blob)

- 👉 Store files in cloud object storage designed for massive scale.

Example: 50GB/hour = 1.2TB/day → S3 handles this easily with 99.999999999% durability.

AWS S3 Google Cloud Storage Azure Blob Storage

2 Direct Upload to S3 (Pre-signed URLs)

- 👉 Let users upload directly to S3, bypassing your servers.

Example: Generate pre-signed URL → client uploads 2GB video directly to S3 → no server bandwidth used.

S3 pre-signed URLs direct upload pattern

3 Multipart Upload

- 👉 Split large files into chunks for parallel, resumable uploads.

Example: 5GB video split into 100MB chunks → upload in parallel → auto-retry failed chunks.

S3 multipart upload resumable uploads

4 CDN Integration

- 👉 Serve uploaded files through CDN for fast global access.

Example: Video uploaded to S3 us-east-1 → served via CloudFront edge locations worldwide.

CloudFront Cloudflare R2

5 Lifecycle Policies

- 👉 Automatically move old files to cheaper storage tiers.

Example: Videos >90 days old → moved to Glacier (\$0.004/GB vs S3's \$0.023/GB).

S3 Lifecycle storage classes archival

6 Metadata in Database

- 👉 Store file metadata (URL, size, owner) in database, not files themselves.

Example: uploads table with {file_id, s3_key, user_id, size, created_at}.

Metadata separation database for search

7 Virus Scanning

- 👉 Scan uploaded files asynchronously before making public.

Example: Upload triggers Lambda → ClamAV scan → if clean, move to public bucket.

ClamAV async processing quarantine bucket

💡 Problem 117: Choosing S3 Storage Classes

INTERVIEWER: You store 500TB of user uploads — how do you minimize S3 costs?

Question:

Solution:

S3 storage classes balance access speed, durability, and cost for different use cases.

1 S3 Standard

- 👉 Frequently accessed data, highest cost.

Example: \$0.023/GB/month → use for recent uploads accessed daily.

S3 Standard frequent access

2 S3 Intelligent-Tiering

- 👉 Automatically moves objects between tiers based on access patterns.

Example: File not accessed for 30 days → auto-moved to cheaper tier → saves cost.

Intelligent-Tiering automatic optimization

3 S3 Standard-IA (Infrequent Access)

- 👉 Lower storage cost, but charges for retrieval.

Example: \$0.0125/GB/month + \$0.01/GB retrieval → for backups accessed monthly.

Standard-IA infrequent access

4 S3 One Zone-IA

- 👉 Cheaper than Standard-IA, single AZ (lower durability).

Example: \$0.01/GB/month → for reproducible data like thumbnails.

One Zone-IA single AZ

5 S3 Glacier Instant Retrieval

- 👉 Archive storage with millisecond retrieval.

Example: \$0.004/GB/month → for archives accessed quarterly.

Glacier Instant archive

6 S3 Glacier Flexible/Deep Archive

- 👉 Lowest cost, slow retrieval (hours).

Example: \$0.00099/GB/month (Deep Archive) → for compliance archives, 7-year retention.

Glacier Deep Archive long-term archive

7 Lifecycle Policies

- 👉 Automatically transition objects between classes.

Example: Standard (30 days) → Standard-IA (90 days) → Glacier (1 year) → Deep Archive.

Lifecycle rules automated transitions

Problem 117: Choosing S3 Storage Classes

INTERVIEWER: You store 500TB of user uploads — how do you minimize S3 costs?

Question:

Solution:

S3 storage classes balance access speed, durability, and cost for different use cases.

1 S3 Standard

- 👉 Frequently accessed data, highest cost.

Example: \$0.023/GB/month → use for recent uploads accessed daily.

S3 Standard frequent access

2 S3 Intelligent-Tiering

- 👉 Automatically moves objects between tiers based on access patterns.

Example: File not accessed for 30 days → auto-moved to cheaper tier → saves cost.

Intelligent-Tiering automatic optimization

3 S3 Standard-IA (Infrequent Access)

- 👉 Lower storage cost, but charges for retrieval.

Example: \$0.0125/GB/month + \$0.01/GB retrieval → for backups accessed monthly.

Standard-IA infrequent access

4 S3 One Zone-IA

- 👉 Cheaper than Standard-IA, single AZ (lower durability).

Example: \$0.01/GB/month → for reproducible data like thumbnails.

One Zone-IA single AZ

5 S3 Glacier Instant Retrieval

- 👉 Archive storage with millisecond retrieval.

Example: \$0.004/GB/month → for archives accessed quarterly.

Glacier Instant archive

6 S3 Glacier Flexible/Deep Archive

- 👉 Lowest cost, slow retrieval (hours).

Example: \$0.00099/GB/month (Deep Archive) → for compliance archives, 7-year retention.

Glacier Deep Archive long-term archive

\n

7 Lifecycle Policies

- 👉 Automatically transition objects between classes.

Example: Standard (30 days) → Standard-IA (90 days) → Glacier (1 year) → Deep Archive.

Lifecycle rules automated transitions



Problem 117: Building Sub-50ms Real-Time Chat

INTERVIEWER: Your chat needs messages in <50ms — how would you build that?

Question:

Solution:

Real-time communication requires WebSockets and in-memory message passing.

1 WebSocket Connections

- 👉 Use WebSockets for bidirectional, low-latency communication.

Example: HTTP polling: 500ms latency; WebSocket: 10–30ms latency.

WebSockets Socket.io SignalR

2 In-Memory Message Broker

- 👉 Use Redis Pub/Sub for real-time message distribution.

Example: User sends message → published to Redis → all chat server instances get it in <5ms.

Redis Pub/Sub in-memory messaging

3 Connection Management

- 👉 Keep persistent WebSocket connections open per user.

Example: 10K concurrent users = 10K open WebSocket connections across server pool.

Connection pooling stateful connections

4 Horizontal Scaling with Sticky Sessions

- 👉 Route user's WebSocket connection to same server instance.

Example: Load balancer uses user_id hash → consistent server routing.

Sticky sessions consistent hashing

5 Message Persistence (Asynchronous)

- 👉 Store messages in database asynchronously, don't block delivery.

Example: Message delivered via WebSocket in 20ms → saved to DB in background (200ms).

Async writes eventual persistence

6 Presence System

- 👉 Track online/offline status with Redis or similar.

Example: User connects → SET user:123:online with 60s TTL → heartbeat every 30s.

Redis TTL heartbeat mechanism

7 Geographic Distribution

👉 Deploy chat servers in multiple regions for lower latency.

Example: US users connect to us-east servers, EU users to eu-west → <50ms regional latency.

Multi-region edge servers

💡 Problem 118: Scaling WebSocket Connections

INTERVIEWER: You have 100K concurrent WebSocket connections — how do you scale them?

Question:

Solution:

WebSocket connections are stateful and long-lived, requiring different scaling patterns than HTTP.

1 Stateful Connection Challenge

👉 Each user connected to specific server instance.

Example: User on server1 → can't handle message from user on server2 without coordination.

Stateful sticky connections

2 Redis Pub/Sub for Message Broadcasting

👉 Publish messages to all server instances via Redis.

Example: User on server1 sends message → published to Redis → all servers receive → broadcast to connected users.

Redis Pub/Sub message broadcasting

3 Sticky Sessions (Load Balancer)

- 👉 Route user to same server for connection lifetime.

Example: ALB uses source IP hash → user always routed to same server.

Sticky sessions session affinity

4 Connection Limits Per Server

- 👉 Each server has max connection capacity.

Example: Node.js server handles ~10K connections → 100K users = 10 servers minimum.

Connection limits capacity planning

5 Horizontal Scaling

- 👉 Add more servers as connection count grows.

Example: Auto-scale based on connection count metric → add server when > 8K connections/server.

Auto-scaling horizontal scaling

6 Dedicated WebSocket Servers

- 👉 Separate WebSocket servers from HTTP API servers.

Example: HTTP API handles REST, WebSocket servers handle only real-time → optimized separately.

Service separation optimization

7 Connection State Management

- 👉 Store connection metadata in shared storage.

Example: Track which user connected to which server in Redis → route messages correctly.

State management connection tracking

Problem 118: GraphQL vs REST Trade-offs

INTERVIEWER: When should you choose GraphQL over REST, and what are the trade-offs?

Question:

Solution:

GraphQL solves specific API problems but introduces new complexity.

1 Over-fetching Problem in REST

👉 REST endpoints return fixed data structures, often with unnecessary fields.

Example: GET /users returns 20 fields, but mobile app only needs name and email → wasted bandwidth.

REST limitations data over-fetching

2 Under-fetching Problem in REST

👉 Need multiple REST calls to get related data.

Example: Get user → separate call for posts → separate call for comments → 3 round trips.

N+1 problem multiple requests

3 GraphQL Single Endpoint

👉 One endpoint with flexible queries for exact data needs.

Example: Single query fetches user + posts + comments in one request with only needed fields.

Flexible queries single endpoint

4 GraphQL Complexity

👉 Requires schema definition, resolvers, and query complexity analysis.

Example: Must prevent malicious deep queries that could DoS your server.

[Query complexity](#) [depth limiting](#)

5 Caching Challenges

- 👉 GraphQL makes HTTP caching harder due to POST requests.

Example: REST GET /users/123 → easily cached; GraphQL POST with query → needs custom caching.

[Caching complexity](#) [Apollo cache](#)

6 When to Use REST

- 👉 Simple CRUD operations, public APIs, heavy caching needs.

Example: Public weather API with standard endpoints → REST is simpler.

[REST advantages](#) [simple APIs](#)

7 When to Use GraphQL

- 👉 Complex data relationships, multiple clients with different needs, real-time updates.

Example: Social media app with web, mobile, smart TV → each needs different data subsets.

[GraphQL advantages](#) [flexible clients](#)

💡 Problem 118: Scaling WebSocket Connections

INTERVIEWER: You have 100K concurrent WebSocket connections — how do you scale them?

Question:

Solution:

WebSocket connections are stateful and long-lived, requiring different scaling patterns than HTTP.

1 Stateful Connection Challenge

- 👉 Each user connected to specific server instance.

Example: User on server1 → can't handle message from user on server2 without coordination.

Stateful sticky connections

2 Redis Pub/Sub for Message Broadcasting

- 👉 Publish messages to all server instances via Redis.

Example: User on server1 sends message → published to Redis → all servers receive → broadcast to connected users.

Redis Pub/Sub message broadcasting

3 Sticky Sessions (Load Balancer)

- 👉 Route user to same server for connection lifetime.

Example: ALB uses source IP hash → user always routed to same server.

Sticky sessions session affinity

4 Connection Limits Per Server

- 👉 Each server has max connection capacity.

Example: Node.js server handles ~10K connections → 100K users = 10 servers minimum.

Connection limits capacity planning

5 Horizontal Scaling

- 👉 Add more servers as connection count grows.

Example: Auto-scale based on connection count metric → add server when > 8K connections/server.

Auto-scaling horizontal scaling

6 Dedicated WebSocket Servers

- 👉 Separate WebSocket servers from HTTP API servers.

Example: HTTP API handles REST, WebSocket servers handle only real-time → optimized separately.

Service separation optimization

7 Connection State Management

👉 Store connection metadata in shared storage.

Example: Track which user connected to which server in Redis → route messages correctly.

State management connection tracking

💡 Problem 118: GraphQL vs REST Trade-offs

INTERVIEWER: When should you choose GraphQL over REST, and what are the trade-offs?

Question:

Solution:

GraphQL solves specific API problems but introduces new complexity.

1 Over-fetching Problem in REST

👉 REST endpoints return fixed data structures, often with unnecessary fields.

Example: GET /users returns 20 fields, but mobile app only needs name and email → wasted bandwidth.

REST limitations data over-fetching

2 Under-fetching Problem in REST

👉 Need multiple REST calls to get related data.

Example: Get user → separate call for posts → separate call for comments → 3 round trips.

N+1 problem multiple requests

3 GraphQL Single Endpoint

- 👉 One endpoint with flexible queries for exact data needs.

Example: Single query fetches user + posts + comments in one request with only needed fields.

Flexible queries single endpoint

4 GraphQL Complexity

- 👉 Requires schema definition, resolvers, and query complexity analysis.

Example: Must prevent malicious deep queries that could DoS your server.

Query complexity depth limiting

5 Caching Challenges

- 👉 GraphQL makes HTTP caching harder due to POST requests.

Example: REST GET /users/123 → easily cached; GraphQL POST with query → needs custom caching.

Caching complexity Apollo cache

6 When to Use REST

- 👉 Simple CRUD operations, public APIs, heavy caching needs.

Example: Public weather API with standard endpoints → REST is simpler.

REST advantages simple APIs

\n\n\n

7 When to Use GraphQL

- 👉 Complex data relationships, multiple clients with different needs, real-time updates.

Example: Social media app with web, mobile, smart TV → each needs different data subsets.

GraphQL advantages flexible clients

Problem 116: Handling Large File Uploads at Scale

INTERVIEWER: Your users upload 50GB of videos every hour — where do you store them?

Question:

Solution:

Scalable file storage requires object storage, not local disks or databases.

1 Object Storage (S3, GCS, Azure Blob)

👉 Store files in cloud object storage designed for massive scale.

Example: 50GB/hour = 1.2TB/day → S3 handles this easily with 99.999999999% durability.

AWS S3 Google Cloud Storage Azure Blob Storage

2 Direct Upload to S3 (Pre-signed URLs)

👉 Let users upload directly to S3, bypassing your servers.

Example: Generate pre-signed URL → client uploads 2GB video directly to S3 → no server bandwidth used.

S3 pre-signed URLs direct upload pattern

3 Multipart Upload

👉 Split large files into chunks for parallel, resumable uploads.

Example: 5GB video split into 100MB chunks → upload in parallel → auto-retry failed chunks.

S3 multipart upload resumable uploads

4 CDN Integration

- 👉 Serve uploaded files through CDN for fast global access.

Example: Video uploaded to S3 us-east-1 → served via CloudFront edge locations worldwide.

CloudFront Cloudflare R2

5 Lifecycle Policies

- 👉 Automatically move old files to cheaper storage tiers.

Example: Videos >90 days old → moved to Glacier (\$0.004/GB vs S3's \$0.023/GB).

S3 Lifecycle storage classes archival

6 Metadata in Database

- 👉 Store file metadata (URL, size, owner) in database, not files themselves.

Example: uploads table with {file_id, s3_key, user_id, size, created_at}.

Metadata separation database for search

\n

7 Virus Scanning

- 👉 Scan uploaded files asynchronously before making public.

Example: Upload triggers Lambda → ClamAV scan → if clean, move to public bucket.

ClamAV async processing quarantine bucket

💡 Problem 117: Building Sub-50ms Real-Time Chat

INTERVIEWER: Your chat needs messages in <50ms — how would you build that?

Question:

Solution:

Real-time communication requires WebSockets and in-memory message passing.

1 WebSocket Connections

- 👉 Use WebSockets for bidirectional, low-latency communication.

Example: HTTP polling: 500ms latency; WebSocket: 10-30ms latency.

WebSockets Socket.io SignalR

2 In-Memory Message Broker

- 👉 Use Redis Pub/Sub for real-time message distribution.

Example: User sends message → published to Redis → all chat server instances get it in <5ms.

Redis Pub/Sub in-memory messaging

3 Connection Management

- 👉 Keep persistent WebSocket connections open per user.

Example: 10K concurrent users = 10K open WebSocket connections across server pool.

Connection pooling stateful connections

4 Horizontal Scaling with Sticky Sessions

- 👉 Route user's WebSocket connection to same server instance.

Example: Load balancer uses user_id hash → consistent server routing.

Sticky sessions consistent hashing

5 Message Persistence (Asynchronous)

- 👉 Store messages in database asynchronously, don't block delivery.

Example: Message delivered via WebSocket in 20ms → saved to DB in background (200ms).

Async writes eventual persistence

6 Presence System

- 👉 Track online/offline status with Redis or similar.

Example: User connects → SET user:123:online with 60s TTL → heartbeat every 30s.

Redis TTL heartbeat mechanism

\n\n

7 Geographic Distribution

- 👉 Deploy chat servers in multiple regions for lower latency.

Example: US users connect to us-east servers, EU users to eu-west → <50ms regional latency.

Multi-region edge servers

Topic 41: Auto Scaling Strategies

💡 Problem 116: Application Auto Scaling

INTERVIEWER: How do you automatically scale your application based on load?

Question:

Solution:

Auto scaling adds/removes instances based on metrics — scale out for traffic, scale in to save cost.

Here's how auto scaling works 👉

1 Metric-Based Scaling

- 👉 Scale based on CPU, memory, request count, or custom metrics.

Example: CPU > 70% for 5 minutes → add 2 instances. CPU < 30% → remove 1 instance.

CloudWatch metrics threshold-based

2 Target Tracking Scaling

- 👉 Maintain metric at target value.

Example: Target 1000 requests per instance → traffic increases → auto scaling adds instances to maintain target.

AWS target tracking automatic calculation

3 Scheduled Scaling

- 👉 Scale based on known patterns.

Example: Scale up every weekday at 8 AM, scale down at 6 PM.

Predictable traffic cron-based

4 Step Scaling for Rapid Changes

- 👉 Add/remove multiple instances based on alarm severity.

Example: CPU 70-80% → add 1 instance, 80-90% → add 3 instances, >90% → add 5 instances.

Aggressive scaling step policies

5 Cooldown Period

- 👉 Wait after scaling action before next scaling.

Example: Scale up → wait 5 minutes → check if more scaling needed → prevents flapping.

Stabilization prevent thrashing

\n

6 Scale Out Faster Than Scale In

- 👉 Add capacity quickly, remove slowly.

Example: Scale out immediately on high load, wait 10 minutes before scaling in.

Aggressive scale out conservative scale in

💡 Problem 117: Handling Overnight Viral Traffic

INTERVIEWER: Your app goes viral overnight — what's the first thing that breaks?

Question:

Solution:

Sudden traffic spikes expose bottlenecks in least scalable components.

1 Database Connection Pool Exhaustion

👉 Fixed-size connection pools can't scale with traffic.

Example: 10K concurrent users → 10K app server connections → but DB pool only 100 connections → timeouts.

Connection pooling PgBouncer read replicas

2 Rate Limit External APIs

👉 Third-party API rate limits hit unexpectedly.

Example: 10x traffic → 10x payment API calls → hit Stripe rate limit → checkout fails.

Circuit breakers queue external calls upgrade API tiers

3 Memory Exhaustion

👉 Application servers run out of memory with increased load.

Example: Each user session = 10MB → 10K users = 100GB RAM → servers crash.

Stateless design external session storage horizontal scaling

4 CDN Origin Overload

- 👉 CDN cache misses overwhelm origin servers.

Example: Viral post links to uncached page → 100K users → all hit origin → origin crashes.

Cache warming higher CDN cache TTL origin rate limiting

5 Database Write Bottleneck

- 👉 Single primary database can't handle write traffic.

Example: 10K sign-ups/minute → single DB primary maxed out → sign-up fails.

Write sharding queue writes batch inserts

6 Logging/Monitoring Overload

- 👉 Excessive logging consumes resources.

Example: 100x traffic → 100x logs → disk full → app crashes trying to write logs.

Log sampling async logging log levels

\n\n

7 Auto-Scaling Lag

- 👉 Auto-scaling reacts too slowly to sudden spike.

Example: Traffic spikes in 30 seconds → auto-scaling takes 5 minutes to launch instances → app down in between.

Pre-warming faster scaling over-provisioning buffer

 **Problem 117: Handling Overnight Viral Traffic**

INTERVIEWER: Your app goes viral overnight — what's the first thing that breaks?

Question:

Solution:

Sudden traffic spikes expose bottlenecks in least scalable components.

1 Database Connection Pool Exhaustion

👉 Fixed-size connection pools can't scale with traffic.

Example: 10K concurrent users → 10K app server connections → but DB pool only 100 connections → timeouts.

Connection pooling PgBouncer read replicas

2 Rate Limit External APIs

👉 Third-party API rate limits hit unexpectedly.

Example: 10x traffic → 10x payment API calls → hit Stripe rate limit → checkout fails.

Circuit breakers queue external calls upgrade API tiers

3 Memory Exhaustion

👉 Application servers run out of memory with increased load.

Example: Each user session = 10MB → 10K users = 100GB RAM → servers crash.

Stateless design external session storage horizontal scaling

4 CDN Origin Overload

👉 CDN cache misses overwhelm origin servers.

Example: Viral post links to uncached page → 100K users → all hit origin → origin crashes.

Cache warming higher CDN cache TTL origin rate limiting

5 Database Write Bottleneck

- 👉 Single primary database can't handle write traffic.

Example: 10K sign-ups/minute → single DB primary maxed out → sign-up fails.

Write sharding queue writes batch inserts

6 Logging/Monitoring Overload

- 👉 Excessive logging consumes resources.

Example: 100x traffic → 100x logs → disk full → app crashes trying to write logs.

Log sampling async logging log levels

\n

7 Auto-Scaling Lag

- 👉 Auto-scaling reacts too slowly to sudden spike.

Example: Traffic spikes in 30 seconds → auto-scaling takes 5 minutes to launch instances → app down in between.

Pre-warming faster scaling over-provisioning buffer

💡 Problem 117: Kubernetes Horizontal Pod Autoscaler (HPA)

INTERVIEWER: How does Kubernetes auto-scale pods based on load?

Question:

Solution:

HPA adjusts pod replicas based on metrics — CPU, memory, or custom metrics from application.

Here's how Kubernetes HPA works 👉

1 CPU-Based Scaling

- 👉 Scale pods when average CPU usage crosses threshold.

Example: Target 50% CPU → actual 80% → HPA adds pods to bring average down to 50%.

Resource requests metrics server

2 Memory-Based Scaling

- 👉 Scale based on memory utilization.

Example: Memory > 70% → add pods.

Memory metrics resource requests required

3 Custom Metrics from Application

- 👉 Scale based on application-specific metrics.

Example: Queue length > 100 → scale up workers → queue length < 20 → scale down.

Custom metrics API Prometheus adapter

4 Min and Max Replicas

- 👉 Set bounds to prevent under/over-provisioning.

Example: Min 2 replicas (high availability), max 50 replicas (cost control).

Replica limits safety bounds

5 Stabilization Window

- 👉 Wait before scaling down to prevent flapping.

Example: Load drops → wait 5 minutes → if still low → scale down.

Scale-down delay stability

6 Vertical Pod Autoscaler (VPA)

- 👉 Adjust resource requests/limits instead of replica count.

Example: Pod using 2GB RAM but requested 4GB → VPA adjusts to 2.5GB.

Resource optimization right-sizing

💡 Problem 118: Advanced HPA Scaling Strategies

INTERVIEWER: HPA scales based on CPU — but your app is memory-bound. How do you fix that?

Question:

Solution:

Advanced HPA configurations handle complex scaling scenarios beyond simple CPU metrics.

1 Memory-Based Scaling

👉 Scale based on memory utilization instead of CPU.

Example: scale when memory > 80% → prevents OOM kills.

Memory metrics resource limits

2 Custom Metrics (Prometheus)

👉 Scale on application-specific metrics.

Example: Scale based on queue depth, active connections, request latency from Prometheus.

Custom metrics Prometheus adapter

3 Multiple Metrics

👉 Combine multiple metrics for scaling decisions.

Example: Scale if CPU > 70% OR memory > 80% OR queue > 1000 messages.

Multiple metrics OR conditions

4 Scale-Up vs Scale-Down Policies

👉 Aggressive scale-up, conservative scale-down.

Example: Scale up immediately on spike, scale down slowly over 5 minutes.

Scale policies stabilization window

5 Min/Max Replicas

- 👉 Set bounds to prevent over-scaling or under-provisioning.

Example: Min 3 replicas (HA), max 50 replicas (cost control).

Replica limits cost control

6 Handling Cold Starts

- 👉 Pre-warm pods to handle sudden spikes.

Example: Keep 5 idle pods ready → instant capacity for traffic spike.

Pre-warming cold start

7 VPA (Vertical Pod Autoscaler) Alternative

- 👉 Adjust pod resources instead of replica count.

Example: App needs more memory → VPA increases memory limit → restart pod.

VPA resource adjustment



Problem 118: Implementing Blue-Green Deployments

INTERVIEWER: How do you deploy new code with zero downtime and instant rollback?

Question:

Solution:

Blue-green deployments enable zero-downtime releases with instant rollback capability.

1 Two Identical Environments

- 👉 Maintain two production environments: blue (current) and green (new).

Example: Blue environment serves traffic; Deploy new version to green environment.

Blue-green parallel environments

2 Switch Traffic with Load Balancer

- 👉 Flip traffic from blue to green instantly.

Example: Update ALB target group from blue to green → instant switch.

Load balancer traffic switch

3 Instant Rollback

- 👉 If issues detected, switch back to blue immediately.

Example: Green has bug → switch traffic back to blue in seconds.

Rollback safety

4 Testing Green Before Switch

- 👉 Verify green environment before sending production traffic.

Example: Run smoke tests against green → if pass, switch traffic.

Smoke tests verification

5 Database Migration Challenge

- 👉 Ensure schema compatible with both versions.

Example: Add column in backward-compatible way → deploy code → run migration.

Schema compatibility migrations

6 Cost Consideration

- 👉 Running two environments doubles infrastructure cost during deploy.

Example: 20 servers × 2 environments = 40 servers during deployment window.

Cost resource doubling

7 Canary as Alternative

👉 Gradually shift traffic instead of instant switch.

Example: Route 5% → 25% → 50% → 100% to new version → lower risk.

Canary deployment gradual rollout



Problem 118: Advanced HPA Scaling Strategies

INTERVIEWER: HPA scales based on CPU — but your app is memory-bound. How do you fix that?

Question:

Solution:

Advanced HPA configurations handle complex scaling scenarios beyond simple CPU metrics.

1 Memory-Based Scaling

👉 Scale based on memory utilization instead of CPU.

Example: scale when memory > 80% → prevents OOM kills.

Memory metrics resource limits

2 Custom Metrics (Prometheus)

👉 Scale on application-specific metrics.

Example: Scale based on queue depth, active connections, request latency from Prometheus.

Custom metrics Prometheus adapter

3 Multiple Metrics

- 👉 Combine multiple metrics for scaling decisions.

Example: Scale if CPU > 70% OR memory > 80% OR queue > 1000 messages.

Multiple metrics OR conditions

4 Scale-Up vs Scale-Down Policies

- 👉 Aggressive scale-up, conservative scale-down.

Example: Scale up immediately on spike, scale down slowly over 5 minutes.

Scale policies stabilization window

5 Min/Max Replicas

- 👉 Set bounds to prevent over-scaling or under-provisioning.

Example: Min 3 replicas (HA), max 50 replicas (cost control).

Replica limits cost control

6 Handling Cold Starts

- 👉 Pre-warm pods to handle sudden spikes.

Example: Keep 5 idle pods ready → instant capacity for traffic spike.

Pre-warming cold start

7 VPA (Vertical Pod Autoscaler) Alternative

- 👉 Adjust pod resources instead of replica count.

Example: App needs more memory → VPA increases memory limit → restart pod.

VPA resource adjustment

Problem 118: Implementing Blue-Green Deployments

INTERVIEWER: How do you deploy new code with zero downtime and instant rollback?

Question:

Solution:

Blue-green deployments enable zero-downtime releases with instant rollback capability.

1 Two Identical Environments

👉 Maintain two production environments: blue (current) and green (new).

Example: Blue environment serves traffic; Deploy new version to green environment.

Blue-green parallel environments

2 Switch Traffic with Load Balancer

👉 Flip traffic from blue to green instantly.

Example: Update ALB target group from blue to green → instant switch.

Load balancer traffic switch

3 Instant Rollback

👉 If issues detected, switch back to blue immediately.

Example: Green has bug → switch traffic back to blue in seconds.

Rollback safety

4 Testing Green Before Switch

👉 Verify green environment before sending production traffic.

Example: Run smoke tests against green → if pass, switch traffic.

Smoke tests verification

5 Database Migration Challenge

- 👉 Ensure schema compatible with both versions.

Example: Add column in backward-compatible way → deploy code → run migration.

Schema compatibility migrations

6 Cost Consideration

- 👉 Running two environments doubles infrastructure cost during deploy.

Example: $20 \text{ servers} \times 2 \text{ environments} = 40 \text{ servers}$ during deployment window.

Cost resource doubling

7 Canary as Alternative

- 👉 Gradually shift traffic instead of instant switch.

Example: Route 5% → 25% → 50% → 100% to new version → lower risk.

Canary deployment gradual rollout



CDN & Edge Caching



Topic 42: CDN & Edge Caching



💡 Problem 118: Using CDN for Static and Dynamic Content

INTERVIEWER: How do you reduce latency for users across the globe?

Question:

Solution:

CDN caches content at edge locations near users — reduces latency and origin server load.

Here's how CDN works 

1 Edge Locations Cache Content

👉 CDN servers in many geographic locations cache content.

Example: User in Tokyo requests image → served from Tokyo edge location → 10ms latency instead of 200ms from US origin.

CloudFront Cloudflare Akamai edge caching

2 Static Content Caching

👉 Images, CSS, JS, fonts served from CDN with long cache TTL.

Example: 'Cache-Control: max-age=31536000' → cached for 1 year.

Static assets immutable content cache forever

3 Dynamic Content Caching

👉 Cache API responses with shorter TTL.

Example: Product page cached for 5 minutes → most users hit cache → origin gets 1 request per 5 minutes.

API caching Cache-Control headers

4 Cache Invalidation

👉 Purge or invalidate CDN cache when content updates.

Example: Deploy new website version → invalidate CDN cache → users get new version immediately.

Cache purge versioned URLs cache busting

5 Origin Shield for Origin Protection

👉 Additional caching layer between edge and origin to reduce origin load.

Example: 100 edge locations → origin shield → origin. Edge requests consolidated at shield.

CloudFront origin shield cache hierarchy

6 Gzip/Brotli Compression

👉 CDN compresses responses before sending to clients.

Example: 1MB HTML file → compressed to 200KB → 5x bandwidth savings.

Compression Accept-Encoding faster delivery

💡 Problem 119: CDN Cache Control and Strategies

INTERVIEWER: How do you control what gets cached and for how long in a CDN?

Question:

Solution:

Use Cache-Control headers and cache keys to control CDN caching behavior.

Here's how to control CDN caching 👉

1 Cache-Control Header

👉 Tell CDN how long to cache content.

Example: 'Cache-Control: public, max-age=3600' → cache for 1 hour.

max-age s-maxage public/private

2 Vary Header for Multiple Versions

- 👉 Cache different versions based on request headers.

Example: 'Vary: Accept-Encoding' → cache gzip and non-gzip versions separately.

Vary header multiple cache entries

3 Query String in Cache Key

- 👉 Include or exclude query parameters from cache key.

Example: '/api/products?id=123' → cache per ID. '/api/products?utm_source=email' → ignore utm params.

Cache key configuration query string handling

4 Cookie-Based Caching

- 👉 Include cookies in cache key for personalized content.

Example: Logged-in users see personalized page → cache keyed by session cookie.

Personalization cache fragmentation

5 Stale-While-Revalidate

- 👉 Serve stale content while fetching fresh content in background.

Example: 'Cache-Control: max-age=60, stale-while-revalidate=300' → serve stale for 5 min while revalidating.

Improved availability eventual freshness

6 Bypass Cache for Dynamic Content

- 👉 Set 'Cache-Control: no-cache' for content that shouldn't be cached.

Example: User account page → no-cache → always fetched from origin.

No-cache dynamic content user-specific data

Problem 120: Response Compression Strategies

INTERVIEWER: Your API response is 2MB — how do you reduce bandwidth costs?

Question:



Solution:

Response compression reduces bandwidth and improves load times, especially on slow networks.

1 Gzip Compression

👉 Standard HTTP compression algorithm.

Example: 2MB JSON → gzip → 200KB → 90% size reduction.

Gzip deflate

2 Brotli Compression

👉 Modern algorithm with better compression than gzip.

Example: Same 2MB JSON → Brotli → 150KB → better than gzip's 200KB.

Brotli better compression

3 Content-Encoding Header

👉 Server indicates compression used via HTTP header.

Example: Content-Encoding: br → browser auto-decompresses Brotli.

HTTP headers Content-Encoding

4 Compression Level Trade-offs

👉 Higher compression = smaller size but more CPU.

Example: Brotli level 4 (fast) vs level 11 (max compression, slow).

CPU usage compression levels

5 Pre-compression for Static Assets

👉 Compress files at build time, serve pre-compressed.

Example: Webpack generates .js.gz and .js.br files → nginx serves pre-compressed.

Pre-compression static assets

6 Don't Compress Images/Video

👉 Already-compressed formats see minimal benefit.

Example: JPEG, PNG, MP4 already compressed → gzipping adds overhead, no gain.

Skip compression binary files

7 Selective Compression

👉 Compress only responses above size threshold.

Example: Only compress responses > 1KB → small responses not worth CPU overhead.

Threshold selective compression

💡 Problem 120: Response Compression Strategies

INTERVIEWER: Your API response is 2MB — how do you reduce bandwidth costs?

Question:

Solution:

Response compression reduces bandwidth and improves load times, especially on slow networks.

1 Gzip Compression

👉 Standard HTTP compression algorithm.

Example: 2MB JSON → gzip → 200KB → 90% size reduction.

Gzip deflate

2 Brotli Compression

- 👉 Modern algorithm with better compression than gzip.

Example: Same 2MB JSON → Brotli → 150KB → better than gzip's 200KB.

Brotli better compression

3 Content-Encoding Header

- 👉 Server indicates compression used via HTTP header.

Example: Content-Encoding: br → browser auto-decompresses Brotli.

HTTP headers Content-Encoding

4 Compression Level Trade-offs

- 👉 Higher compression = smaller size but more CPU.

Example: Brotli level 4 (fast) vs level 11 (max compression, slow).

CPU usage compression levels

5 Pre-compression for Static Assets

- 👉 Compress files at build time, serve pre-compressed.

Example: Webpack generates .js.gz and .js.br files → nginx serves pre-compressed.

Pre-compression static assets

6 Don't Compress Images/Video

- 👉 Already-compressed formats see minimal benefit.

Example: JPEG, PNG, MP4 already compressed → gzipping adds overhead, no gain.

Skip compression binary files

\n

7 Selective Compression

- 👉 Compress only responses above size threshold.

Example: Only compress responses > 1KB → small responses not worth CPU overhead.

Threshold selective compression

💡 Problem 120: Reducing Global Latency from 3s to <200ms

INTERVIEWER: US loads in 150ms, India in 3 seconds — how do you fix global latency?

Question:

Solution:

Global performance requires geo-distributed infrastructure and edge caching.

1 Multi-Region Deployment

- 👉 Deploy application servers in multiple geographic regions.

Example: Servers in us-east, eu-west, ap-southeast → users routed to nearest region.

AWS regions GCP multi-region geo-routing

2 CDN with Edge Caching

- 👉 Cache static and dynamic content at edge locations worldwide.

Example: India users hit Mumbai edge location (10ms) instead of US origin (300ms).

CloudFront Cloudflare Akamai

3 GeoDNS Routing

- 👉 DNS routes users to geographically closest servers.

Example: Route53 routes Indian users to ap-south servers, US users to us-east servers.

Route53 geolocation traffic policies

4 Database Read Replicas Per Region

- 👉 Place read replicas in each region for faster data access.

Example: India replica serves Indian users → 20ms DB latency vs 300ms to US primary.

Cross-region replication Aurora Global Database

5 Asset Optimization

- 👉 Compress images, minify JS/CSS, use modern formats.

Example: 5MB images → WebP compression → 500KB → 10x faster download over slow connections.

Image optimization WebP Brotli compression

6 Connection Optimization

- 👉 Use HTTP/2, connection pooling, and keep-alive.

Example: HTTP/1.1 = 6 parallel connections; HTTP/2 = multiplexed → faster page loads.

HTTP/2 HTTP/3 connection pooling

7 Pre-loading Critical Resources

- 👉 Use resource hints to load critical assets early.

Example: for fonts, critical CSS → loaded before parser reaches them.

Resource hints critical rendering path

Problem 121: Load Testing Before Production

INTERVIEWER: How do you know your system can handle Black Friday traffic?

Question:



Solution:

Load testing validates system capacity and identifies bottlenecks before they cause outages.

1 Load Testing Tools

- 👉 Simulate realistic user traffic patterns.

Example: k6, JMeter, Gatling, Locust → simulate 10K concurrent users.

k6 JMeter Gatling Locust

2 Realistic Traffic Patterns

- 👉 Model actual user behavior, not just constant load.

Example: Ramp up over 10 mins → hold 10K users for 30 mins → ramp down.

Traffic patterns ramp-up

3 Soak Testing

- 👉 Run sustained load to find memory leaks and resource exhaustion.

Example: Run 5K users for 24 hours → detect slow memory leak in session storage.

Soak testing memory leaks

4 Spike Testing

- 👉 Sudden traffic bursts to test auto-scaling.

Example: 1K users → spike to 20K in 1 minute → verify auto-scaling responds in time.

Spike testing auto-scaling

5 Identify Bottlenecks

- 👉 Monitor all components during load test.

Example: Database CPU at 90%, app servers at 40% → database is bottleneck.

Bottleneck analysis monitoring

6 Test in Production-Like Environment

- 👉 Staging should match production capacity.

Example: Prod has 10 servers → staging has 10 servers → realistic results.

Staging production parity

7 Continuous Load Testing

- 👉 Automate load tests in CI/CD pipeline.

Example: Every deploy → run 5-minute load test → fail build if p95 latency > 500ms.

CI/CD automated testing

💡 Problem 122: Practicing Chaos Engineering

INTERVIEWER: How do you know your system will survive a production failure before it happens?

Question:

Solution:

Chaos engineering proactively tests system resilience by intentionally injecting failures.

1 Start with Hypotheses

- 👉 Define what you expect to happen during failure.

Example: 'If one AZ goes down, app should continue serving with <1% error rate.'

Hypothesis expected behavior

2 Chaos Monkey (Random Instance Termination)

- 👉 Randomly kill instances to test auto-healing.

Example: Chaos Monkey terminates 1 instance every hour → verify auto-scaling replaces it.

Chaos Monkey Netflix instance failure

3 Network Latency Injection

- 👉 Add artificial latency to test timeout handling.

Example: Inject 5s delay to database calls → verify circuit breaker opens.

Latency injection timeout testing

4 Dependency Failure Testing

- 👉 Simulate external service outages.

Example: Block payment API → verify graceful degradation → queue payments for later.

Dependency failure graceful degradation

5 Resource Exhaustion

- 👉 Consume CPU/memory to test resource limits.

Example: Fill disk to 95% → verify alerting triggers and app handles gracefully.

Resource limits capacity testing

6 Chaos in Production (Carefully)

- 👉 Run controlled chaos experiments in production.

Example: Start with off-peak hours, single AZ → gradually increase scope.

Production chaos controlled experiments

7 Continuous Chaos

- 👉 Automate chaos experiments as part of CI/CD.

Example: Every deploy → run chaos suite → fail build if resilience tests fail.

Automated chaos

CI/CD integration



💡 Problem 121: Load Testing Before Production

INTERVIEWER: How do you know your system can handle Black Friday traffic?

Question:



Solution:

Load testing validates system capacity and identifies bottlenecks before they cause outages.

1 Load Testing Tools

- 👉 Simulate realistic user traffic patterns.

Example: k6, JMeter, Gatling, Locust → simulate 10K concurrent users.

k6

JMeter

Gatling

Locust



2 Realistic Traffic Patterns

- 👉 Model actual user behavior, not just constant load.

Example: Ramp up over 10 mins → hold 10K users for 30 mins → ramp down.

Traffic patterns

ramp-up



3 Soak Testing

- 👉 Run sustained load to find memory leaks and resource exhaustion.

Example: Run 5K users for 24 hours → detect slow memory leak in session storage.

Soak testing memory leaks

4 Spike Testing

- 👉 Sudden traffic bursts to test auto-scaling.

Example: 1K users → spike to 20K in 1 minute → verify auto-scaling responds in time.

Spike testing auto-scaling

5 Identify Bottlenecks

- 👉 Monitor all components during load test.

Example: Database CPU at 90%, app servers at 40% → database is bottleneck.

Bottleneck analysis monitoring

6 Test in Production-Like Environment

- 👉 Staging should match production capacity.

Example: Prod has 10 servers → staging has 10 servers → realistic results.

Staging production parity

7 Continuous Load Testing

- 👉 Automate load tests in CI/CD pipeline.

Example: Every deploy → run 5-minute load test → fail build if p95 latency > 500ms.

CI/CD automated testing

💡 Problem 122: Practicing Chaos Engineering

INTERVIEWER: How do you know your system will survive a production failure before it happens?

Question:**Solution:**

Chaos engineering proactively tests system resilience by intentionally injecting failures.

1 Start with Hypotheses

- 👉 Define what you expect to happen during failure.

Example: 'If one AZ goes down, app should continue serving with <1% error rate.'

Hypothesis expected behavior

2 Chaos Monkey (Random Instance Termination)

- 👉 Randomly kill instances to test auto-healing.

Example: Chaos Monkey terminates 1 instance every hour → verify auto-scaling replaces it.

Chaos Monkey Netflix instance failure

3 Network Latency Injection

- 👉 Add artificial latency to test timeout handling.

Example: Inject 5s delay to database calls → verify circuit breaker opens.

Latency injection timeout testing

4 Dependency Failure Testing

- 👉 Simulate external service outages.

Example: Block payment API → verify graceful degradation → queue payments for later.

Dependency failure graceful degradation

5 Resource Exhaustion

- 👉 Consume CPU/memory to test resource limits.

Example: Fill disk to 95% → verify alerting triggers and app handles gracefully.

Resource limits capacity testing

6 Chaos in Production (Carefully)

- 👉 Run controlled chaos experiments in production.

Example: Start with off-peak hours, single AZ → gradually increase scope.

Production chaos controlled experiments

\n\n

7 Continuous Chaos

- 👉 Automate chaos experiments as part of CI/CD.

Example: Every deploy → run chaos suite → fail build if resilience tests fail.

Automated chaos CI/CD integration

💡 Problem 120: Reducing Global Latency from 3s to <200ms

INTERVIEWER: US loads in 150ms, India in 3 seconds — how do you fix global latency?

Question:

(Green box placeholder for question text)

Solution:

Global performance requires geo-distributed infrastructure and edge caching.

1 Multi-Region Deployment

- 👉 Deploy application servers in multiple geographic regions.

Example: Servers in us-east, eu-west, ap-southeast → users routed to nearest region.

AWS regions GCP multi-region geo-routing

2 CDN with Edge Caching

- 👉 Cache static and dynamic content at edge locations worldwide.

Example: India users hit Mumbai edge location (10ms) instead of US origin (300ms).

CloudFront Cloudflare Akamai

3 GeoDNS Routing

- 👉 DNS routes users to geographically closest servers.

Example: Route53 routes Indian users to ap-south servers, US users to us-east servers.

Route53 geolocation traffic policies

4 Database Read Replicas Per Region

- 👉 Place read replicas in each region for faster data access.

Example: India replica serves Indian users → 20ms DB latency vs 300ms to US primary.

Cross-region replication Aurora Global Database

5 Asset Optimization

- 👉 Compress images, minify JS/CSS, use modern formats.

Example: 5MB images → WebP compression → 500KB → 10x faster download over slow connections.

Image optimization WebP Brotli compression

6 Connection Optimization

- 👉 Use HTTP/2, connection pooling, and keep-alive.

Example: HTTP/1.1 = 6 parallel connections; HTTP/2 = multiplexed → faster page loads.

HTTP/2 HTTP/3 connection pooling

\n

7 Pre-loading Critical Resources

- 👉 Use resource hints to load critical assets early.

Example: for fonts, critical CSS → loaded before parser reaches them.

Resource hints critical rendering path



💡 Problem 120: Edge Computing and Lambda@Edge

INTERVIEWER: How can you run code at CDN edge locations?

Question:



Solution:

Edge computing runs code close to users — enables personalization, A/B testing, authentication at edge.

Here's how edge computing works 🤖

1 Lambda@Edge for Request/Response Manipulation

- 👉 Run code on CloudFront edge to modify requests/responses.

Example: Viewer request → Lambda@Edge adds authentication check → forwards to origin or returns 403.

CloudFront Functions Lambda@Edge edge logic



2 A/B Testing at Edge

- 👉 Route users to different origins based on cookies or randomization.

Example: 50% of users see version A, 50% see version B → routed at edge.

A/B testing traffic splitting edge decision



3 URL Rewriting and Redirects

- 👉 Modify URLs or redirect at edge without hitting origin.

Example: '/old-url' → edge function redirects to '/new-url'.

URL rewrite redirect edge routing

4 Bot Detection and Security

- 👉 Block malicious requests at edge before reaching origin.

Example: Edge function checks user-agent → blocks known bots → protects origin.

WAF bot detection DDoS protection

5 Personalization with Edge Functions

- 👉 Customize content based on location, device, or headers.

Example: User in France → edge function returns French content → no origin request.

Personalization localization edge logic

6 Edge Compute Limitations

- 👉 Limited execution time, memory, and no persistent storage.

Example: Lambda@Edge max 5 seconds, 128-256MB → lightweight logic only.

Constraints stateless cold starts

📚 Asynchronous Processing

Topic 43: Asynchronous Processing

💡 Problem 121: Background Job Processing

INTERVIEWER: How do you handle long-running tasks without blocking the user?

Question:

Solution:

Move long tasks to background workers — return immediately to user, process asynchronously.

Here's how background job processing works 🤝

1 Job Queue Pattern

👉 Enqueue job, worker picks up and processes asynchronously.

Example: User uploads video → return success immediately → background worker transcodes video.

RabbitMQ AWS SQS Redis Queue

2 Worker Pool for Parallel Processing

👉 Multiple workers process jobs concurrently.

Example: 100 jobs in queue, 10 workers → each worker processes 10 jobs in parallel.

Horizontal scaling worker instances

3 Job Retries with Exponential Backoff

👉 Retry failed jobs with increasing delay.

Example: Job fails → retry after 1 second → fails again → retry after 2 seconds → 4 seconds → 8 seconds.

Retry mechanism exponential backoff max retries

4 Job Status Tracking

👉 Store job status so user can check progress.

Example: Job ID returned to user → user polls '/jobs/123/status' → returns "processing", "completed", or "failed".

Status API polling webhooks

5 Priority Queues

👉 High-priority jobs processed before low-priority.

Example: Paid users → high priority queue, free users → low priority queue.

Priority levels SLA-based processing

6 Dead Letter Queue for Failed Jobs

👉 Jobs that fail max retries moved to DLQ for investigation.

Example: Job fails 5 times → moved to DLQ → manual inspection.

DLQ error handling manual intervention

💡 Problem 122: Webhook and Callback Patterns

INTERVIEWER: How do you notify clients when async jobs complete?

Question:

Solution:

Use webhooks to push results to clients — avoid polling overhead.

Here's how webhook patterns work 👉

1 Webhook Registration

- 👉 Client provides callback URL when creating job.

Example: POST `/jobs` with `{"callback_url": "https://client.com/webhook"}` → job completes → POST to callback URL.

[Callback URLs](#) [push notifications](#)

2 Webhook Retries

- 👉 Retry webhook delivery if client is unavailable.

Example: POST to webhook fails → retry after 1 min → 5 min → 15 min → mark as failed.

[Retry policy](#) [exponential backoff](#)

3 Webhook Signatures for Security

- 👉 Sign webhook payload so client can verify authenticity.

Example: HMAC-SHA256(payload + secret) → include in header → client verifies signature.

[Signature verification](#) [security](#)

4 Polling as Fallback

- 👉 Allow clients to poll job status if webhook delivery fails.

Example: Client polls `/jobs/123/status` every 5 seconds until completion.

[Polling](#) [fallback mechanism](#)

5 Webhook Ordering

- 👉 Include sequence number or timestamp for ordering.

Example: Multiple webhooks for same job → sequence numbers ensure correct order.

[Event ordering](#) [sequence numbers](#)

6 Idempotent Webhook Processing

- 👉 Client must handle duplicate webhook deliveries.

Example: Webhook delivered twice → client processes idempotently → no duplicate side effects.

Idempotency duplicate handling

Bloom Filters

Topic 44: Bloom Filters

Problem 123: Understanding Bloom Filters

INTERVIEWER: How do you check if an element might be in a set without storing the entire set?

Question:

Solution:

Bloom filters are space-efficient probabilistic data structures — can have false positives, never false negatives.

Here's how Bloom filters work 

1 What is a Bloom Filter?

 Bit array with multiple hash functions — space-efficient set membership test.

Example: Store 1M URLs in 1MB Bloom filter instead of 100MB hash set.

Probabilistic space-efficient false positives possible

2 False Positives, No False Negatives

- 👉 Bloom filter might say "maybe in set" but never wrong about "not in set".

Example: Check if URL visited → "definitely not visited" (accurate) or "probably visited" (might be false positive).

False positive rate no false negatives

3 How It Works

- 👉 Hash element with K hash functions → set K bits in bit array.

Example: Hash "hello" with 3 functions → set bits 5, 23, 47 → to check membership, verify all 3 bits are set.

Multiple hash functions bit array

4 Trade-off: Size vs Accuracy

- 👉 Larger bit array = lower false positive rate, more hash functions = lower rate.

Example: 1% false positive rate needs 9.6 bits per element, 0.1% needs 14.4 bits.

Size calculation accuracy tuning

5 Can't Delete Elements

- 👉 Clearing a bit might affect other elements.

Example: Deleting one element clears bits → might cause false negatives for other elements.

Immutable counting Bloom filter variant

6 Use Cases

- 👉 Cache filters, spell checkers, malicious URL detection, database query optimization.

Example: Check Bloom filter before expensive database query → skip query if definitely not in DB.

Optimization pre-filter cache efficiency

Congratulations! You've completed all backend problems from ebook!

mission_compile