

SYSTEM DESIGN & ARCHITECTURE



Design Principles, Patterns, and
Real-World Architectures for
Modern Systems



Message to Developers !!

Hey Techies,

Over the years, I've seen many talented developers build great code but struggle when it comes to *designing systems*. Writing code is an art — but designing a scalable, reliable, and maintainable system is *architecture*. It's what transforms a good developer into a great engineer.

In this book, I've curated the most essential **system design principles, patterns, and real-world concepts** that are not only *commonly asked in interviews* but also *crucial for building production-grade systems*. Each topic is simplified, structured, and supported with practical insights drawn from real projects — the kind of lessons that go beyond textbooks and straight into what industry experts use every day.

Whether you're preparing for a **System Design interview**, aiming to **improve your architectural thinking**, or simply **want to understand how large-scale systems are built**, this book is for you. Don't just read it — *walk through it, sketch the diagrams, imagine the data flow*, and see how each decision impacts scalability, performance, and reliability.

System design is not about memorizing answers — it's about understanding *why* and *how* things work together. Once you master that mindset, you'll design solutions that don't just work — they *last*.

Keep learning. Keep building.
And most importantly — keep *designing with intent*.

– Amitesh Kumar Ray

Who Is This Book For?

This book is for **every software professional and learner** who wants to design better systems.

If you're a **Software Engineer**, it will help you go beyond coding and understand how real systems are built, scaled, and maintained.

If you're a **Designer or Architect**, it will sharpen your ability to make informed design choices and balance scalability, performance, and simplicity.

If you're an **Engineering Manager**, it will give you a framework to guide your team in making sound architectural decisions.

And if you're an **Engineering Student**, it will bridge the gap between classroom concepts and real-world software design.

If you want to design a BETTER system, you must know what's inside this book.

– Amitesh Kumar Ray

About the Author

Amitesh Kumar Ray is a seasoned **Software Engineer** with over a decade of experience in building scalable, high-performance systems. He currently serves as a **Lead SDE at M2P Fintech**, where he leads product design and architecture initiatives, guiding teams to create resilient and distributed enterprise systems.

With deep expertise in **Java, Spring Boot, Microservices, and System Architecture**, Amitesh has been instrumental in transforming large-scale applications into modern, scalable solutions. His practical, hands-on approach to engineering reflects years of experience designing systems from the ground up.

A passionate **author and mentor**, Amitesh has written **15+ programming books** on topics including Java, functional programming, microservices, and prompt engineering. He has conducted **100+ tech sessions and workshops**, helping developers bridge the gap between theory and real-world application.

He holds a **certification in Software Design and Architecture** from the **University of Alberta** and is currently **pursuing an 11-month Advanced AI and ML Certificate Program from IIT Madras (IITM)**, reflecting his commitment to continuous learning and innovation.

“True engineering lies not just in writing code, but in designing systems that evolve gracefully over time.”

Connect with him on [LinkedIn](#) to explore his latest work and insights on modern software design and architecture.

– Amitesh Kumar Ray

Chapter 1: Fundamentals of System Design

- 1.1 Fundamentals of System Design
 - 1.2 Scalability
 - 1.3 Availability
 - 1.4 Reliability
 - 1.5 Latency
 - 1.6 Throughput
 - 1.7 Performance vs Efficiency
 - 1.8 Fault Tolerance
 - 1.9 Resiliency and Recovery
 - 1.10 CAP Theorem
 - 1.11 ACID vs BASE
 - 1.12 Durability and Consistency Models
 - 1.13 Data Partitioning and Replication Basics
-

Chapter 2: Core Infrastructure Components

- 2.1 Load Balancer (L4/L7)
 - 2.2 CDN (Content Delivery Network)
 - 2.3 Proxy Server & Reverse Proxy
 - 2.4 API Gateway
 - 2.5 Message Queue (Kafka, RabbitMQ, SQS)
 - 2.6 Rate Limiting and Throttling
 - 2.7 Circuit Breaker & Bulkhead Pattern
 - 2.8 Service Discovery (Eureka, Consul, Zookeeper)
 - 2.9 Caching (Client, Server, Distributed)
 - 2.10 Cache Strategies: Read-Through, Write-Through, Write-Back
 - 2.11 Consistent Hashing
 - 2.12 Heartbeats & Health Checks
 - 2.13 Checksum and Data Validation
 - 2.14 Consensus Algorithms (Raft, Paxos, ZAB)
 - 2.15 Gossip Protocol
 - 2.16 Bloom Filter
 - 2.17 Leader Election
 - 2.18 Distributed Locking
-

Chapter 3: Database and Storage Design

- 3.1 Databases (RDBMS, NoSQL)
 - 3.2 SQL vs NoSQL
 - 3.3 DB Indexing and Query Optimization
 - 3.4 Database Sharding and Partitioning
 - 3.5 Database Replication
 - 3.6 Normalization vs Denormalization
 - 3.7 Read Replica and Write Master Design
 - 3.8 Transaction Management and Isolation Levels
 - 3.9 Distributed Transactions
 - 3.10 Data Modeling Techniques
 - 3.11 Caching with Databases (Redis, Memcached)
 - 3.12 Event Sourcing & CQRS (for complex data workflows)
-

Chapter 4: System Design Trade-offs

- 4.1 Vertical vs Horizontal Scaling
 - 4.2 Strong vs Eventual Consistency
 - 4.3 Stateful vs Stateless Design
 - 4.4 Synchronous vs Asynchronous Communication
 - 4.5 Batch vs Stream Processing
 - 4.6 Long Polling vs WebSockets vs Server-Sent Events
 - 4.7 Read-Through vs Write-Through vs Write-Back Caching
 - 4.8 REST vs RPC (gRPC, Thrift)
 - 4.9 SQL vs NoSQL Trade-offs
 - 4.10 Normalization vs Denormalization
 - 4.11 TCP vs UDP
 - 4.12 Push vs Pull Architecture
 - 4.13 Throughput vs Latency Trade-off
 - 4.14 Cost vs Performance Optimization
-

Chapter 5: Architectural Patterns

- 5.1 Client–Server Architecture
 - 5.2 Layered (N-tier) Architecture
 - 5.3 Microservices Architecture
 - 5.4 Monolithic Architecture (and evolution to microservices)
 - 5.5 Event-Driven Architecture
 - 5.6 Serverless Architecture
 - 5.7 Peer-to-Peer Architecture
 - 5.8 Service-Oriented Architecture (SOA)
 - 5.9 Hexagonal (Ports and Adapters) Architecture
 - 5.10 Domain-Driven Design (DDD)
 - 5.11 CQRS and Event Sourcing Pattern
 - 5.12 Saga Pattern (Transaction Management in Microservices)
 - 5.13 Bulkhead & Circuit Breaker Pattern
 - 5.14 Sidecar and Ambassador Patterns
 - 5.15 Strangler Fig Pattern (for legacy migration)
 - 5.16 API Gateway Pattern
-

Chapter 6: Distributed Systems & Communication

- 6.1 Fundamentals of Distributed Systems
 - 6.2 Distributed Caching
 - 6.3 Consensus & Coordination (Zookeeper, etcd)
 - 6.4 Gossip Protocol in Node Communication
 - 6.5 Consistent Hashing for Node Distribution
 - 6.6 Quorum Reads/Writes
 - 6.7 Idempotency and Retry Mechanisms
 - 6.8 Leader Election and Replication
 - 6.9 Eventual Consistency in Distributed Systems
 - 6.10 Heartbeats, Health Checks, Failover Mechanisms
 - 6.11 Split-Brain Handling
 - 6.12 Distributed Logging and Monitoring
 - 6.13 Message Ordering and Exactly-once Semantics
-

Chapter 7: Performance, Reliability, and Observability

- 7.1 Fault Detection and Recovery
 - 7.2 Redundancy and Failover Mechanisms
 - 7.3 Circuit Breaker and Bulkhead Design
 - 7.4 Graceful Degradation
 - 7.5 Autoscaling (Reactive Scaling)
 - 7.6 Load Shedding and Backpressure
 - 7.7 Observability (Logs, Metrics, Traces)
 - 7.8 Monitoring Stack (Prometheus, Grafana, ELK, OpenTelemetry)
 - 7.9 Alerting and Incident Management
 - 7.10 SLO, SLI, and SLA
 - 7.11 Blue-Green and Canary Deployment Patterns
 - 7.12 Capacity Planning and Cost Optimization
-

Chapter 8: Implementation in Modern Tech Stack

- 8.1 Designing Scalable Systems with Java and Spring Boot
 - 8.2 Microservices with Spring Cloud (Eureka, Config Server, Gateway, Feign)
 - 8.3 Resilience with Resilience4j and Hystrix
 - 8.4 Containerization with Docker
 - 8.5 Kubernetes for Orchestration and Autoscaling
 - 8.6 Service Mesh (Istio, Linkerd)
 - 8.7 CI/CD Pipeline (Jenkins, GitHub Actions, ArgoCD)
 - 8.8 SonarQube for Code Quality and Static Analysis
 - 8.9 Distributed Tracing (Sleuth, Zipkin, Jaeger)
 - 8.10 Security and Authentication (OAuth2, JWT, mTLS)
 - 8.11 Infrastructure as Code (Terraform, Helm)
-

Chapter 9: Real-World System Design Examples

- 9.1 Design a URL Shortener
 - 9.2 Design a Social Media Feed System
 - 9.3 Design a Ride-Sharing Platform (Uber)
 - 9.4 Design a Chat Application (WhatsApp)
 - 9.5 Design a Video Streaming Platform (YouTube/Netflix)
 - 9.6 Design an E-Commerce System
 - 9.7 Design a Payment Gateway
 - 9.8 Design a Notification/Alert System
 - 9.9 Design a Logging and Monitoring System
 - 9.10 Design a Distributed File Storage (Google Drive / Dropbox Style)
-

A Note at the End

You've reached the final page — but this isn't the end. It's the beginning of your journey as a system thinker.

System design isn't about memorizing patterns — it's about understanding *why* things work and *how* to make them better. Every line of code, every design choice, and every trade-off teaches you something deeper about building systems that last.

Keep learning. Keep experimenting. Keep designing.

Because true engineers don't just build systems — they **create impact**.

"Don't just aim to be a better coder. Aim to be a better designer of solutions."

– Amitesh Kumar Ray

